

# Lecture 03: Concurrency Decomposition

Vivek Kumar

Computer Science and Engineering

IIIT Delhi

[vivekk@iiitd.ac.in](mailto:vivekk@iiitd.ac.in)

# Last Lecture

- Why multicores?
  - Moore's law and Dennard scaling
  - Multicore saves power
- Pthread based implementation of parallel Fibonacci
- Wait/Notify sequence using pthread

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x = fib(n-1);
        uint64_t y = fib(n-2);
        return (x + y);
    }
}

typedef struct {
    uint64_t input;
    uint64_t output;
} thread_args;

void *thread_func(void *ptr) {
    uint64_t i =
        ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    uint64_t result;

    if (argc < 2) { return 1; }
    uint64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);



        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate.
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRIu64 " is %" PRIu64 ".\n",
          n, result);
    return 0;
}
```

# Today's Class

- Decomposition of sequential program into parallel program
  - Tasks and decomposition
  - Amdahl's law
  - Tasks and mapping
  - Decomposition techniques
    - Recursive
    - Data
    - Exploratory
    - Speculative

# Concurrency v/s Parallelism

- Concurrency
  - “**Dealing**” with lots of things at once
- Parallelism
  - “**Doing**” with lots of things at once



# Concurrency v/s Parallelism

- Concurrency

- Refers to tasks that appear to be running simultaneously, but which may, in fact, actually be running serially

- Parallelism


- Refers to concurrent tasks that actually run at the same time
- Always implies multiple processors
- Parallel tasks always run concurrently, but not all concurrent tasks are parallel

# Recipe to Solve a Problem using Parallel Programming

- **Typical** steps for constructing a parallel algorithm
  - identify what pieces of work can be performed concurrently
  - partition concurrent work onto independent processors
  - distribute a program's input, output, and intermediate data
  - coordinate accesses to shared data: avoid conflicts
  - ensure proper order of work using synchronization
- Why “**typical**”? Some of the steps may be omitted.
  - if data is in shared memory, distributing it may be unnecessary
  - the mapping of work to processors can be done statically by the programmer or dynamically by the runtime


# Decomposing Work for Parallel Execution

- Divide work into tasks that can be executed concurrently
- Many different decompositions possible for any computation
- Tasks may be same, different, or even indeterminate sizes
- Tasks may be independent or have non-trivial order
  - Conceptualize tasks and ordering as computation graph
    - Node = task
    - Edge = control dependency



# Example: Dense Matrix Vector Product


- Computing each element of output vector  $y$  is independent
- Easy to decompose dense matrix-vector product into tasks
  - one per element in  $y$
- Observations
  - task size is uniform
- no control dependences between tasks
  - tasks share  $b$





# Granularity of Task Decomposition


- Granularity = task size
  - depends on the number of tasks
- Fine-grain = large number of tasks
- Coarse-grain = small number of tasks
- Granularity examples for dense matrix-vector multiply
  - fine-grain: each task represents an individual element in  $y$
  - coarser-grain: each task computes 3 elements in  $y$



# Critical Path


- Edge in computation graph represents task serialization
- Critical path = longest weighted path through graph
- Critical path length = lower bound on parallel execution time

# Critical Path Length



Note: number in vertex represents task cost

# Critical Path Length



Note: number in vertex represents task cost

## Questions:

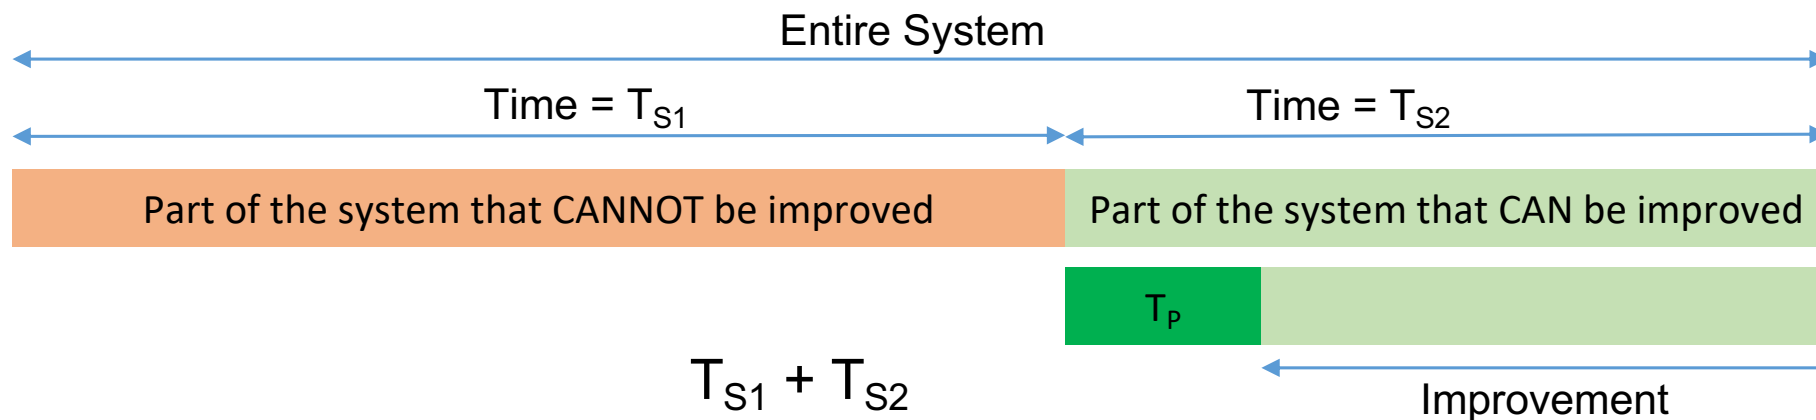
- What are the tasks on the critical path for each dependency graph?
- What is the shortest parallel execution time for each decomposition?

# Limits on Parallel Performance

- What bounds parallel execution time?
  - minimum task granularity
    - e.g. dense matrix-vector multiplication  $\leq n^2$  concurrent tasks
  - dependencies between tasks
  - parallelization overheads
    - e.g., cost of communication between tasks
  - fraction of application work that can't be parallelized
    - Amdahl's law

# Amdahl's Law

- Gives an estimate of maximum expected improvement  $S$  to an overall system when only part of the system  $F_E$  is improved by a factor  $F_I$



$$\text{Speedup using } P \text{ processes} = \frac{T_{S1} + T_{S2}}{T_{S1} + T_P}$$

# Speedup Analysis

1. Do disproportionately less work
2. Harness disproportionately more resources


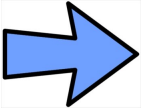


Fig. source: <http://www.drdobbs.com/cpp/going-superlinear/206100542>

# Today's Class

- Decomposition of sequential program into parallel program
  - Tasks and decomposition
  - Amdahl's law
  -  ○ Tasks and mapping
  - Decomposition techniques
    - Recursive
    - Data
    - Exploratory
    - Speculative




# Mapping Tasks to Cores

- Generally

- # of tasks > # threads available
- parallel algorithm must map tasks to threads
- schedule independent tasks on separate threads (consider computation graph)
- threads should have minimum interaction with one another

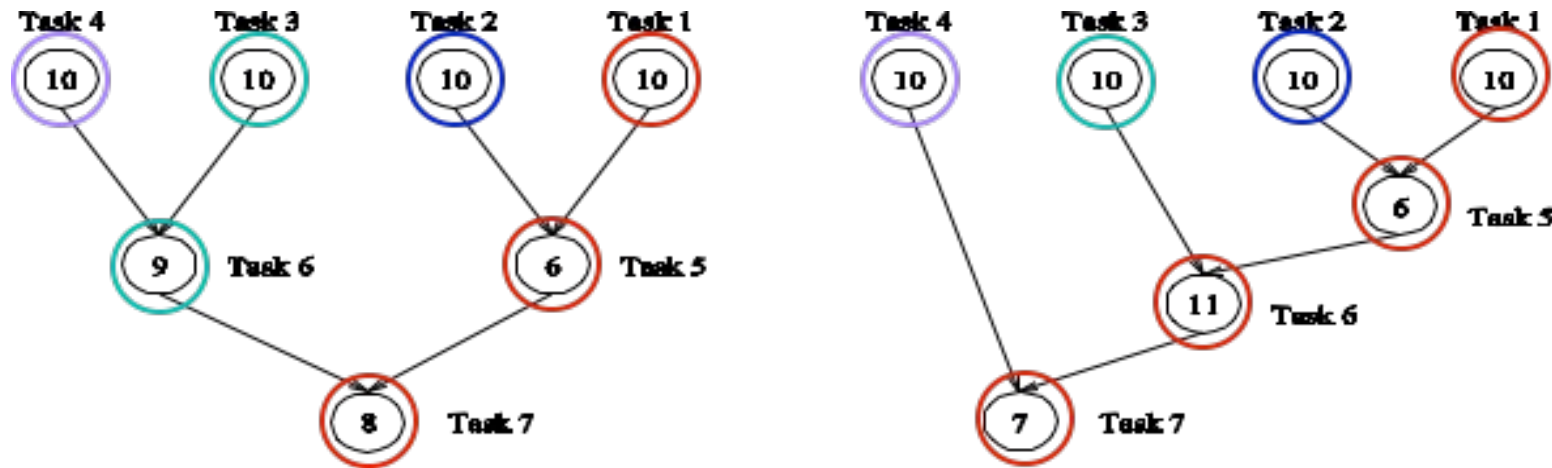
# Tasks, Threads, and Mapping Example



Note: number in vertex represents task cost

- How to best map these tasks on threads?

# Tasks, Threads, and Mapping Example



- No tasks in a level depend upon each other
- Assign all tasks within a level to different threads

# Mapping Techniques

## Static vs. dynamic mappings

- Static mapping
  - *a-priori* mapping of tasks to threads or processes
    - requirements
      - a good estimate of task size
      - even so, computing an optimal mapping may be hard
- Dynamic mapping
  - map tasks to threads or processes at runtime
  - why?
    - tasks are generated at runtime, or
    - their sizes are unknown

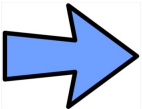
# Static Mapping

- Data partitioning
- Computation graph partitioning

# Dynamic Mapping

- Dynamic mapping AKA dynamic load balancing
  - load balancing is the primary motivation for dynamic mapping
- Styles
  - centralized
  - distributed

# Today's Class

- Decomposition of sequential program into parallel program
  - Tasks and decomposition
  - Amdahl's law
  - Tasks and mapping
  -  ○ Decomposition techniques
    - Recursive
    - Data
    - Exploratory
    - Speculative

# Decomposition Techniques

*How should one decompose a task into various subtasks?*

- No single universal recipe
- In practice, a variety of techniques are used including
  - Data decomposition
  - Recursive decomposition
  - Exploratory decomposition
  - Speculative decomposition



# Data Decomposition


- Steps

1. identify the data on which computations are performed
2. partition the data across various tasks
  - partitioning induces a decomposition of the problem


# Data Decomposition Example

- If each element of the output can be computed independently
- Partition the output data across tasks
- Have each task perform the computation for its outputs

**Example:  
dense matrix-vector  
multiply**



# Recursive Decomposition



**Question:** what kind of mapping is suited for this scenario?

```
uint64_t fib(uint64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        uint64_t x = fib(n-1);  
        uint64_t y = fib(n-2);  
        return (x + y);  
    }  
}
```

DAG Source: <http://www.cs.ucsb.edu/projects/jicos/tutorial/fibonacci/index.html>


# Exploratory Decomposition

- Exploration (search) of a state space of solutions
  - Problem decomposition reflects shape of execution
  - Parallel formulation may perform a different amount of work



# Exploratory Decomposition Speedup

- Parallel formulation may perform a different amount of work
  - **Can cause super- or sub-linear speedup**
- Assume each vertex of the triangles represents a computation that takes 'T' unit of time to compute and execution begins from leftmost triangle to the rightmost



- Serial execution time =  $7T$
- Parallel execution time using 4 threads to compute each triangle in parallel =  $T$
- Speedup (4 threads) =  $7T/T = 7$
- **Super-linear speedup**

- Serial execution time =  $3T$
- Parallel execution time using 4 threads to compute each triangle in parallel =  $3T$
- Speedup (4 threads) =  $3T/3T = 1$
- **Sub-linear speedup**

# Question

- How exploratory decomposition (ED) differs from data decomposition (DD)?
  1. Unlike ED, **all** partial tasks contribute to final result in DD
  2. Unlike DD, unfinished tasks in ED can be terminated once final solution is found

# Speculative Decomposition

- Example: when program may take one of many possible compute-intensive branches depending on the output of preceding computation

```
int val = T1      //compute intensive
switch(val) {    // cases may be computed speculatively
    case 0: T2; break;
    case 1: T3; break;
    .....
    case n: Tn; break;
}
```

# Next Lecture (#04)

- Productivity in parallel programming (tasks based parallel programming model)
- Quiz-1
  - Syllabus: Lectures 02 and 03