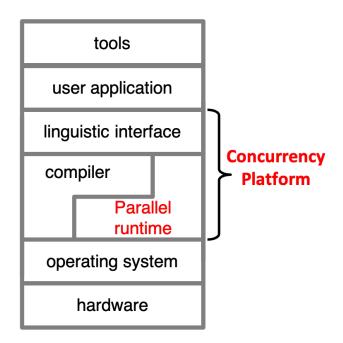
Lecture 05: Performance in Parallel Programming

Vivek Kumar
Computer Science and Engineering
IIIT Delhi
vivekk@iiitd.ac.in



Last Lecture



```
async { Wash your clothes in washing machine

async { Complete your PRMP project deadline

async { Watch movies on laptop }

async { Talk to father

Talk to mother }

async { Buy fruits online using your smartphone

async { Make your bed }

Post on Facebook that you are done with all your tasks!
```

```
T_1
                                                     T_0
// To(Parent task)
STMT0;
                                                   STMTO
finish {
            //Begin finish
                                                   fork
  async {
    STMT1; //T, (Child task)
                                                   STMT2
                                     STMT1
            //Continue in To
  STMT2;
            //Wait for T<sub>1</sub>
                                                  → join
            //End finish
            //Continue in To
STMT3;
                                                   STMT3
```

Today's Lecture

- Parallel runtime system for task-scheduling
 - Work-stealing

Tasks-based parallel programming model and its underlying runtime system would be referred throughout in this course

Mapping the Linguistic Interface to the Parallel Runtime

- Compiler based runtimes
 - User code translated to runtime code and then compiled using a native compiler (e.g., gcc)
 - Compiler maintenance is a costly affair and it is not so easy to use new features from mainstream languages
 - Using standard debugger (e.g., gdb) is not possible as the line number information inside the symbol table is w.r.t. the compiler generated code and not w.r.t. the user written code
 - However, compiler based approach provide several opportunities for code optimizations and doing smart things
- Library based runtimesOur focus
 - Removes all the drawbacks of a compiler based approach

Tasks Based Parallel Programming Model

```
uint64_t fib(uint64_t n) {
  if (n < 2) {
    return n;
  } else {
    uint64_t x, y;
    finish([&]() {
        async([&](){ x = fib(n-1);});
        y = fib(n-2);
    });
    return (x + y);
  }
}</pre>
```

- High productivity due to serial elision
 - Removing all async and finish constructs results in a valid sequential program
 - Several existing frameworks support this programming model, although the name of the APIs for tasking would be different
- Uses an underlying high performance parallel runtime system for load balancing of dynamically created asynchronous tasks

	Java Fork/Join	Cilk	OpenMP	HClib[1]	TBB	C++11
Serial Elision	NO	Yes spawn-sync	Yes #pragma omp task #pragma omp taskwait	Yes async-finish	NO	Yes async-future
Performance	Limited	High	Limited	High	High	NO

Popular options for simple tasks based parallel programming model

[1] http://habanero-rice.github.io/hclib/



Mapping the Linguistic Interface to Library

Based Parallel Runtime

Runtime APIs

```
#include <runtime-API.h>
main() {
   init_runtime();
   finish {
     async (S1);
     S2;
   }
   finalize_runtime();
}
```

Initialize runtime and associated data-structures

Release runtime resources

Runtime equivalent of starting a finish scope

Runtime equivalent of closing a finish scope

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
  async (S1);
  S2;
  end_finish();
  finalize_runtime();
```

Runtime APIs

Initialize runtime and associated data-structures

Release runtime resources

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
   async (S1);
   S2;
  end_finish();
  finalize_runtime();
}
```

```
volatile boolean shutdown = false;
void init_runtime() {
   int size = runtime_pool_size();
   for(int i=1; i<size; i++) {
     pthread_create(worker_routine);
   }
}

void worker_routine() {
   while(!shutdown) {
     find_and_execute_task();
   }
}</pre>
```

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
   async (S1);
  S2;
  end_finish();
  finalize_runtime();
}
```

```
volatile int finish_counter = 0;
void start_finish() {
  finish_counter = 0; //reset
}
```

Note: in case of nested finish (e.g., Fibonacci), we need a better way to manage finish scopes. Recall, in Fibonacci every fib(n) call created a new finish, which ultimately creates a tree of finishes

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
   async (S1);
   S2;
  end_finish();
  finalize_runtime();
}
```

Note: Runtime stores pointer to the tasks passed in the async. To ensure valid pointer during task execution, we heap allocate the task and store pointer to the task on heap.

```
void async(task) {
  lock_finish();
  finish_counter++;//concurrent access
  unlock_finish();
  // copy task on heap
  void* p = malloc(task_size);
  memcpy(p, task, task_size);
  //thread-safe push_task_to_runtime
  push_task_to_runtime(&p);
  return;
}
```

Note: there are better ways to increment finish counter rather than doing it inside locks

Mapping the Linguistic Interface to Library

Based Parallel Runtime

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
  async (S1);
  S2;
  end_finish();
  finalize_runtime();
}
```

Note: there are better ways to decrement finish counter rather than doing it inside locks

```
while(finish counter != 0) {
           find and execute task();
void find and execute task() {
   //pop from runtime is thread-safe
   task = pop_task_from_runtime();
   if(task != NULL) {
      execute task(task);
      free(task);
      lock finish();
      finish_counter--;
      unlock finish();
```

void end_finish() {

```
#include <runtime-API.h>
main() {
  init_runtime();
  start_finish();
   async (S1);
  S2;
  end_finish();
  finalize_runtime();
}
```

```
void finalize_runtime() {
    //all spinning workers
    //will exit worker_routine
    shutdown = true;
    int size = runtime_pool_size();
    // master waits for helpers to join
    for(int i=1; i<size; i++) {
        pthread_join(thread[i]);
    }
}</pre>
```

How to Store Tasks in Runtime?

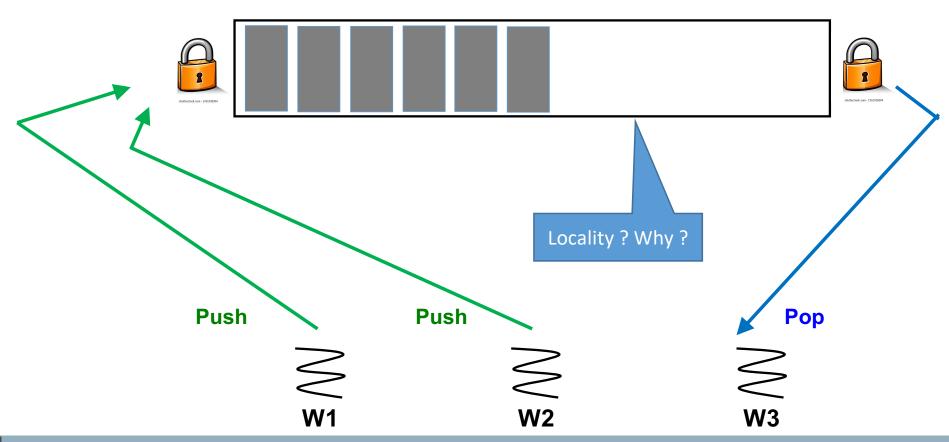
- push_task_to_runtime()
- pop_task_from_runtime()

Data-structures for storing tasks in a thread pool based runtime plays a very important role in determining the scalability and performance of the runtime

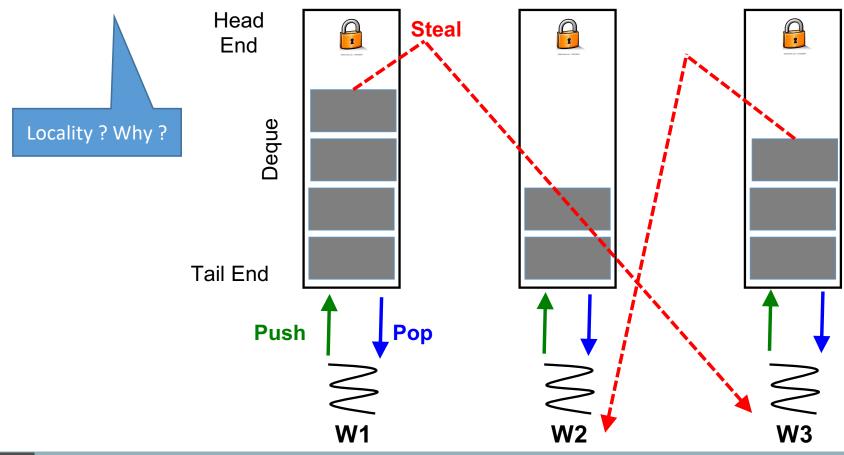
Parallel Runtime for Task Scheduling

- There are several different implementations of parallel runtimes, but at the core almost all of them uses either a work-sharing a work-stealing runtime underneath
- Tasks based parallel runtime systems primarily use workstealing runtime only

Work-Sharing Runtime System



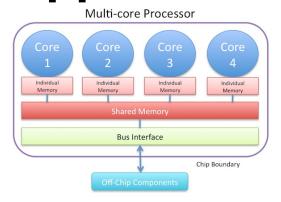
Work-Stealing Runtime System

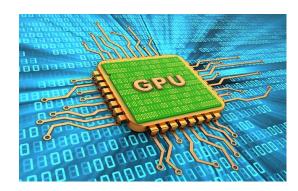


Work-Sharing v/s Work-Stealing

- Work-sharing
 - Busy worker re-distributes the task eagerly
 - Easy implementation through global task pool
 - Access to the global pool needs to be synchronized: scalability bottleneck
- Work-stealing
 - Busy worker pays little overhead to enable stealing
 - A lock is required for pop and steal only in case single task remaining on deque (only feasible by using atomic operations)
 - Idle worker steals the tasks from busy workers
 - Distributed task pools
 - Better scalability

Supported on Wide Range of Architectures



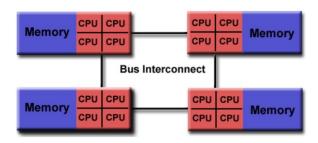


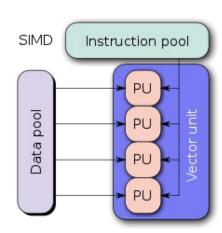


Multiprocessor System-on-Chip



Shared Memory (NUMA)

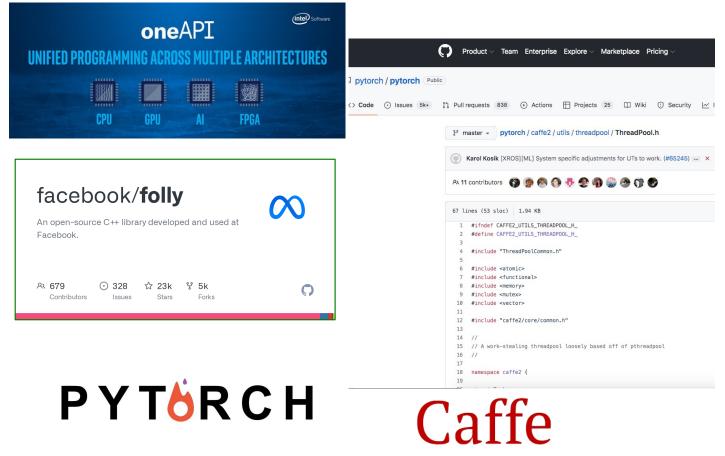






Supercomputers

Supported/Used by Several Companies/Projects



Futures

A Non-actor re-implementation of Scala Futures.

import com.twitter.conversions.DurationOps. import com.twitter.util.{Await, Future, Promise} val f = new Promise[Int] val g = f.map { result => result + 1 } f.setValue(1) Await.result(g, 1.second) // => this blocks for the futures result (and eventually returns 2) // Another option: g.onSuccess { result => println(result) // => prints "2" // Using for expressions: val xFuture = Future(1) val yFuture = Future(2) **Twitter** x <- xFuture y <- yFuture println(x + y) // => prints "3"

Future interrupts

Method raise on Future (def raise(cause: Throwable)) raises the interrupt described by cause to the producer of this Future. Interrupt handlers are installed on a Promise using setInterruptHandler, which takes a partial function:

Interrupts differ in semantics from cancellation in important ways: there can only be one interrupt handler per promise, and interrupts are only delivered if the promise is not yet complete.

Object Pool

The pool order is FIFO.

Reading Materials

- https://doi.org/10.1007/s11227-018-2238-4
- https://gee.cs.oswego.edu/dl/papers/fj.pdf

Next Lecture (#06)

- Performance in parallel programming (contd.)
- Project deadline-1 will be announced tonight with a deadline of ten days