

Lecture 06: Sequential Overheads from Task Granularity

Vivek Kumar

Computer Science and Engineering

IIIT Delhi

vivekk@iiitd.ac.in

```

void async(task) {
    lock_finish();
    finish_counter++; // concurrent access
    unlock_finish();
    // copy task on heap
    void* p = malloc(task_size);
    memcpy(p, task, task_size);
    // thread-safe push_task_to_runtime
    push_task_to_runtime(&p);
    return;
}

```


```

void find_and_execute_task() {
    // pop_from_runtime is thread-safe
    task = pop_task_from_runtime();
    if (task != NULL) {
        execute_task(task);
        free(task);
        lock_finish();
        finish_counter--;
        unlock_finish();
    }
}

```

overheads from task granularity


Last Lecture (Recap)



Today's Class

- ➔● Performance in task based parallel programming models (contd.)
 - Work-sharing
- Sequential overheads of work-stealing
 - Controlling task granularity


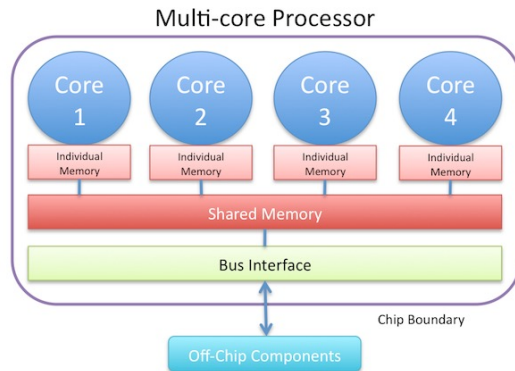
Work-Sharing Runtime System



Work-Sharing v/s Work-Stealing

- Work-sharing
 - Busy worker re-distributes the task eagerly
 - Easy implementation through global task pool
 - Access to the global pool needs to be synchronized: **scalability bottleneck**
- Work-stealing
 - Busy worker pays little overhead to enable stealing
 - A lock is required for pop and steal only in case single task remaining on deque (only feasible by using atomic operations)
 - Idle worker steals the tasks from busy workers
 - Distributed task pools
 - **Better scalability**



Supported on Wide Range of Architectures



Multiprocessor System-on-Chip

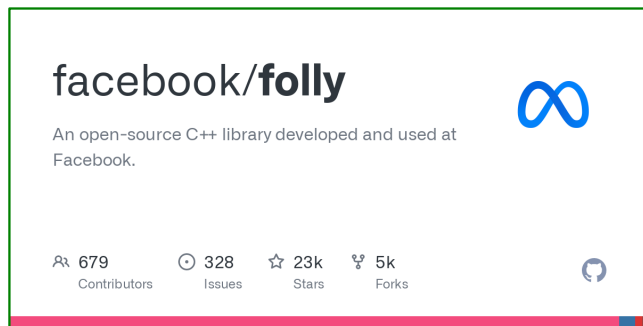


Shared Memory (NUMA)

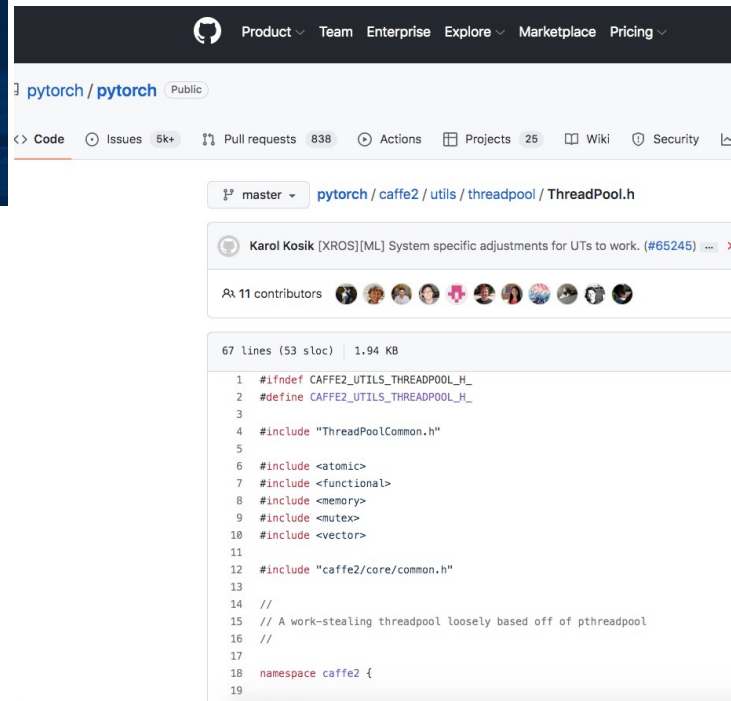


Supercomputers

Supported/Used by Several Companies/Projects



PYTORCH



Caffe

Futures

A Non-actor re-implementation of Scala Futures.

```
import com.twitter.conversions.DurationOps._
import com.twitter.util.{Await, Future, Promise}

val f = new Promise[Int]
val g = f.map { result => result + 1 }
f.setValue(1)
Await.result(g, 1.second) // => this blocks for the futures result (and eventually returns 2)

// Another option:
g.onSuccess { result =>
  println(result) // => prints "2"
}

// Using for expressions:
val xFuture = Future(1)
val yFuture = Future(2)

for {
  x <- xFuture
  y <- yFuture
} {
  println(x + y) // => prints "3"
}
```

Twitter

Future interrupts

Method `raise` on `Future` (`def raise(cause: Throwable)`) raises the interrupt described by `cause` to the producer of this `Future`. Interrupt handlers are installed on a `Promise` using `setInterruptHandler`, which takes a partial function:

```
val p = new Promise[T]
p.setInterruptHandler {
  case exc: MyException =>
    // deal with interrupt..
}
```

Interrupts differ in semantics from cancellation in important ways: there can only be one interrupt handler per promise, and interrupts are only delivered if the promise is not yet complete.

Object Pool


The pool order is FIFO.

Today's Class

- Performance in task based parallel programming models
 - Work-sharing
- ➔ ● Sequential overheads of work-stealing
 - Controlling task granularity

Sequential Overheads (1/2)

```
uint64_t fib(uint64_t n) {
    if (n < THRESHOLD) {
        return fib_sequential(n);
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            async([&y]() { y = fib(n-2); })
        })
        return (x + y);
    }
}
```



Running parallel recursive parallel Fib(40) using HClb as its async won't launch thread unlike std::async

- Sequential overhead = $\text{Ratio Time}_{\text{seq}} / \text{Time}_{\text{Par}}$
 - Time_{seq} is time for Fibonacci with serial elision
 - Time_{seq} is for the corresponding parallel version, but by only using a single thread (sequential execution)
- **Observation**
 - Overheads can be controlled using **optimal task granularity**
 - Neither too many tasks, nor too few!
- Options to control task granularity?
 1. Calculate Task-2 (fib of n-2) sequentially
 2. Don't create async tasks when N is less than certain threshold
 - What threshold is optimal?
 - What depth in a recursion tree is optimal to execute sequentially?
 3. Use memoization
 - Saving and reusing previously computed values of a function rather than recomputing them

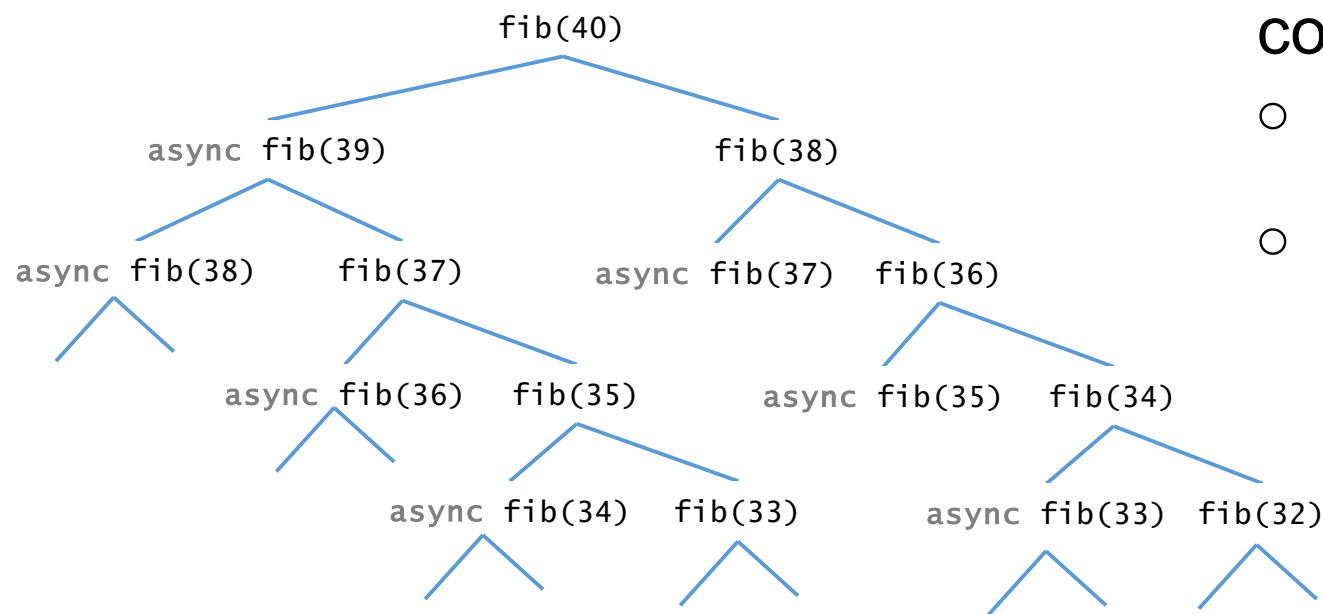
Sequential Overheads (2/2)

- Creating an async is not same as executing it sequentially
 - Each async has some metadata associated with it
 - Recall, coping user lambda on heap is important so that it can be used later even if the function that created that task has gone out of scope
 - **It is important to control task granularity**
 - We will discuss three different solutions in this lecture

Today's Class


- Performance in task based parallel programming models
 - Work-sharing
- Sequential overheads of work-stealing
 - ➡ ○ Controlling task granularity

First and Foremost, Work-Stealing is Best Suited for What Kind of Applications?




- Recursive divide-and-conquer style
 - Leads to fine granular task creation
 - **How its helpful?**
 1. Nested task creation

First and Foremost, Work-Stealing is Best Suited for What Kind of Applications?




- Recursive divide-and-conquer style
 - Leads to fine granular task creation
 - **How its helpful?**
 1. Nested task creation
 2. Stealing an async will eventually give birth to several new asyncs at the thief
 - It will keep the thief busy and reduce steal attempts

First and Foremost, Work-Stealing is Best Suited for What Kind of Applications?




- Recursive divide-and-conquer style
 - Leads to fine granular task creation
 - **Disadvantages?**
 1. Tasks created near the bottom of the tree are too small in computation, and wouldn't be able to keep a thief busy once stolen

First and Foremost, Work-Stealing is Best Suited for What Kind of Applications?



Stack Growth

Stack Bottom



Stack Top

- Recursive divide-and-conquer style
 - Leads to fine granular task creation
 - **Disadvantages?**
 1. Tasks created near the bottom of the tree are too small in computation, and wouldn't be able to keep a thief busy once stolen
 2. Thread stack too deep
 - Too many context switches for moving back and forth between caller and callee stack frames (although in user space)


How to Avoid Those Disadvantages

1. **Tasks near the bottom of the tree are small computations**
 - Automatic granularity control
 - Stop creating new async after some “**depth**” is reached
 - Async created after that “depth” is executed sequentially
2. **Deep thread stack due to recursion**
 - Using two versions of the parallel code
 - Convert recursion into iterative call after appropriate “depth”

Solution-1: Automatic Granularity Control


- Runtime can perform dynamic task aggregations

Aggregation of this task will not create any new async in this subtree, and the async will be executed sequentially

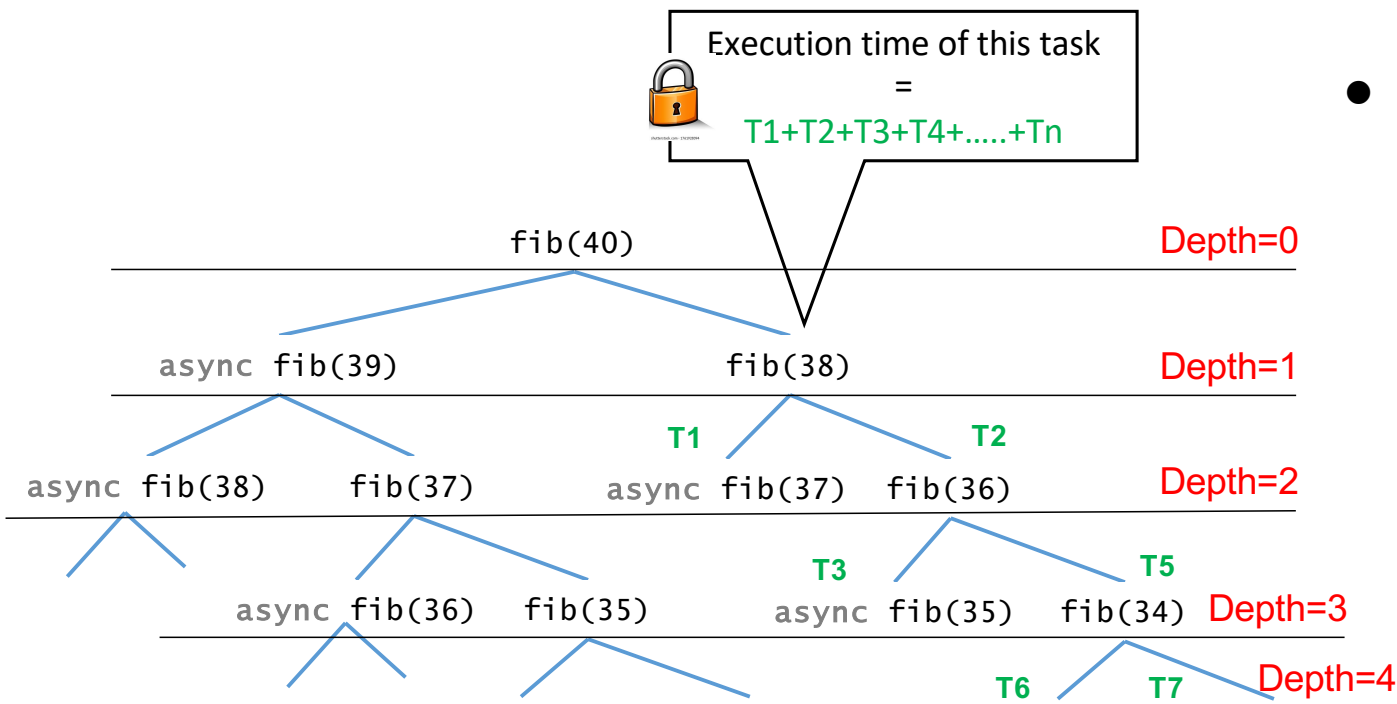


Solution-1: Automatic Granularity Control

- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task



Solution-1: Automatic Granularity Control




- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required

Solution-1: Automatic Granularity Control



Depth=0 Time=0	Depth=1 Time=0	Depth=2 Time=0	Depth=3 Time=0	
-------------------	-------------------	-------------------	-------------------	--




- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required
 - Update the execution time at its depth in a shared global hash map (key=depth, value=time)

Solution-1: Automatic Granularity Control



Depth=0 Time=0	Depth=1 Time=0	Depth=2 Time=10	Depth=3 Time=0	
-------------------	-------------------	--------------------	-------------------	--


- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required
 - Update the execution time at its depth in a shared global hash map (key=depth, value=time)



Solution-1: Automatic Granularity Control



Depth=0 Time=0	Depth=1 Time=0	Depth=2 Time=10	Depth=3 Time=0	
-------------------	-------------------	--------------------	-------------------	--




- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required
 - Update the execution time at its depth in a shared global hash map (key=depth, value=time)
 - Averaging of value (time) for a given key (depth) when more than one tasks complete its execution
 - Averaging would be stopped after enough samples collected at a depth

Solution-1: Automatic Granularity Control



Depth=0 Time=0	Depth=1 Time=0	Depth=2 Time=11	Depth=3 Time=0	
-------------------	-------------------	--------------------	-------------------	--

Average value of all time




- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required
 - Update the execution time at its depth in a shared global hash map (key=depth, value=time)
 - Averaging of value (time) for a given key (depth) when more than one tasks complete its execution
 - Averaging would be stopped after enough samples collected at a depth

Solution-1: Automatic Granularity Control



Key=0 Value=14	Key=1 Value=12	Key=2 Value=11	Key=3 Value=9	
-------------------	-------------------	-------------------	------------------	--



- Runtime can perform dynamic task aggregations
- Each task keeps track of its depth in the recursion tree, and its execution time
 - Depth is stored locally inside the task
- Whenever a task complete its execution, it does two things
 - It add its execution time to the parent task's execution time
 - Mutual exclusion required
 - Update the execution time at its depth in a shared global hash map (key=depth, value=time)
 - Averaging of value (time) for a given key (depth) when more than one tasks complete its execution
 - Averaging would be stopped after enough samples collected at a depth
- Depth threshold decided based on the execution time of tasks at each depth
 - Beyond this depth threshold tasks would be aggregated

Solution-2: Using Two Versions of the Code

```
uint64_t fib(uint64_t n) {  
    if (n < 2) {  
        return n;  
    } else {  
        uint64_t x = fib(n-1);  
        uint64_t y = fib(n-2);  
        return (x + y);  
    }  
}
```



```
uint64_t fib(uint64_t n) {  
    uint64_t f1=1;  
    uint64_t f2=1;  
    uint64_t m=2;  
    while(m < n) {  
        uint64_t temp = f2+f1;  
        f1=f2;  
        f2=temp;  
        m=m+1;  
    }  
    return f2;  
}
```

- When depth threshold is reached, switch to an iterative version of the recursive algorithm
 - Most of the recursive algorithms can be converted into iterative algorithm
 - Although, asking the user to provide an iterative version is breaking the support for serial elision


There is a general format for converting tail recursion into iterative version: <https://www.baeldung.com/cs/convert-recursion-to-iteration>

Solution-3: Using Memoization

- **Memoization** is about saving and reusing previously computed values of a function rather than recomputing them
- An optimization technique with space-time tradeoff

Memoization: Functional Programming

- Functional programming emphasizes functions whose results that depend only on their inputs and not on any other program state
- Revisiting recursive Fibonacci
 - Calling a function, `fib(x)`, twice with the same value for the argument 'x' will produce the same result both times



```

uint64_t fib(uint64_t n) {
    if (n < THRESHOLD) {
        return fib_sequential(n);
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            async([&y]() { y = fib(n-2); })
        })
        return (x + y);
    }
}
  
```

Solution-3: Applying Memoization on Fib (1/2)

```
uint64_t fib(uint64_t n) {
    int value = getValue(n);
    if(value != -1) return value;
    else if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish ([&]( ) {
            async ([&]( ) {x = fib(n-1);});
            y = fib(n-2);
        });
        int result = x + y;
        storeValue(n, result);
        return result;
    }
}
```

```
My_hashmap<int, int> my_cache;
int getValue(int key) {
    /* return value if available */
    /* else return -1 */
}
void storeValue(int key, value) { ..... }
```

- A function can only be memoized if it is functional
- Related to caching
 - memoized function "remembers" the results corresponding to some set of specific inputs
 - memoized function populates its cache of results transparently on the fly, as needed, rather than in advance

Solution-3: Applying Memoization on Fib (2/2)

```
uint64_t fib(uint64_t n) {
    int value = getValue(n);
    if(value != -1) return value;
    else if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish ([&]( ) {
            async ([&]( ) {x = fib(n-1);});
            y = fib(n-2);
        });
        int result = x + y;
        lock(); {storeValue(n, result);} unlock();
        return result;
    }
}
```

```
My_hashmap<int, int> my_cache;
int getValue(int key) {
    /* return value if available */
    /* else return -1 */
}
void storeValue(int key, value) { ..... }
```

- A function can only be memoized if it is functional
- Related to caching
 - memoized function "remembers" the results corresponding to some set of specific inputs
 - memoized function populates its cache of results transparently on the fly, as needed, rather than in advance

Reading Materials

- Automatic granularity control
 - An adaptive cut-off for task parallelism, SC 2008
 - https://www.academia.edu/download/35796885/1234120839604_a36-duran.pdf
- Using multiple versions of the code
 - A static cut-off for task parallel programs, PACT 2016
 - https://www.eidos.ic.i.u-tokyo.ac.jp/~iwasaki/files/PACT2016_slides.pdf
- You may only read the implementation section and skip theorem/proofs (if any)

Next Lecture (L #07)

- Sequential overheads from concurrent deque