

Lecture 07: Sequential Overheads from Concurrent Deque

Vivek Kumar

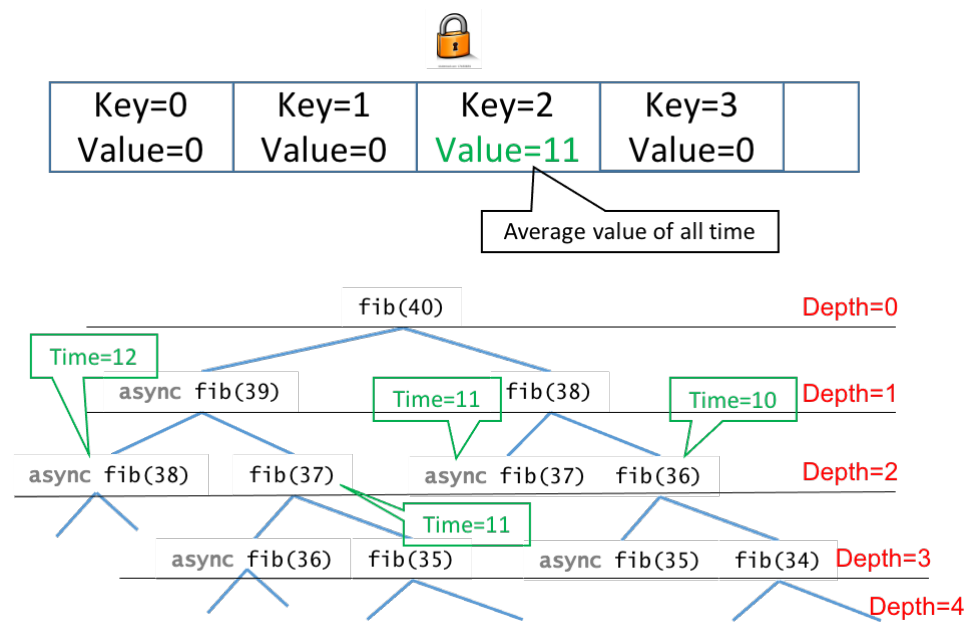
Computer Science and Engineering

IIIT Delhi

vivekk@iiitd.ac.in

Last Lecture (Recap)

- Sequential overheads from fine granular task creation
 - Tasks near the bottom of tree are smaller computations
 - Deep procedure calling stack in thread due to recursion
- Automatically controlling task granularity in recursive task decomposition
 - Assumption is that the tree (computation graph) is well balanced
 - Dynamic task aggregation
 - Each task records its depth and the execution time at that depth
 - Above information is used to decide if any more tasks at certain depth has to be created or should be executed serially

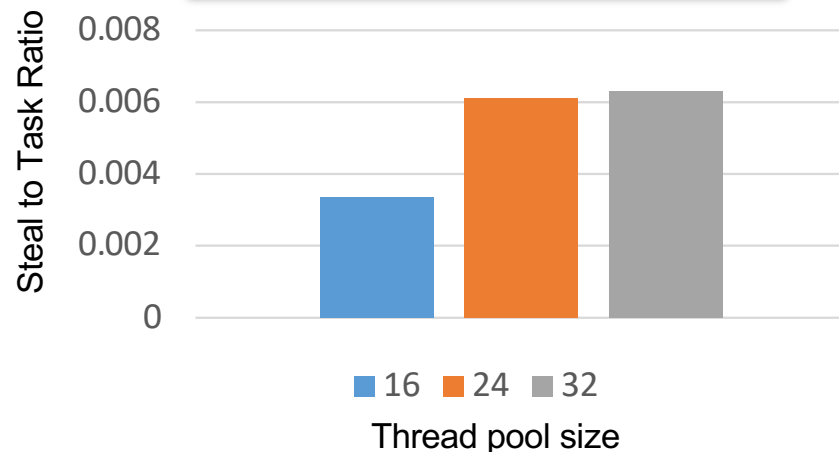


Today's Class

- ➔ ● Minimizing overheads from deque operations
 - Using a mix of list and deque
 - Using private deques

Sequential Overheads (1/2)

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```



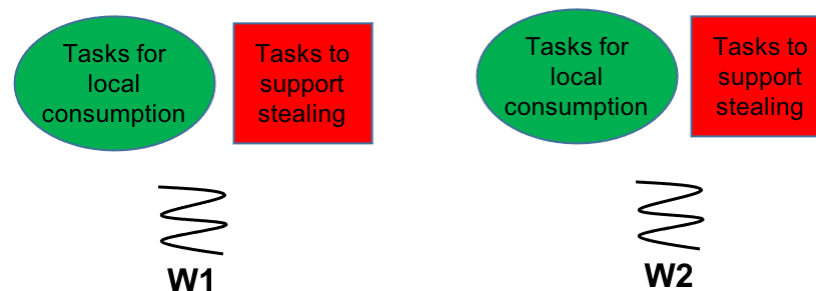
- Graph shows the ratio of total tasks stolen to total tasks created while executing Fibonacci(30) on a 32-core processor at different thread counts (16, 24, and 32)
 - Using HCLib implementation of Fib that spawns task for every fib(n-1) recursive call until n<2
- We can observe the steal ratio is extremely low
 - Implies that most of the tasks created by a victim is consumed by itself

Sequential Overheads (2/2)

- Creating an async is not same as executing it sequentially
 - Each async has some metadata associated with it
 - Copying user lambda on heap so that it can be used later even if the function that created that task has gone out of scope
 - **It is important to control task granularity**
- Deque operations are costly*
 - For implementing any thread-safe (concurrent) data structure we always have to use some sort of mutual exclusion that avoids the race condition
 - Imagine using an integer counter that is private to a thread v/s using an integer counter that is to be updated concurrently by several threads
 - **It is important to use an efficient deque implementation**

Reducing Concurrent Access: General Idea

- Steals are rare
 - Majority of the tasks produced by the victim are consumed by itself
- Recall, deques are concurrent data structure, hence to reduce the overheads, each victim should minimize accessing its “concurrent” deque for push/pop
 - Then where to store async tasks at victims?
 - Use a mix of private and shared task pools
 - Push/pop from private pool, but ensure task(s) availability in shared pool to support stealing




Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```

fib(40)




Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list
- Victim checks total tasks on its deque during each push and pop operations

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list
- Victim checks total tasks on its deque during each push and pop operations

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list
- Victim checks total tasks on its deque during each push and pop operations

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```





Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list
- Victim checks total tasks on its deque during each push and pop operations
- Thief always steal from the deque as it was doing in default case

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```

pop(task)



Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>


Reducing Concurrent Access: Using List & Deque

- Each worker uses a private linked list and a concurrent deque
- Victim ensures there are some minimum number of tasks always available in concurrent deque to support steals
- If there are sufficient tasks available in concurrent deque then victim always push/pop from its private list
- Victim checks total tasks on its deque during each push and pop operations
- Thief always steal from the deque as it was doing in default case

```
uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x, y;
        finish([&]() {
            async([&x]() { x = fib(n-1); })
            y = fib(n-2);
        })
        return (x + y);
    }
}
```

fib(40)

pop(task)



Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using List & Deque

```
Task* pop() {
    Task* t = NULL;
    if(current_worker->Tail != NULL) {
        t = current_worker->pop_from_list_tail();
        if(current_worker->deque_size < DEQUE_LIMIT) {
            move_task_from_list_to_deque();
        }
    } else {
        t = current_worker->deque_pop();
    }
    return t;
}
```

```
bool move_task_from_list_to_deque() {
    Task* t = pop_from_list_head();
    if(t) {
        current_worker->deque_push(t);
        return true;
    } else {
        return false;
    }
}
```

Popping items from Head for adding into deque has some benefits with recursive task creation? Why?

```
#define DEQUE_LIMIT /* Some value */

struct Node {
    User_Lambda task;
    Node* next;
}

Node *Head, *Tail; /* Thread local */

void push(T lambda) {
    bool success = true;
    /* Add task to my deque if required */
    if(current_worker->deque_size < DEQUE_LIMIT) {
        success = move_task_from_list_to_deque();
    }
    if(!success) current_worker->deque_push(lambda);
    else current_worker->push_to_list_tail(lambda);
}
```

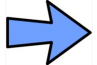
Paper based on a similar idea: <https://terpconnect.umd.edu/~barua/ppopp164.pdf>

Reducing Concurrent Access: Using **List** & **Deque**


● Issues

- Doesn't support stealing more than one tasks at a time
 - Stealing more than one task can reduce the steal frequency
- Maintaining a linked list means more mallocs/frees for adding/removing nodes
 - Tasks are anyway copied on heap

Today's Class

- Minimizing overheads from deque operations
 - Using a mix of list and deque
 -  ○ Using private deques

Reducing Concurrent Access: Using **Private Deque**




- Every worker maintains three data structures
 - A non-concurrent private deque
 - Same as the default deque, but without the support for concurrent (thread-safe) accesses
 - One mailbox
 - That can store one or more tasks
 - Contains a counter indicating total number of stored tasks
 - One shared counter

Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using **Private Deque**

Mailbox

-1

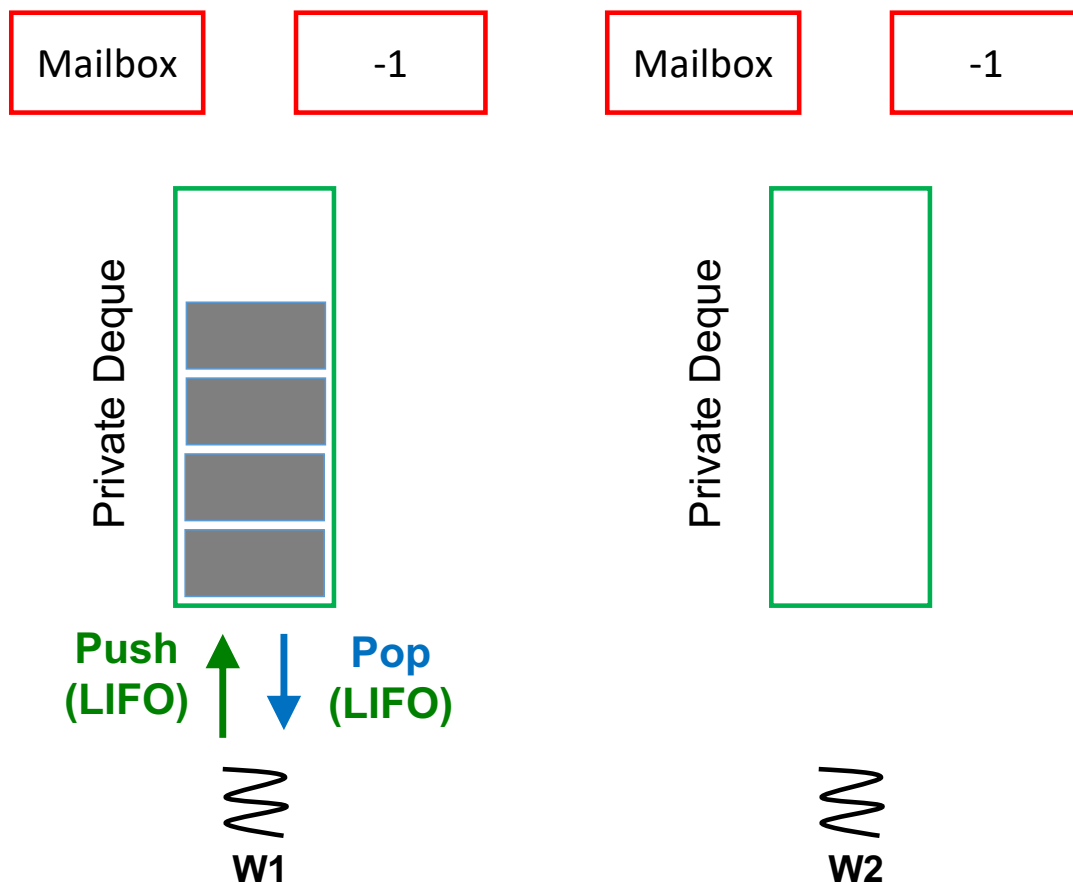


- Victim

- Push/pop tasks into its private deque

Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using **Private Deque**

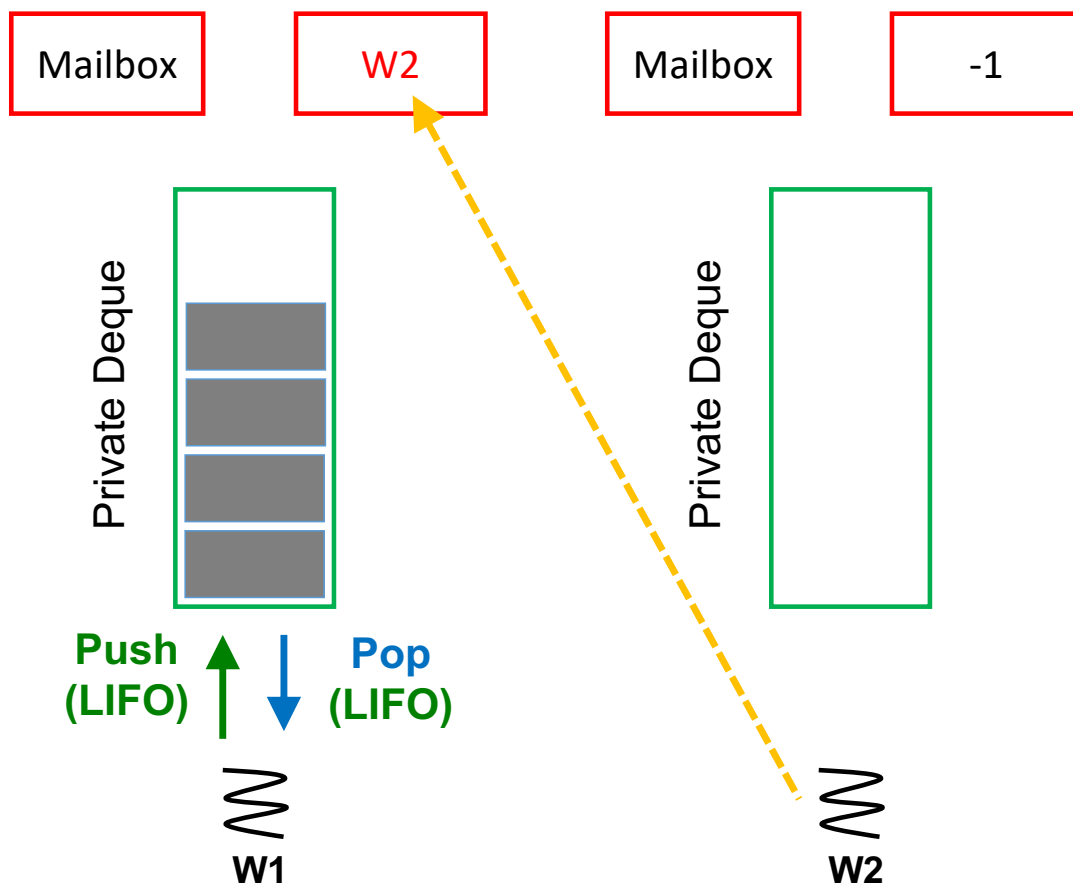


● Thief

- Selects a random victim (W1) who has items on its deque
- Checks victim deque size without any locks
 - The usual way of doing (head-tail)

Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using **Private Deque**

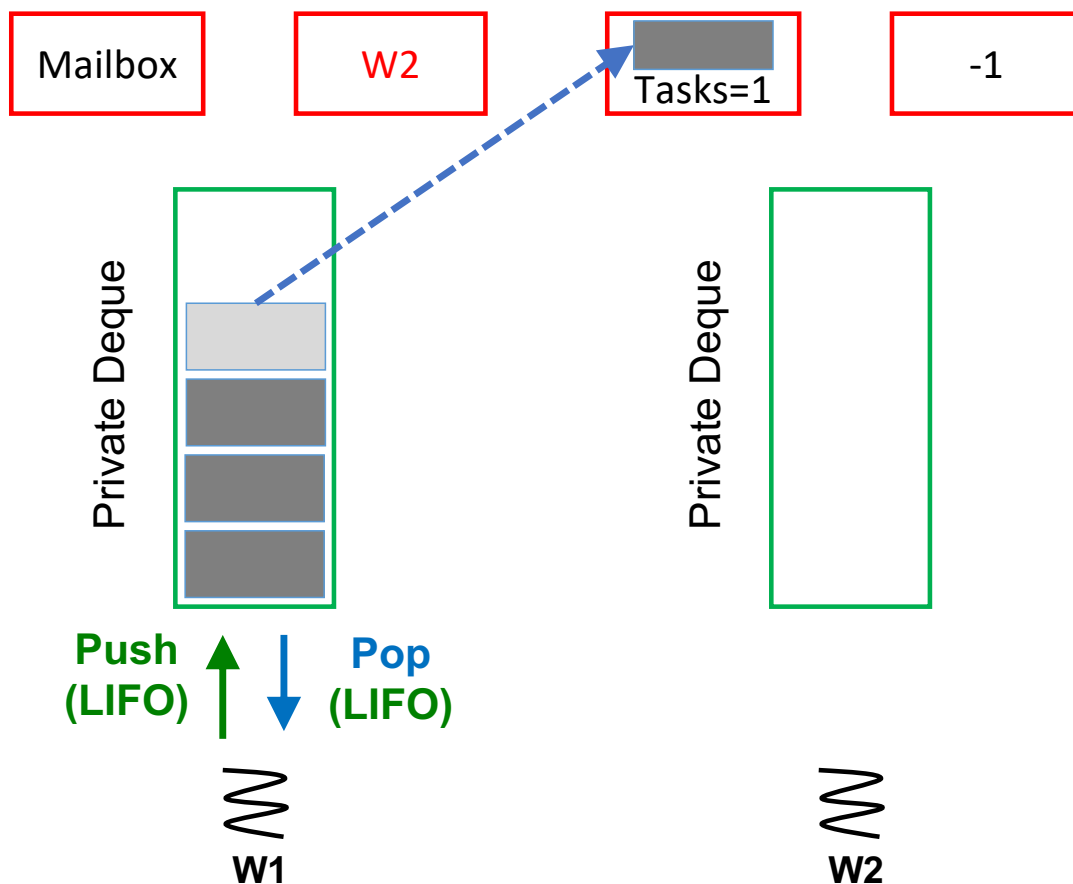


● Thief

- Record its own id inside the request box at W1 (critical section), and goes inside condition wait
- Only one thief at a time

Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using Private Deque



- Victim

- Check its request box inside each push/pop/steal
- If tasks are available on victim's private deque
 - Pop item(s) from the head and copies it into the waiting thief's mailbox (W2)
 - Update W2's mailbox with the total number of tasks copied

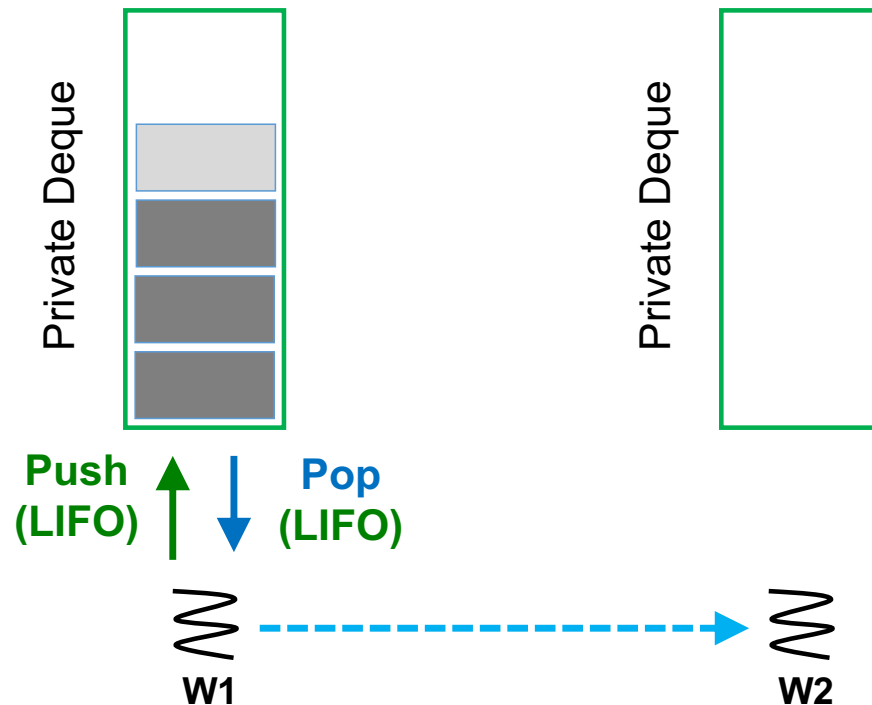
Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using **Private Deque**




● Victim

- Clears its request box
- Signals the waiting thief W2



Paper: <https://hal.inria.fr/hal-00863028/document>

Reducing Deque Access: Using **Private Deque**



● Thief

- Unblocks after being notified by W1
- Steal tasks from its mailbox and start executing them
 - If more than one task received then extra tasks pushed to its private deque
- Failed steal attempt if it did not receive any task (i.e., if W1 ran out of tasks)

Paper: <https://hal.inria.fr/hal-00863028/document>

Reading Materials

- Using list and deques together
 - <https://terpconnect.umd.edu/~barua/ppopp164.pdf>
- Private deques
 - <https://hal.inria.fr/hal-00863028/document>
- You may only read the implementation section and skip theorem/proofs (if any)

Next Lecture (L #08)

- Non uniform memory access architectures