

# Lecture 12: User Level Threads

Vivek Kumar

Computer Science and Engineering

IIIT Delhi

[vivekk@iiitd.ac.in](mailto:vivekk@iiitd.ac.in)

# Today's Class

- ➔ ● Boost fiber library
- Quiz-3

# boost::fibers::fiber

- A fiber is a userland thread unlike the kernel thread (e.g., pthread maps **1x1** with kernel thread in Linux)
  - Several fibers can map with single pthread (**M x N** threading)



- Fiber emulates much of the std::thread
  - Extends the concurrency taxonomy
    - On a single computer, multiple processes can run
    - Within a single process, multiple threads can run
    - Within a single thread, multiple fibers can run
- Builds on top of boost::context
  - Each fiber has its own stack, registers, instruction pointer..
    - It means they can be scheduled cooperatively
- It is super easy to create a fiber
 

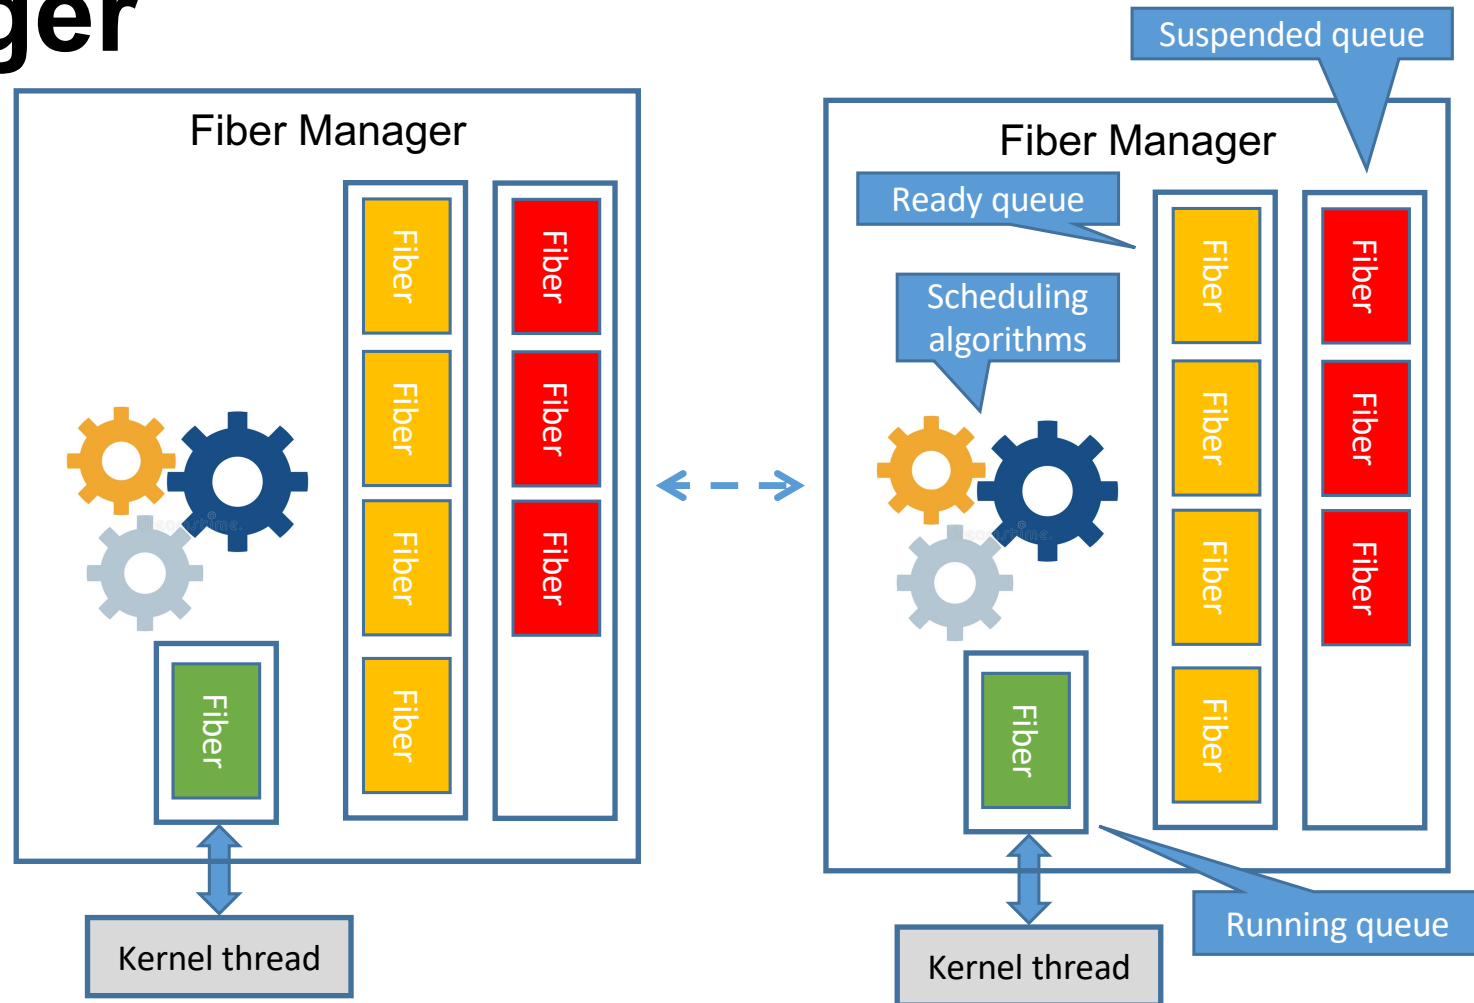
```
boost::fibers::fiber (F, [=]() { /*Do something*/ }); // Spawns a fiber F
```

# Fiber v/s Thread

- A thread can run only one fiber at a time
  - Although several fibers can be queued up for execution at a thread at any given time
- Creating several fibers by a single thread doesn't imply parallelism unlike creating several threads
  - By default fibers created by a thread will run by that thread only, but it can be detached to allow its execution at any other thread

# Fiber Manager

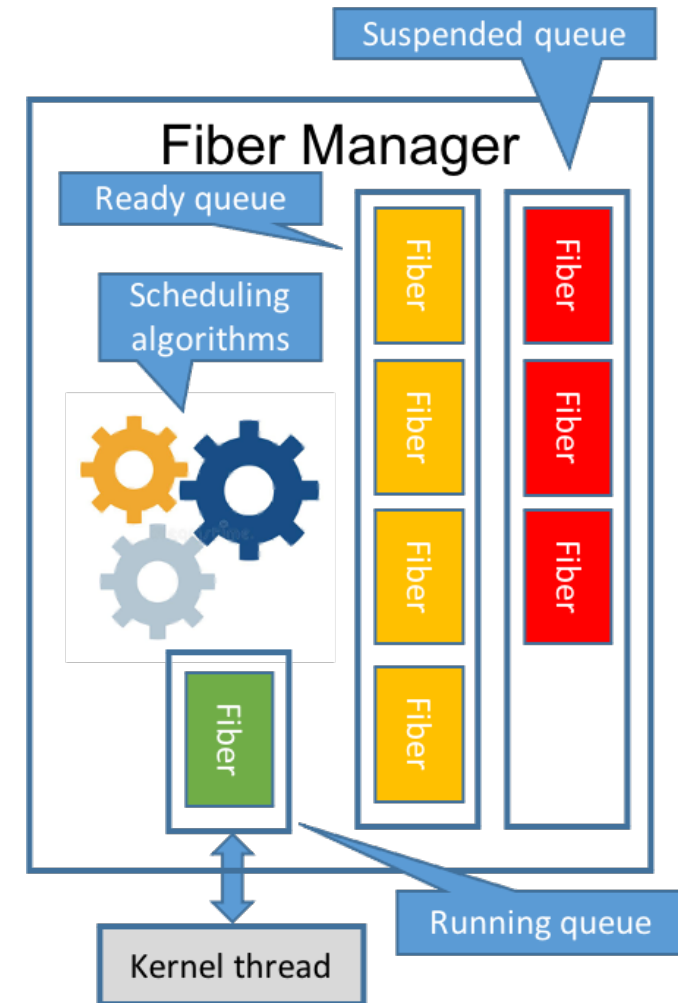
- The fibers in a thread are coordinated by a fiber manager
  - The manager created/managed silently by the fiber library



# Fiber Manager

- Similar to threads, a fiber can be in the running, suspended or ready state
- Fibers trade control with the manager in a cooperative way
  - `boost::this_fiber::yield();`
  - `boost::this_fiber::sleep_for`
  - `boost::this_fiber::sleep_until`
  - `boost::fibers::mutex`
  - `boost::fibers::condition_variable`
  - `some_fiber.join()`
  - ....
- Manager uses a scheduling algorithm to select a ready fiber to run (any similarity with Linux kernel?)
- Manager carries out the context switch to swap between the fibers
  - Kernel thread blocks if there are no ready fibers

These operations will land the fiber into which queue (ready/suspended)?



# Fiber Scheduler

- Manager uses a default round-robin scheduler
  - Scheduling within a thread
- Boost fibers provides `shared_work` and `work_stealing` as alternative schedulers to `round_robin`
  - Scheduling across the threads
- Boost fibers also allow creation of a custom scheduler

```
void thread_function() {  
    boost::fibers::use_scheduling_algorithm<my_own_fiber_scheduler>();  
}
```

# Fiber Context Switching is Extremely Fast

Table 1.3. time per thread (average over 10,000 - unable to spawn 1,000,000 threads)

<code>pthread</code>	<code>std::thread</code>	<code>std::async</code>
54 $\mu$ s - 73 $\mu$ s	52 $\mu$ s - 73 $\mu$ s	106 $\mu$ s - 122 $\mu$ s

Table 1.4. time per fiber (average over 1,000,000)

fiber (16C/32T, work stealing, tcmalloc)	fiber (1C/1T, round robin, tcmalloc)
0.05 $\mu$ s - 0.09 $\mu$ s	1.69 $\mu$ s - 1.79 $\mu$ s

Source: [https://www.boost.org/doc/libs/1\\_80\\_0/libs/fiber/doc/html/fiber/performance.html](https://www.boost.org/doc/libs/1_80_0/libs/fiber/doc/html/fiber/performance.html)



# Question

- Is there any difference(s) between calling `std::future.get()` / `std::future.wait()` on a `std::thread` v/s a `boost::fiber`?

# Creating Fibers

```
#define millisleep(x) std::this_thread::sleep_for(std::chrono::milliseconds(a))
.....
millisleep(500);
millisleep(100);
```

```
#include <boost/fiber/all.hpp>
#define millisleep(x) boost::this_fiber::sleep_for(std::chrono::milliseconds(a))
.....
boost::fibers::fiber f1 ([=]() { millisleep(500); }); // Fiber F1 launched
boost::fibers::fiber f2 ([=]() { millisleep(100); }); // Fiber F2 launched
f1.join(); // Wait for termination of F1
f2.join(); // Wait for termination of F2
```

Method call can be made directly instead of passing lambda, e.g. f1(foo, p1, p2, p3), where 'p' are parameter to method 'foo'

- What would be the execution time of these two programs?
  - Note that it's a single thread execution in each case
- Note that both programs are using different implementations of sleep
  - Fiber manager handle its own sleep, but not the std::sleep

# Fiber Futures

```
uint64_t fib(uint64_t n) {
    if(n<2) {
        return n;
    } else {
        std::future<uint64_t> f1 (std::async([]={}){ return fib(n-1); }));
        std::future<uint64_t> f2 (std::async([]={}){ return fib(n-2); }));
        //get will block until result is ready
        return f1.get() + f2.get();
    }
}
```

```
uint64_t fib(uint64_t n) {
    if(n<2) {
        return n;
    } else {
        boost::fibers::future<uint64_t> f1 (boost::fibers::async([]={}){ return fib(n-1); }));
        boost::fibers::future<uint64_t> f2 (boost::fibers::async([]={}){ return fib(n-2); }));
        //get will block until result is ready
        return f1.get() + f2.get();
    }
}
```

For n=40:

- Which of these programs would be faster?
- Which of these programs is a parallel program?

# Yielding Fibers

```
boost::fibers::fiber f1([=]() {  
    cout << "A ";  
    boost::this_fiber::yield();  
    cout << "B ";  
    boost::this_fiber::yield();  
    cout << "C ";  
});  
  
boost::fibers::fiber f2([=]() {  
    cout << "D ";  
    boost::this_fiber::yield();  
    cout << "E ";  
    boost::this_fiber::yield();  
    cout << "F ";  
});  
  
f1.join();  
f2.join();
```

- `yield()` saves the context of currently running fiber, and places it inside the **ready** queue
  - Manager can schedule it again based on the scheduling algorithm
- What will be the output of this program?

# Producer-Consumer using Fibers

```
boost::fibers::mutex mtx;
boost::fibers::condition_variable cnd;
std::string str;

boost::fibers::fiber f1([&]() {
    std::unique_lock<boost::fibers::mutex> lck(mtx);
    if(str.size() == 0) {
        cnd.wait(lck);
    }
    cout << str << endl;
});

boost::fibers::fiber f2([&]() {
    std::unique_lock<boost::fibers::mutex> lck(mtx);
    str = "Hello Fiber";
    cnd.notify_one();
});

f1.join();
f2.join();
```

- Fiber F1 moving into suspended queue, and then back into ready queue after a notify from F2
  - Single thread execution!

# Fiber Pitfalls

```
std::mutex mtx;
std::condition_variable cnd;
std::string str;


boost::fibers::fiber f1([=]() {
    std::unique_lock<std::mutex> lck(mtx);
    if(str.size() == 0) {
        cnd.wait(lck);
    }
    cout << str << endl;
});

boost::fibers::fiber f2([=]() {
    std::unique_lock<std::mutex> lck(mtx);
    str = "Hello Fiber";
    cnd.notify_one();
});

f1.join();
f2.join();
```

- Can you spot the difference?
  - What effect it would cause, and why?

# Fiber Overheads / Limitations



- Language restriction
  - Fiber library requires C++11
  - Cannot be used in C-based HPC libraries/programs
- (Serious) Runtime overheads
  - Graph shown for calculating recursive Fibonacci that spawns detached fiber for every recursive call until threshold reached ( $n < 10$ )
    - Total nested fibers created
      - Fib [30, 57K], Fib [35, 635K], Fib [40, 7049K]
  - Single worker used!
  - Platform details
    - AMD EPYC 7551 32-core processor
    - Ubuntu 18.04.3 LTS
    - GCC version 7.5.0
      - -O3 flag used
    - Boost version 1\_80\_0

# Reading Materials

- Fibers

- [https://www.boost.org/doc/libs/1\\_80\\_0/libs/fiber/doc/html/index.html](https://www.boost.org/doc/libs/1_80_0/libs/fiber/doc/html/index.html)



# Installing Boost Context and Fiber Library

- Install Boost
  - `wget https://boostorg.jfrog.io/artifactory/main/release/1.80.0/source/boost_1_80_0.tar.gz`
  - `tar xvfz boost_1_80_0.tar.gz`
  - `cd ~/boost_1_80_0/`
  - `./bootstrap.sh --prefix=/absolute/path/to/boost-install --with-libraries=fiber,context`
  - `./b2 install`
- Compile programs
  - `g++ -O3 -I/absolute/path/to/boost-install/include -L/absolute/path/to/boost-install/lib Program.cpp -lboost_fiber -lboost_context -lpthread`
- Execute programs
  - `export LD_LIBRARY_PATH=/absolute/path/to/boost-install/lib:$LD_LIBRARY_PATH`
  - `./a.out`

# Next Lecture (L #13)

- Mid semester review