

Lecture 13: Mid Semester Review

Vivek Kumar

Computer Science and Engineering

IIIT Delhi

vivekk@iiitd.ac.in

Introduction to Parallel Programming (1/2)

- Free lunch is now over!
 - Multicore processors everywhere
- Explicit multithreading
- Amdahl's law

```
#include <inttypes.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

uint64_t fib(uint64_t n) {
    if (n < 2) {
        return n;
    } else {
        uint64_t x = fib(n-1);
        uint64_t y = fib(n-2);
        return (x + y);
    }
}

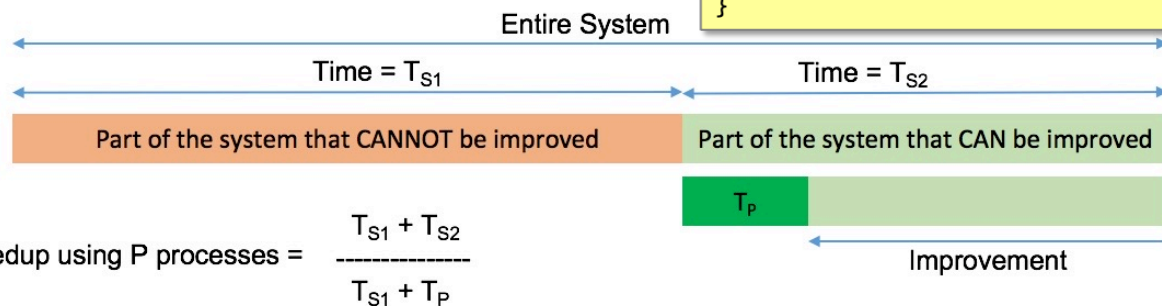
typedef struct {
    uint64_t input;
    uint64_t output;
} thread_args;

void *thread_func(void *ptr) {
    uint64_t i =
        ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t thread;
    thread_args args;
    int status;
    uint64_t result;

    if (argc < 2) { return 1; }
    uint64_t n = strtoul(argv[1], NULL, 0);
    if (n < 30) {
        result = fib(n);
    } else {
        args.input = n-1;
        status = pthread_create(&thread,
                                NULL,
                                thread_func,
                                (void*) &args);

        // main can continue executing
        if (status != NULL) { return 1; }
        result = fib(n-2);
        // wait for the thread to terminate.
        status = pthread_join(thread, NULL);
        if (status != NULL) { return 1; }
        result += args.output;
    }
    printf("Fibonacci of %" PRIu64 " is %" PRIu64 ".\n",
           n, result);
    return 0;
}
```



Introduction to Parallel Programming (2/2)




```
1. pthread_mutex_lock(&mutex);
2. while(task_queue_size() == 0)
3.   pthread_cond_wait(&cond, &mutex);
4. }
5. task = pop_task_queue();
6. pthread_mutex_unlock(&mutex);
7. execute_task(task);
```

```
1. pthread_mutex_lock(&mutex);
2. int queue_size = task_queue_size();
3. push_task_queue(&task);
4. if(queue_size == 0) {
5.   pthread_cond_broadcast(&cond);
6. }
7. pthread_mutex_unlock(&mutex);
```

Concurrency Decomposition

- How should one decompose a task into various subtasks?

- No single universal recipe
 - Recursive decomposition
 - Data decomposition
 - Exploratory decomposition
 - Speculative decomposition



- Serial execution time = $7T$
- Parallel execution time using 4 threads to compute each triangle in parallel = T
- Speedup (4 threads) = $7T/T = 7$
- **Super-linear** speedup

- Serial execution time = $3T$
- Parallel execution time using 4 threads to compute each triangle in parallel = $3T$
- Speedup (4 threads) = $3T/3T = 1$
- **Sub-linear** speedup

```
int val = T1 //compute intensive
switch(val) {
  case 0: T2; break;
  case 1: T3; break;
  ....
  case n: Tn; break;
}
```

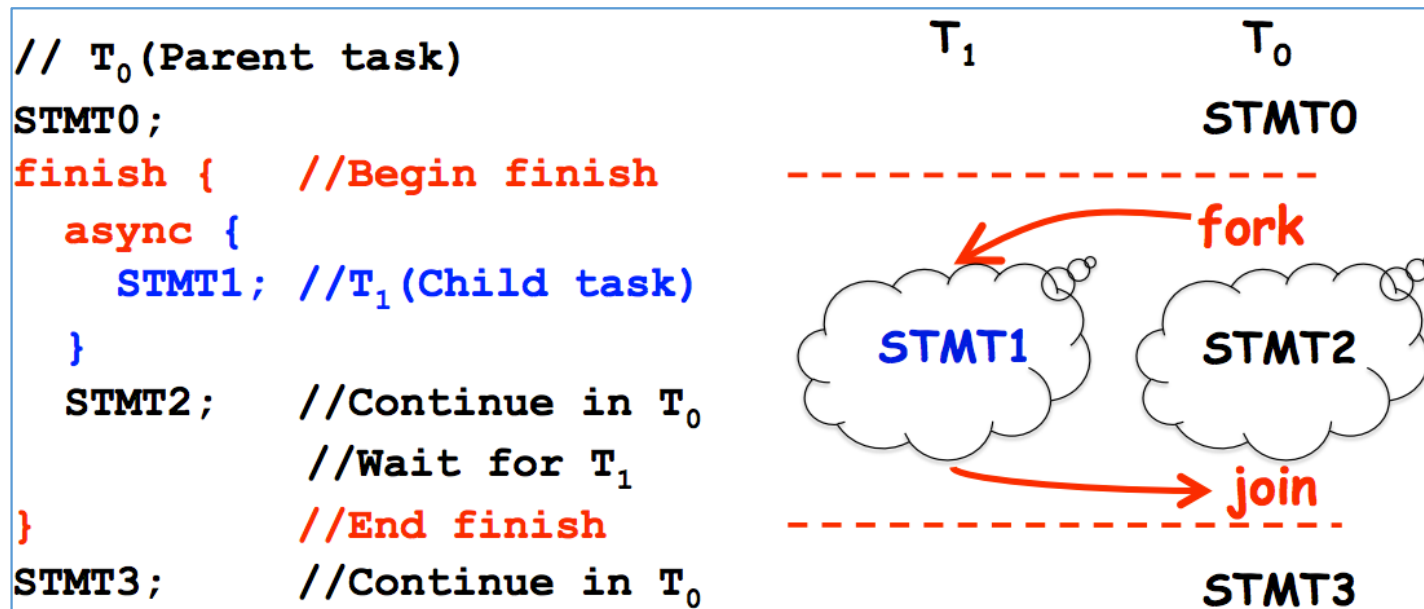
Async and Finish Statements for Task Creation and Termination (Pseudocode)

async S

- Creates a new child task that executes statement **S**

finish S


- Execute **S** but wait until all async in **S**'s scope have terminated



Computation Graph

- Granularity of task decomposition
 - Fine and coarse granular
- Computation graph


finish {		FS
async { Wash your clothes in washing machine }		T1
Complete your PRMP project deadline		T2
async { Watch movies on laptop }		T3
async { Talk to father		T4
Talk to mother }		T5
async { Buy fruits online using your smartphone }		T6
Make your bed		T7
}		FE
Post on Facebook that you are done with all your tasks!		T8




Critical Path

- Critical path is the longest weighted path in computation graph that represents task serialization

finish {		FS
async { Wash your clothes in washing machine }		T1
Complete your PRMP project deadline		T2
async { Watch movies on laptop }		T3
async { Talk to father		T4
Talk to mother }		T5
async { Buy fruits online using your smartphone }		T6
Make your bed		T7
}		FE
Post on Facebook that you are done with all your tasks!		T8




Work-Stealing Runtime System


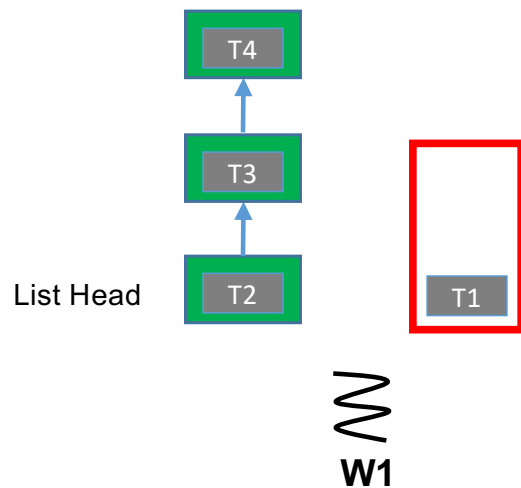


Sequential Overheads: Task Granularity

- Sequential overheads from fine granular task creation
 - Tasks near the bottom of tree are smaller computations
 - Deep procedure calling stack in thread due to recursion
- Automatically controlling task granularity in recursive task decomposition
 - Assumption is that the tree (computation graph) is well balanced
 - Dynamic task aggregation
 - Each task records its depth and the execution time at that depth
 - Above information is used to decide if any more tasks at certain depth has to be created or should be executed serially



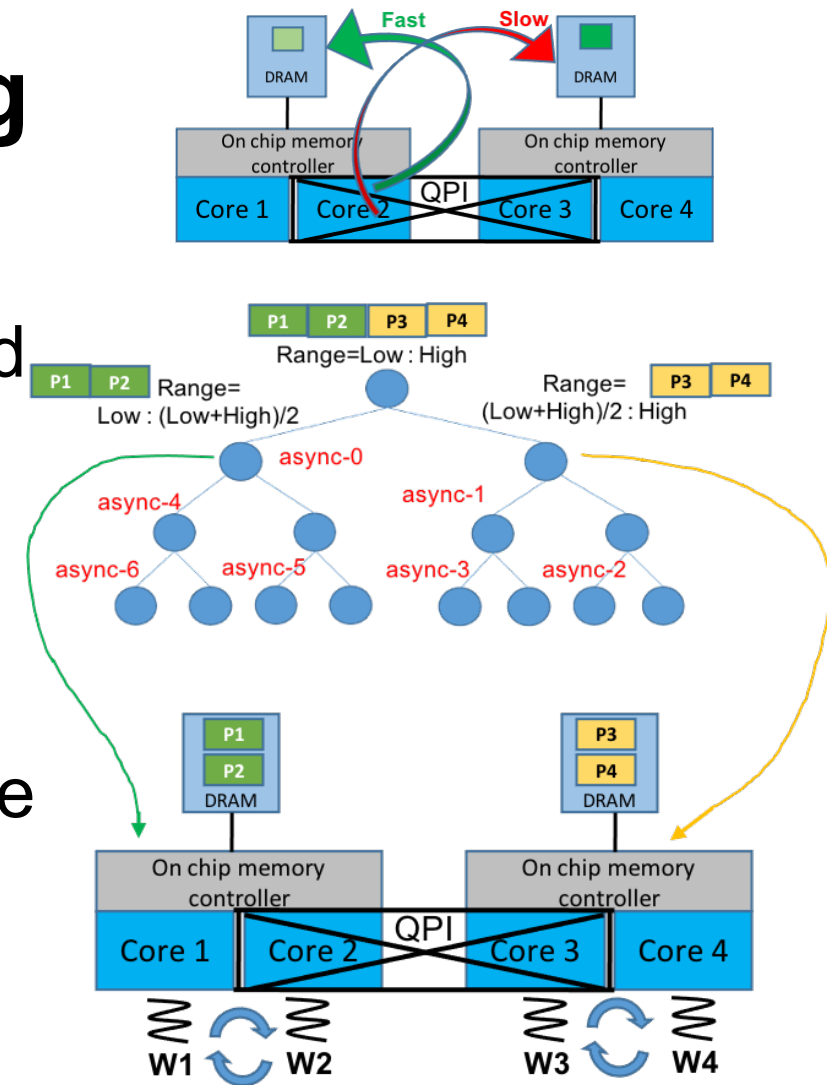
Sequential Overheads: Concurrent Deque



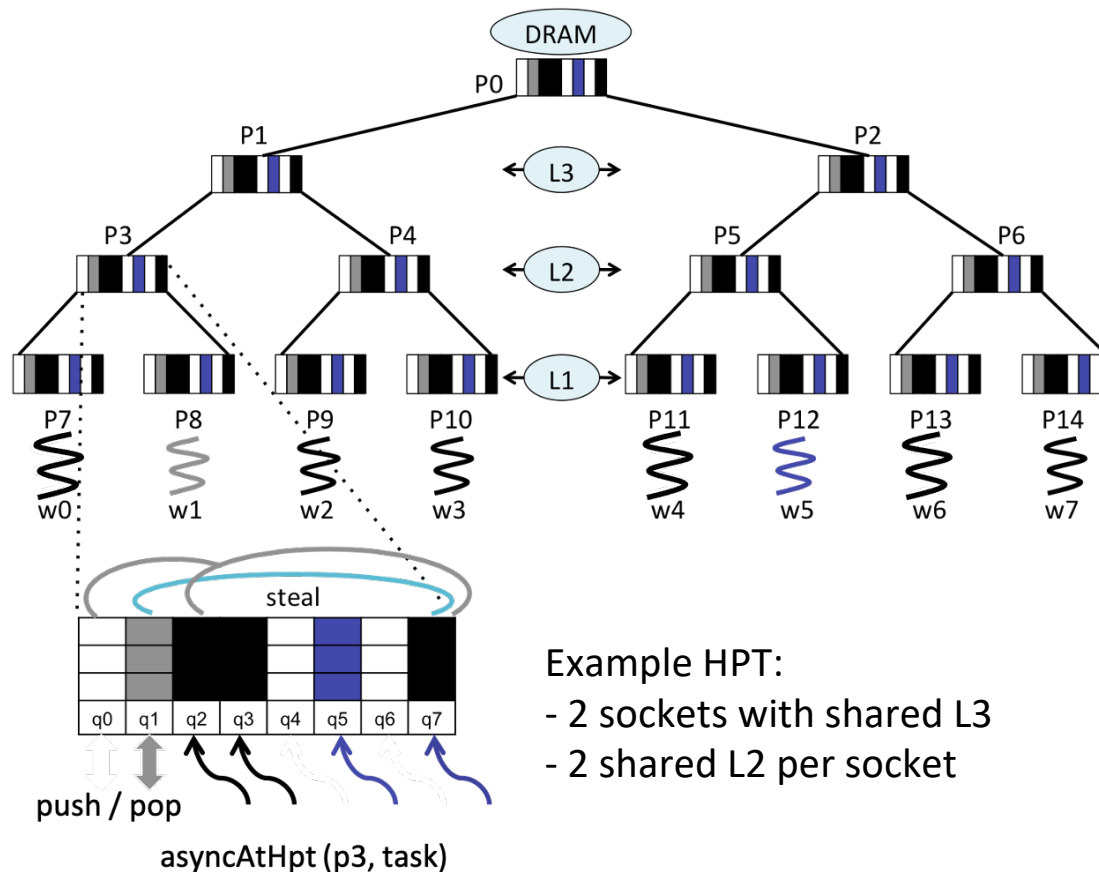
- Minimizing deque overheads
 - Using a mix of list and deque
 - Using private deque

NUMA Aware Work-Stealing

- High performance can be achieved on a NUMA architecture only if the task and its data are collocated, and is local to the worker executing that task
 - By default, Linux uses First-Touch policy for physical page allocation
- Random work-stealing would hurt the locality over NUMA machine due to random victim selection
 - Use hierarchical work-stealing



NUMA Aware Work-Stealing Using HPT




- Round-robin steals instead of random work-stealing
- Workers attach to (own) leaf places
- Each place has one queue per worker
 - Ensures non-synchronized push and pop
- Any worker can push a task at any place
- Pop / steal access permitted to subtree workers
- Workers traverse path from leaf to root
- Tries to pop, then steal, at every place
- After successful pop / steal worker returns to leaf
- Worker threads are bound to cores

Picture credit: Runtime Systems for Extreme Scale Platforms, PhD thesis, Sanjay Chatterjee, Rice University, 2013

Trace/Replay

- Improved locality if each worker executes the exact same set of tasks in each for loop iteration of compute
- Trace/Replay for improving locality
 - Trace (i.e., record) the tasks executed by each worker during the first iteration of for loop inside compute
 - For the rest of iterations of the above for loop of compute, disable random work-stealing and use the information gathered during the Trace (i.e., record) phase to replay the exact set of tasks at each worker



Context Switch Inside Userspace

```
void A() {
    cout<< "IN-A" << endl;
    /* Do something */
    cout<< "OUT-A" << endl;
}
void B() {
    cout<< "IN-B" << endl;
    /* Do something */
    cout<< "OUT-B" << endl;
}
void C() {
    cout<< "IN-C" << endl;
    /* Do something */
    cout<< "OUT-C" << endl;
}
int main() {
    A();
    B();
    C();
}
```

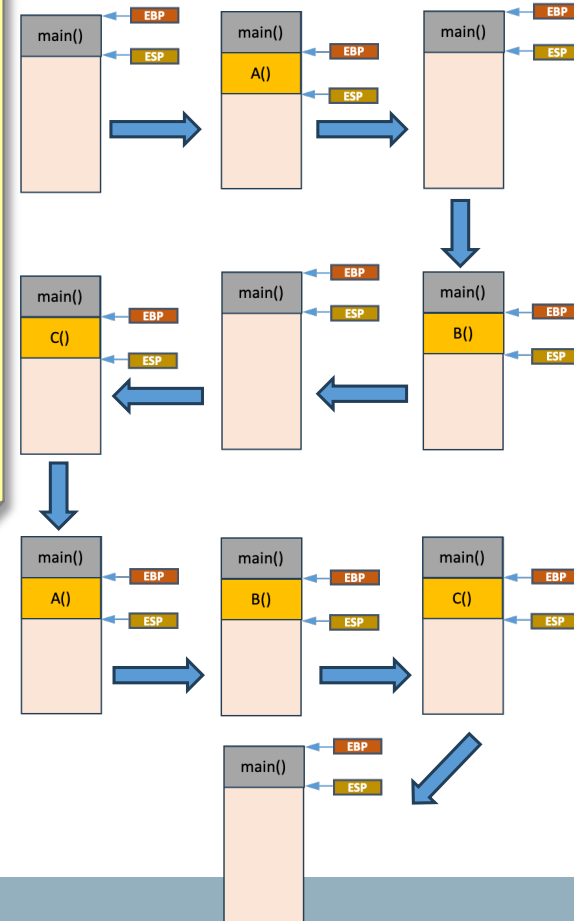
Figure-1

```
#include <boost/context/all.hpp>
ctx::continuation A(ctx::continuation cont) {
    cout<< "IN-A" << endl;
    cont = cont.resume();
    /* Do something */
    cout<< "OUT-A" << endl;
    return std::move(cont);
}
/* Methods B & C rewritten as A above */
int main() {
    ctx::continuation a = ctx::callcc(A);
    ctx::continuation b = ctx::callcc(B);
    ctx::continuation c = ctx::callcc(C);
    a.resume();
    b.resume();
    c.resume();
}
```

Figure-2

Used to switch across different continuations

Call with current continuation. Captures current continuation and triggers a context switch



● **Figure-1**

IN-A
OUT-A
IN-B
OUT-B
IN-C
OUT-C

● **Figure-2**

IN-A
IN-B
IN-C
OUT-A
OUT-B
OUT-C

User Level Threads: Fibers

```
boost::fibers::fiber f1([=]() {
    cout << "A ";
    boost::this_fiber::yield();
    cout << "B ";
    boost::this_fiber::yield();
    cout << "C ";
});
```

```
boost::fibers::fiber f2([=]() {
    cout << "D ";
    boost::this_fiber::yield();
    cout << "E ";
    boost::this_fiber::yield();
    cout << "F ";
});
```


```
f1.join();
f2.join();
```

```
std::mutex mtx;
std::condition_variable cnd;
std::string str;

boost::fibers::fiber f1([=]() {
    std::unique_lock<std::mutex> lck(mtx);
    if(str.size() == 0) {
        cnd.wait(lck);
    }
    cout << str << endl;
});

boost::fibers::fiber f2([=]() {
    std::unique_lock<std::mutex> lck(mtx);
    str = "Hello Fiber";
    cnd.notify_one();
});

f1.join();
f2.join();
```



Midterm Exams

- Midterm exam will be held on 24/02/25 (Monday) 3pm—4pm
 - Total weightage is 20%
 - It is your responsibility to arrive on time. No extra time if you arrive late
 - Closed-notes, closed-book, closed-laptop written exam
 - Syllabus includes Lectures 2–13
 - No penalty for **minor** syntax errors in programming related questions. Minor syntax errors only include **missing semicolon, missing braces, and spell mistakes**.
 - However, you must ensure that your program is: a) clear to understand, and b) has proper indentation. If these two prerequisites are not met, then the marks allocated will be final and reevaluation requests will not be entertained