

Lecture 17: False Sharing

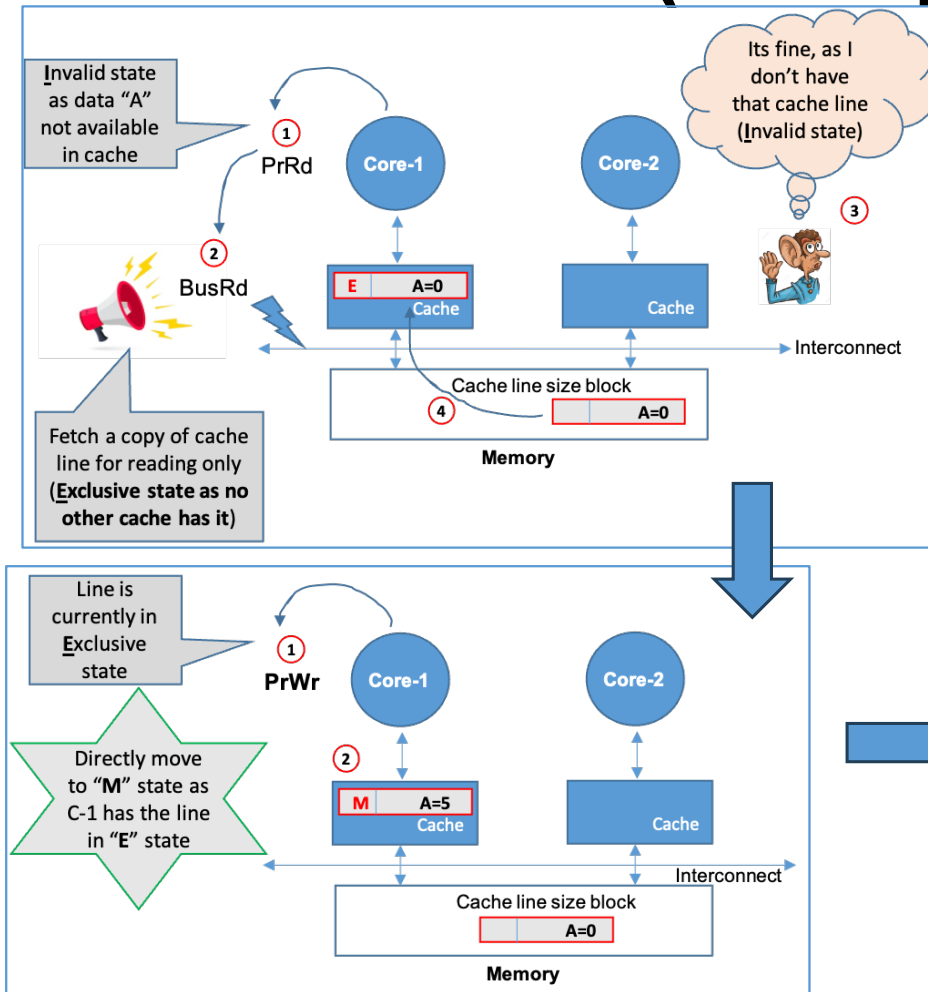
Vivek Kumar

Computer Science and Engineering

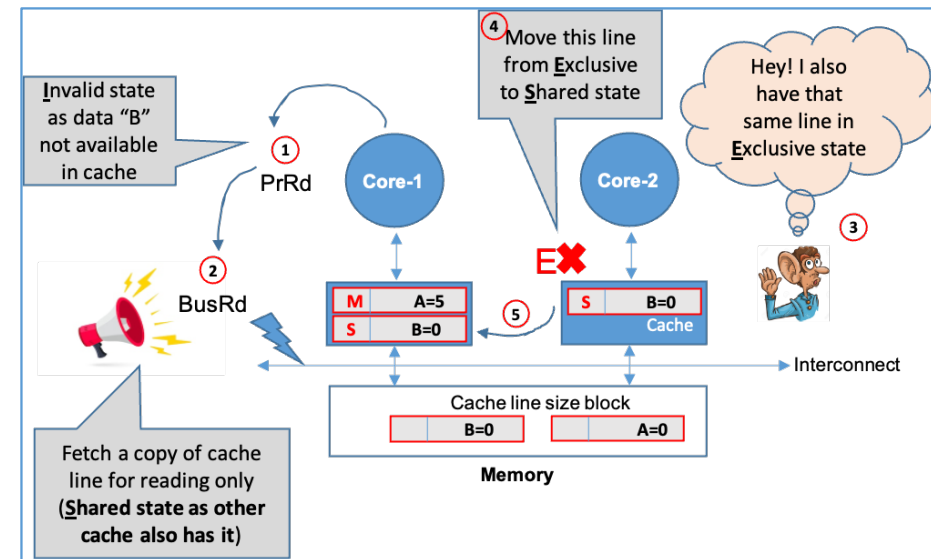
IIIT Delhi

vivekk@iiitd.ac.in

Last Lecture (Recap)



- MESI protocol based cache coherence for write-back private cache



Today's Class

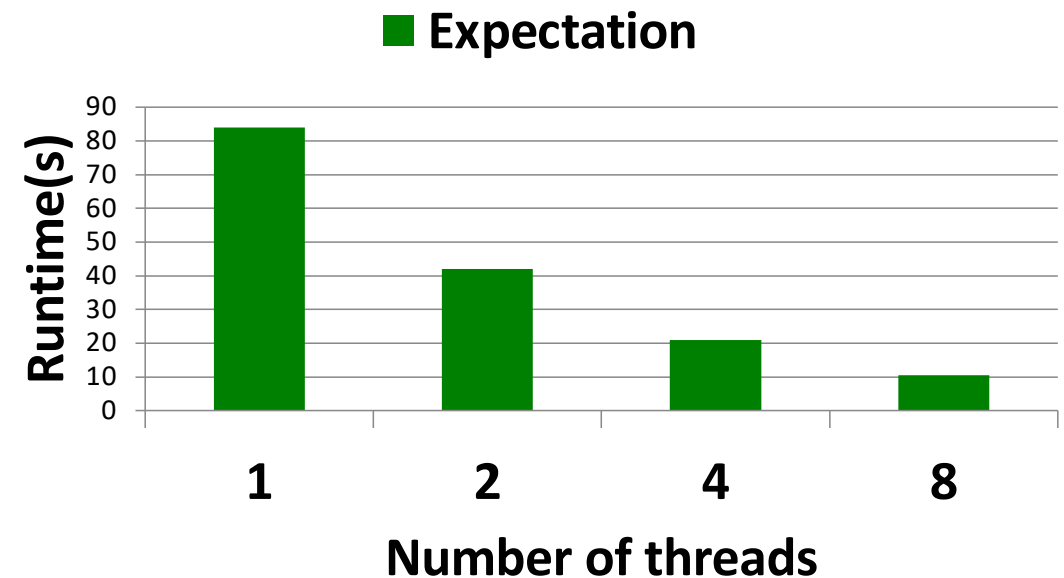
- False sharing
- Runtime solutions for detecting/repairing false sharing
 - Sheriff
 - Featherlight

Acknowledgement: Today's lectures slides are adapted from several conference presentation slides available online on false sharing

Parallel Updates

```
int count[8]; //Global array

thread_func(int id) {
    for(i = 0; i < M; i++)
        count[id]++;
}
```



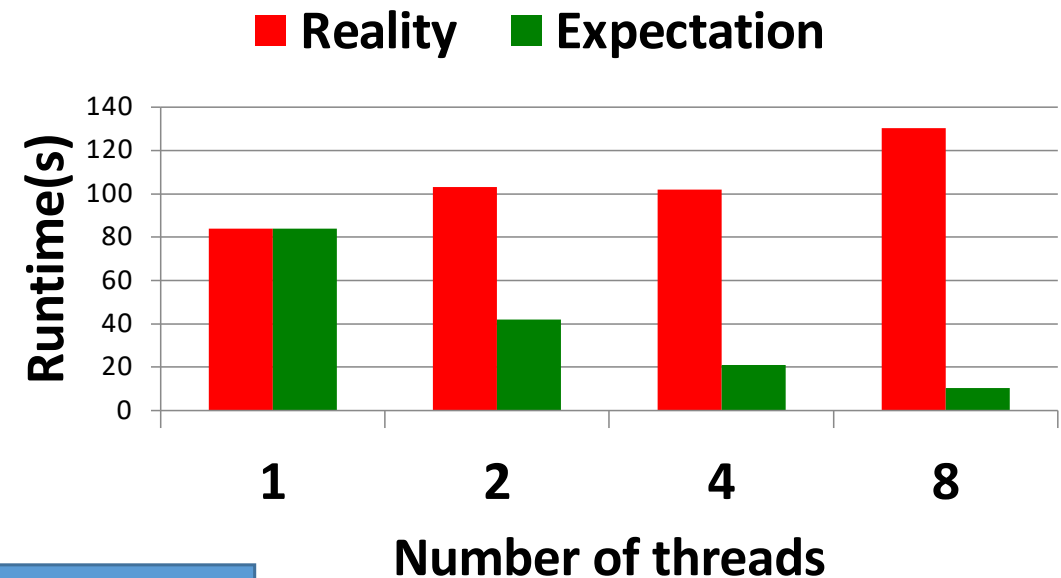
- <https://www.youtube.com/watch?v=NJ46OXN45eU>

False Sharing

```
int count[8]; //Global array

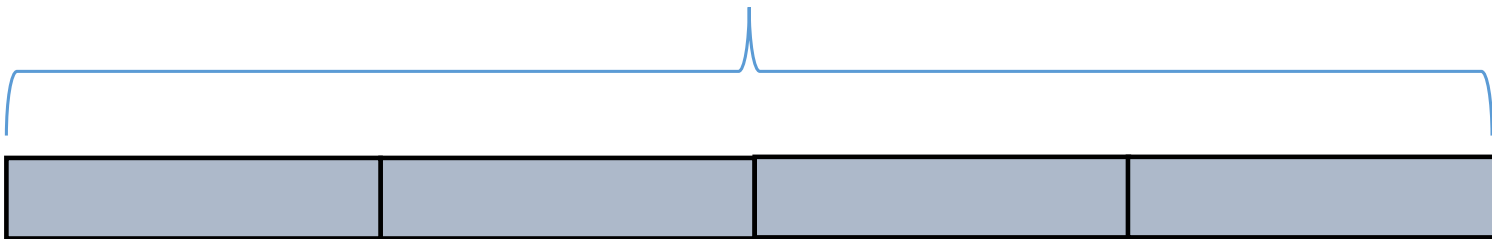
thread_func(int id) {
    for(i = 0; i < M; i++)
        count[id]++;
}
```

Let's try to understand the problem in this code using the MESI coherence protocol

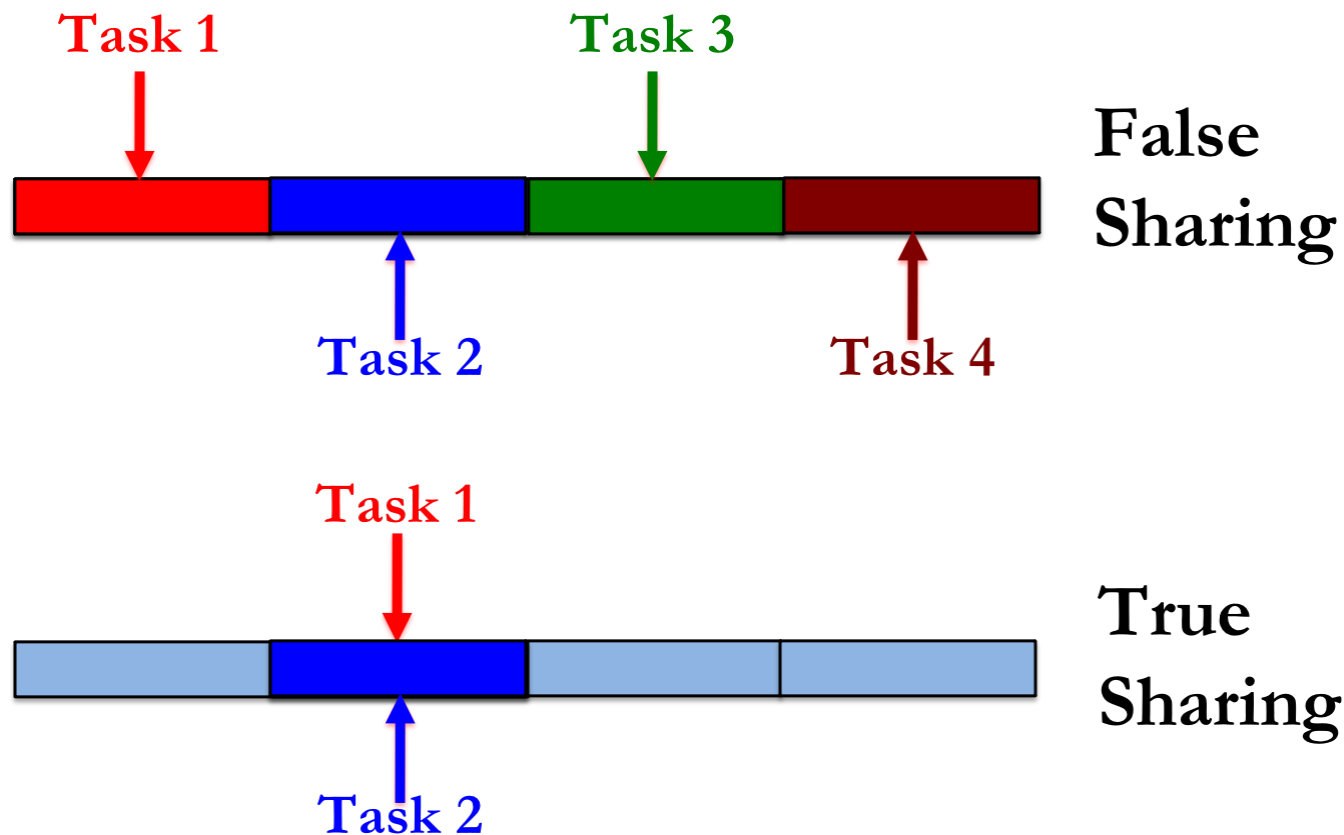


False Sharing vs. True Sharing

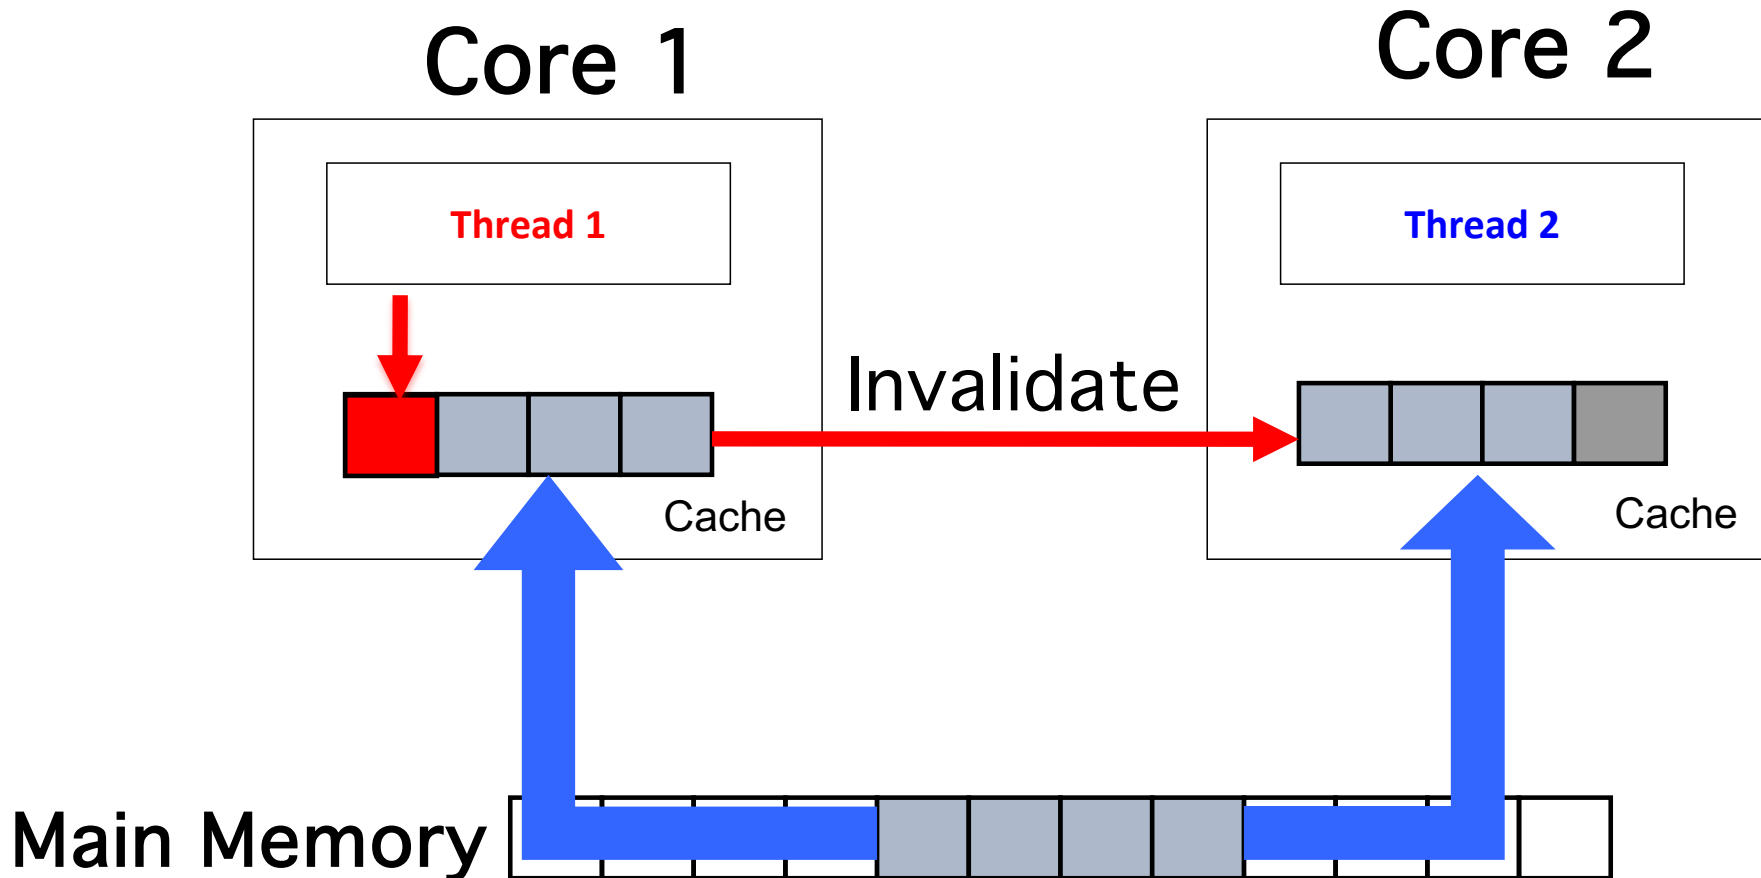
Cache Line



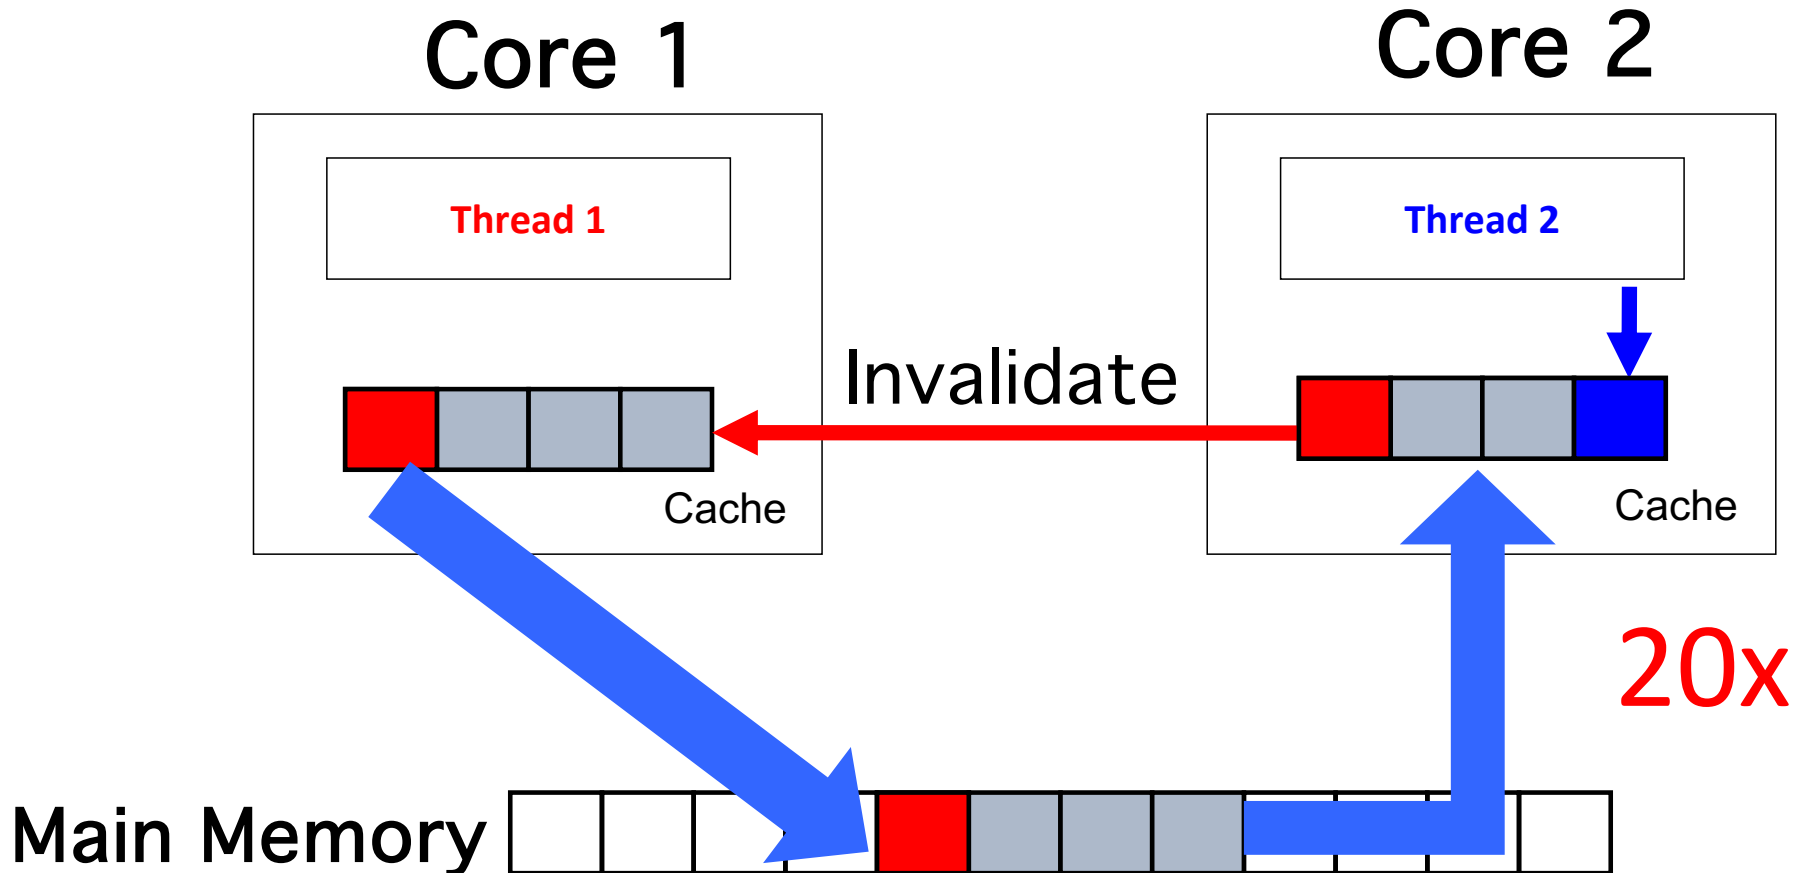
False Sharing vs. True Sharing



False Sharing



False Sharing



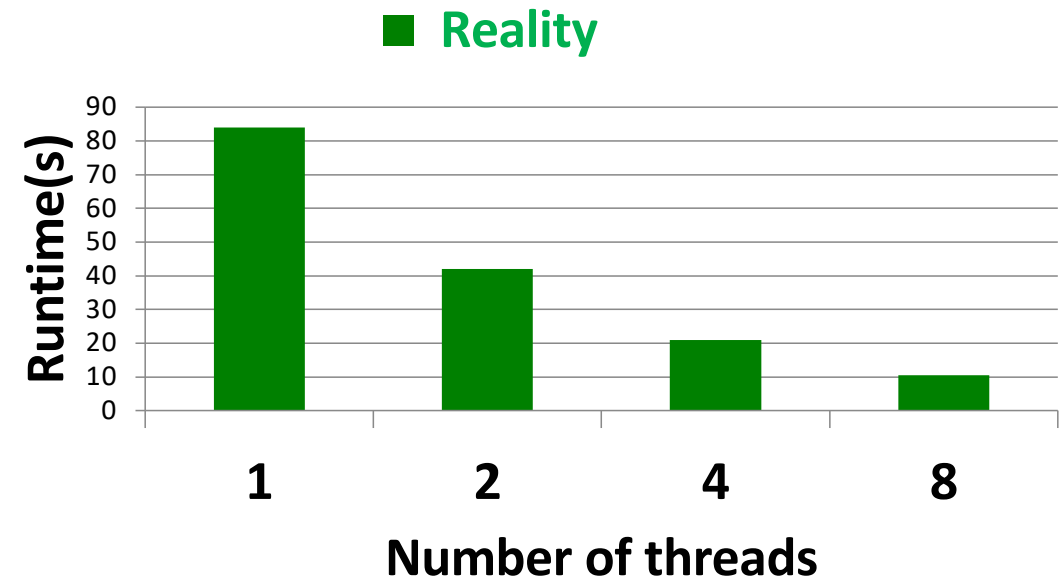
Resource Contention at Cache Line Level



Parallel Updates – How to Fix it Manually?

```
int count[8]; //Global array
```

```
thread_func(int id) {  
    for(i = 0; i < M; i++)  
        count[id]++;  
}
```



False Sharing is Everywhere

```
me = 1;  
you = 1; // globals
```

```
me = new Foo;  
you = new Bar; // heap
```

```
class X {  
    int me;  
    int you;  
}; // fields
```

```
arr[me] = 12;  
arr[you] = 13; // array indices
```

Two different threads
T1 & T2 are involved

Detecting False Sharing

- False sharing on a cache line implies that particular cache line will incur large number of cache invalidations
- A runtime approach can track memory access using hardware performance counters

Automatically Fixing the False Sharing

- Core idea
 - Identify and fix memory locations that could lead to false sharing
 - Approaches
 - Using compilation technique
 - Use static analysis of the application and identify the locations of potential false sharing
 - Can emit memory padding to avoid false sharing
 - Limited to eliminating false sharing for programs where the size and location of data elements can be determined statically
 - Use runtime technique
 - Overcomes the above limitations
 - Runtime checks might inflate the execution time
 - Let us try to understand one of the well-known runtime implementations (**Sheriff**: <https://github.com/plasma-umass/sheriff/tree/master>)

Detour – Process Creation

```

int read_var[1024]; //page aligned (4Kb size)
int write_var[1024]; //page aligned (4Kb size)

int main() {
    int status = fork();
    if(status < 0) {
        printf("Something bad happened\n");
    } else if(status == 0) {
        printf("I am the child process\n");
        write_var[0] = 2 * read_var[0];
    } else {
        printf("I am the parent Shell\n");
        write_var[0] = 4 * read_var[0];
    }
    printf("write_var value = %d\n", write_var[0]);
    ....
    return 0;
}

```

- **fork** is a system call used for creating a new process
- Called once, but returns twice!
 - Return value in child process is zero, whereas child's process PID is returned in parent process
- It creates a replica of the parent process
 - Copy-on-Write (COW) – Initially, both parent and child process have read-only access to parent's address space. Whichever process attempts a write on a memory page in parent's address space, it would get a copy of that page (lazy copy)

What could be our Approach?

```
1. /* global variables */
2. int sum[2];
3. int main() {
4.     /* heap allocations */
5.     int a = new int[size]; //initialized
6.     T1 = new thread(=]() {
7.         for(int i=0; i<size/2; i++) sum[0]+=a[i];
8.     });
9.     T2 = new thread(=]() {
10.        for(int i=size/2; i<size; i++) sum[1]+=a[i];
11.    });
12.    T1.join(); T2.join();
13.    //Cleanups
14. }
```

False sharing

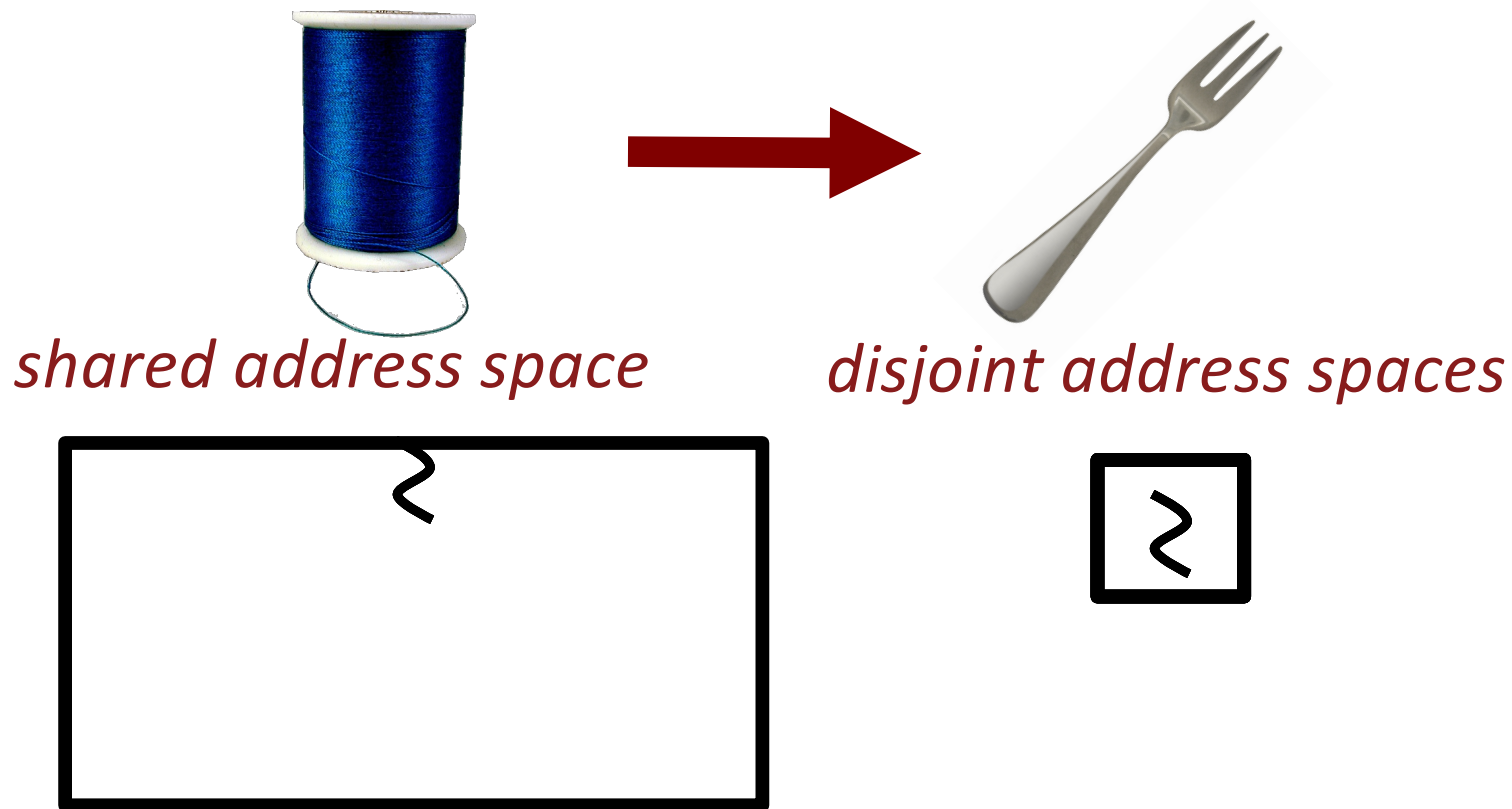
False sharing

Walkthrough of Sheriff Execution

Process creation – creating processes instead of threads

```
1. /* global variables */
2. int sum[2];
3. int main() {
4.     /* heap allocations */
5.     int a = new int[size]; //initialized
6.     T1 = new thread( [=]() {
7.         for(int i=0; i<size/2; i++) sum[0]+=a[i];
8.     });
9.     T2 = new thread( [=]() {
10.        for(int i=size/2; i<size; i++) sum[1]+=a[i];
11.    });
12.    T1.join(); T2.join();
13.    //Cleanups
14. }
```

Sheriff Execution: Process Creation



- In Linux, both pthreads and processes are essentially a KLT, and are created using the same API (`do_fork`)

Walkthrough of Sheriff Execution

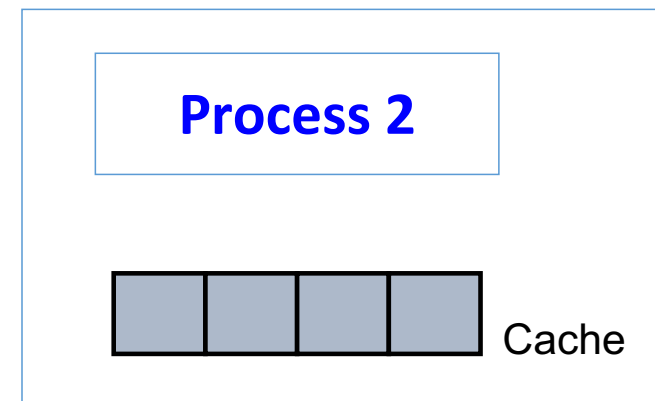
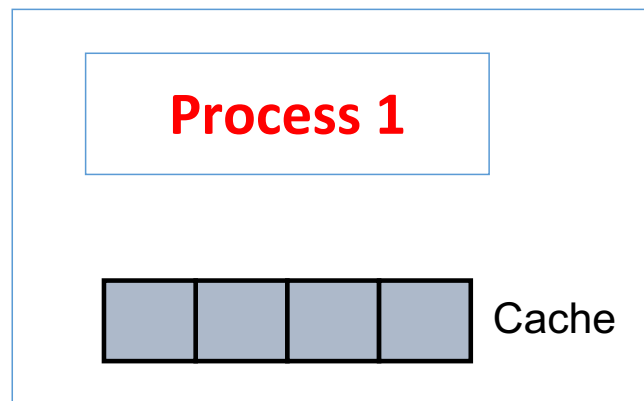
Initialization – creating mapping of global and heap variables for processes

```
1.  /* global variables */
2.  int sum[2];
3.  int main() {
4.      /* heap allocations */
5.      int a = new int[size]; //initialized
6.      T1 = new thread(=]() {
7.          for(int i=0; i<size/2; i++) sum[0]+=a[i];
8.      });
9.      T2 = new thread(=]() {
10.         for(int i=size/2; i<size; i++) sum[1]+=a[i];
11.     });
12.     T1.join(); T2.join();
13.     //Cleanups
14. }
```

Sheriff Execution: Initialization

Core 1

Core 2



Each process operates on private copies of data

Global State



Main Memory

Sheriff Execution: Initialization

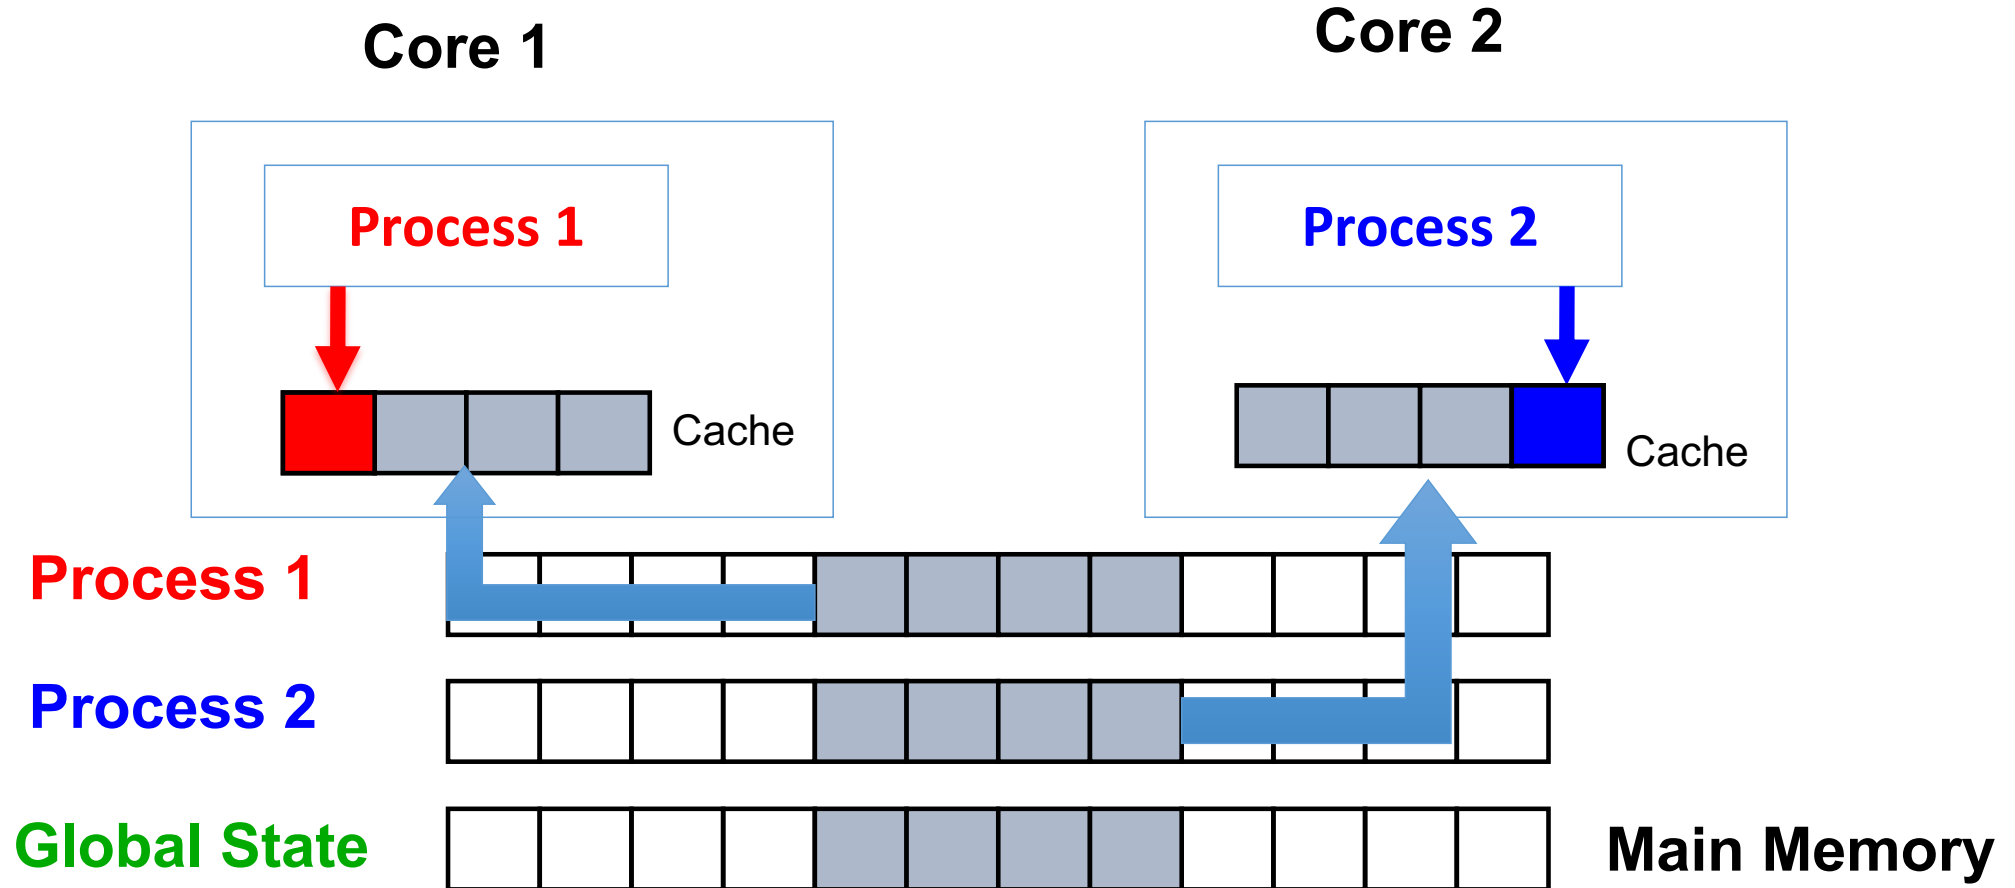
- Advantages of converting threads into processes
 - Enables the use of per-thread page protection, allowing Sheriff to track memory accesses by different threads (processes)
 - Each thread's (process) memory access are isolated, hence they would not update the same cache line
 - No false sharing!
- Memory mapped files are used to share global and heaps across different processes
- Twin copies of the pages for storing the global and heaps
 - Shared mapping for holding shared states
 - Pages storing these shared states are marked copy-on-write
 - Private mapping for per-process updates
 - Private copy of of the above shared pages are created whenever a process would attempt to update a page for the first time

Walkthrough of Sheriff Execution

Execution – copies of memory pages per process for local updates

```
1. /* global variables */
2. int sum[2];
3. int main() {
4.     /* heap allocations */
5.     int a = new int[size]; //initialized
6.     T1 = new thread( [=]() {
7.         for(int i=0; i<size/2; i++) sum[0]+=a[i];
8.     });
9.     T2 = new thread( [=]() {
10.        for(int i=size/2; i<size; i++) sum[1]+=a[i];
11.    });
12.    T1.join(); T2.join();
13.    //Cleanups
14. }
```

Sheriff Execution: Execution



Walkthrough of Sheriff Execution

```
1. /* global variables */
2. int sum[2];
3. int main() {
4.     /* heap allocations */
5.     int a = new int[size]; //initialized
6.     T1 = new thread( [=]() {
7.         for(int i=0; i<size/2; i++) sum[0]+=a[i];
8.     });
9.     T2 = new thread( [=]() {
10.        for(int i=size/2; i<size; i++) sum[1]+=a[i];
11.    });
12.    T1.join(); T2.join();
13.    //Cleanups
14. }
```

Synchronization – merging
diffs in per-process pages
into the global copy

Sheriff Execution: Synchronization

- There are two different types of synchronization points
 - Thread termination
 - End of the critical section (mutex unlock), barriers, etc.
- At each synchronization point, Sheriff commits changes from private pages to the shared pages
 - It commits only the differences between the twin and the modified pages

Sheriff Execution: Synchronization



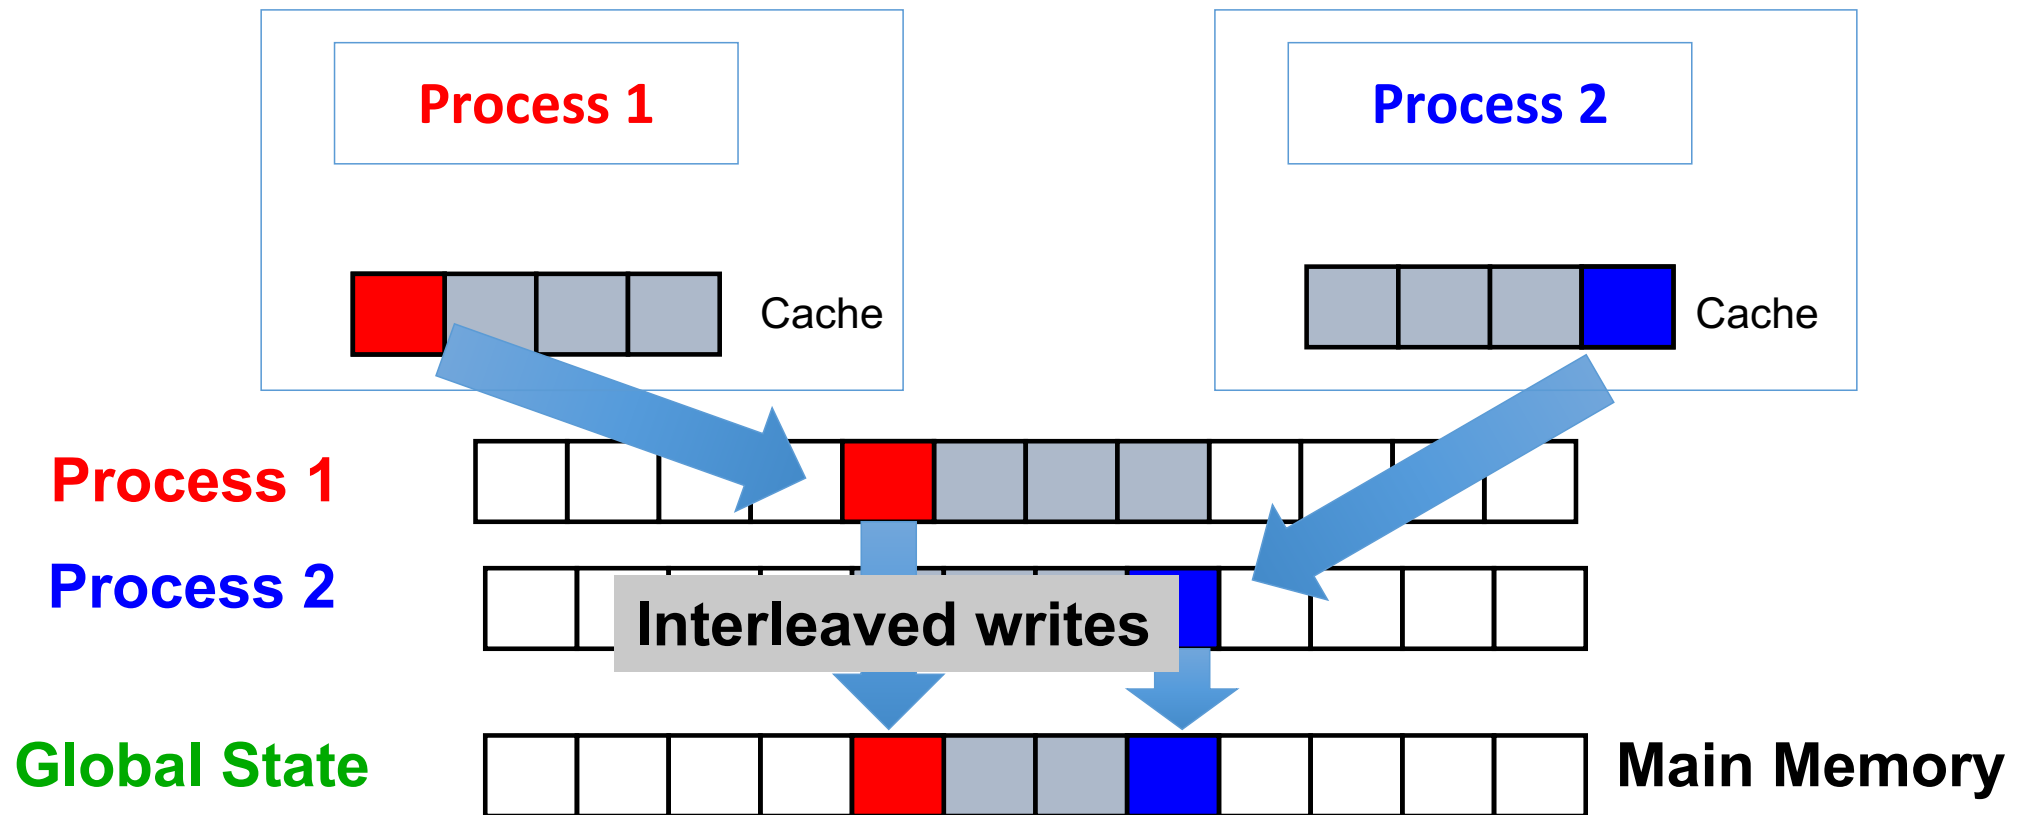
Snapshot and diffing
the local changes



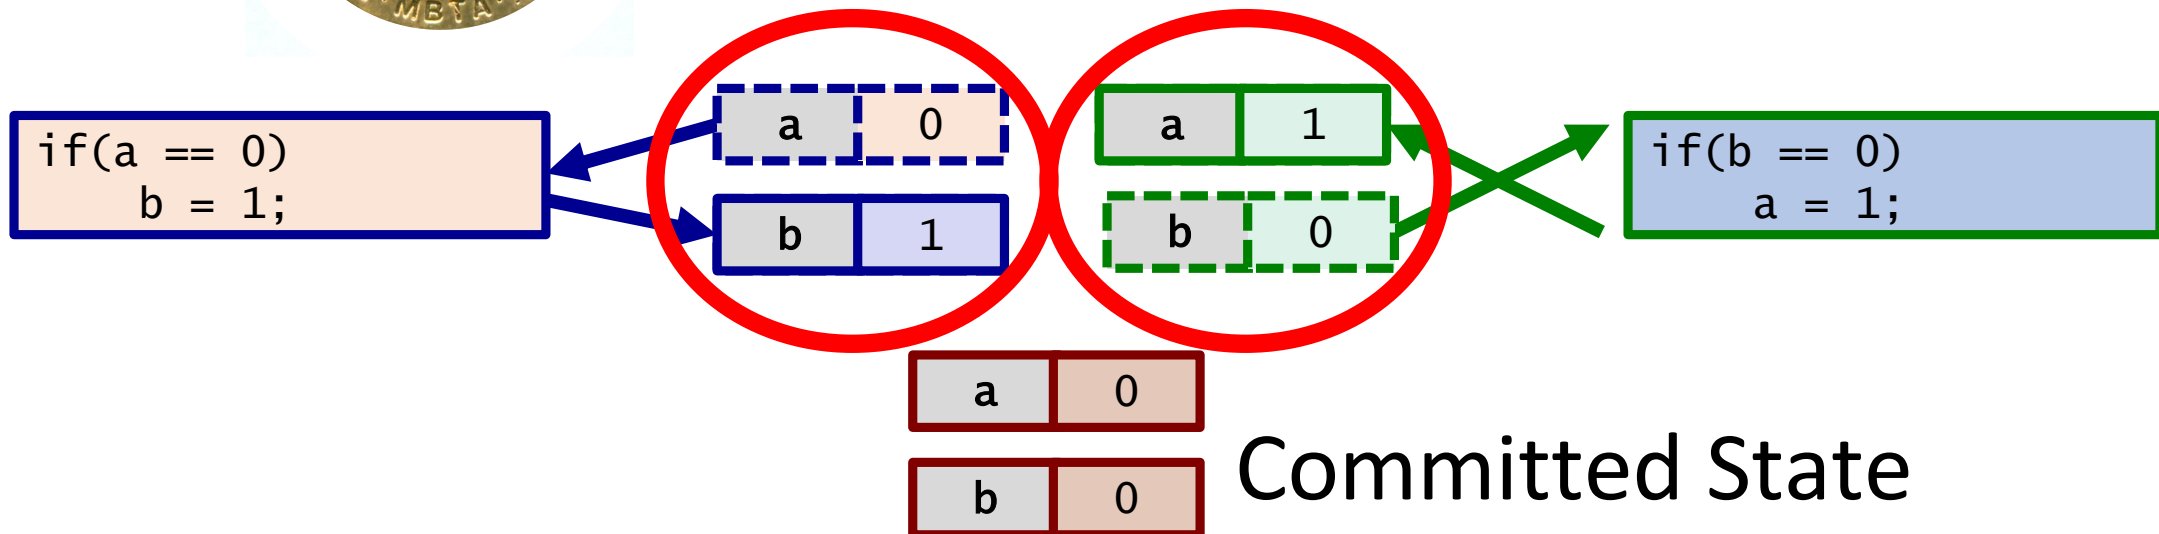
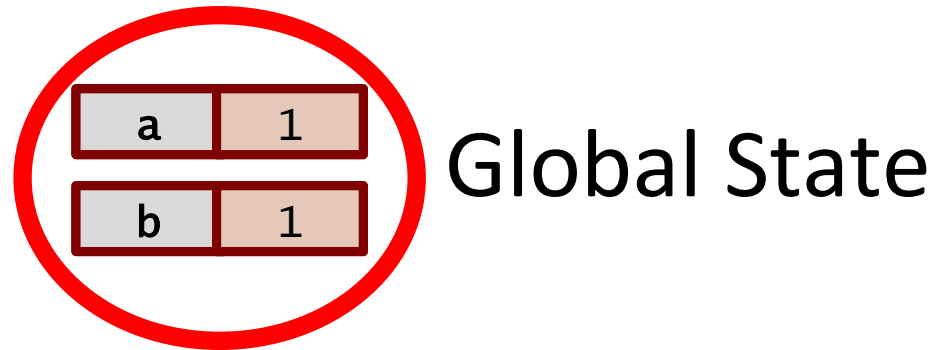
Sheriff Execution: Synchronization

Core 1

Core 2



Sheriff Execution: Synchronization



Reading Materials

- Sheriff
 - <https://people.umass.edu/tongping/pubs/sheriff-oopsla11.pdf>

Next Lecture

- Power management in multicore processors
- **Quiz-4**
 - **Syllabus: Lectures 14-17**