

KarmaPM: Reward-Driven Power Manager

Author Name: **Sunil Kumar** and Vivek Kumar

IIIT Delhi, India

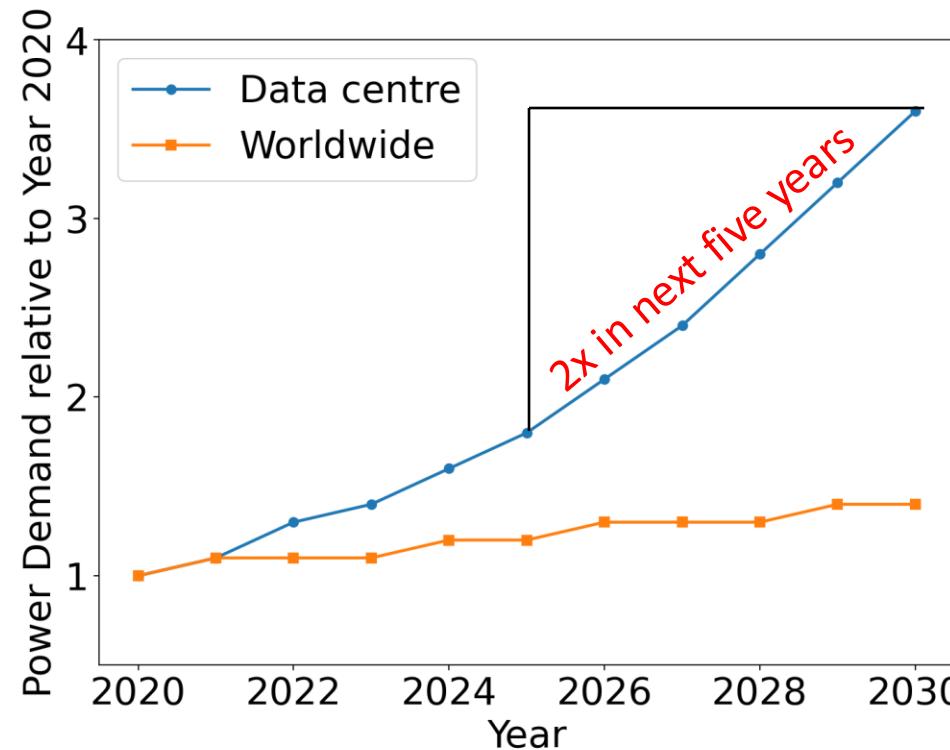


Outline

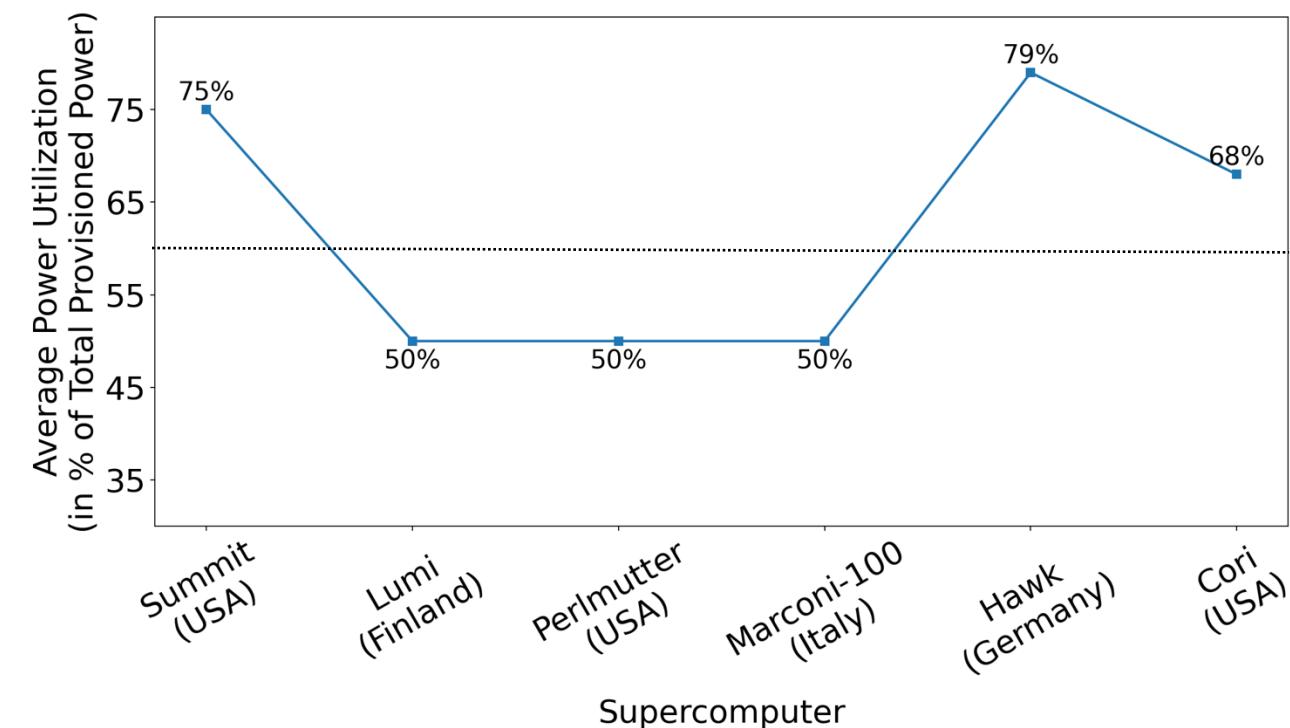
- ✓ Introduction
- ✓ Motivation
- ✓ Related Work
- ✓ Contribution
- ✓ Implementation
- ✓ Results
- ✓ Summary

Global Perspective on Computing Power

Power demand in the data centres



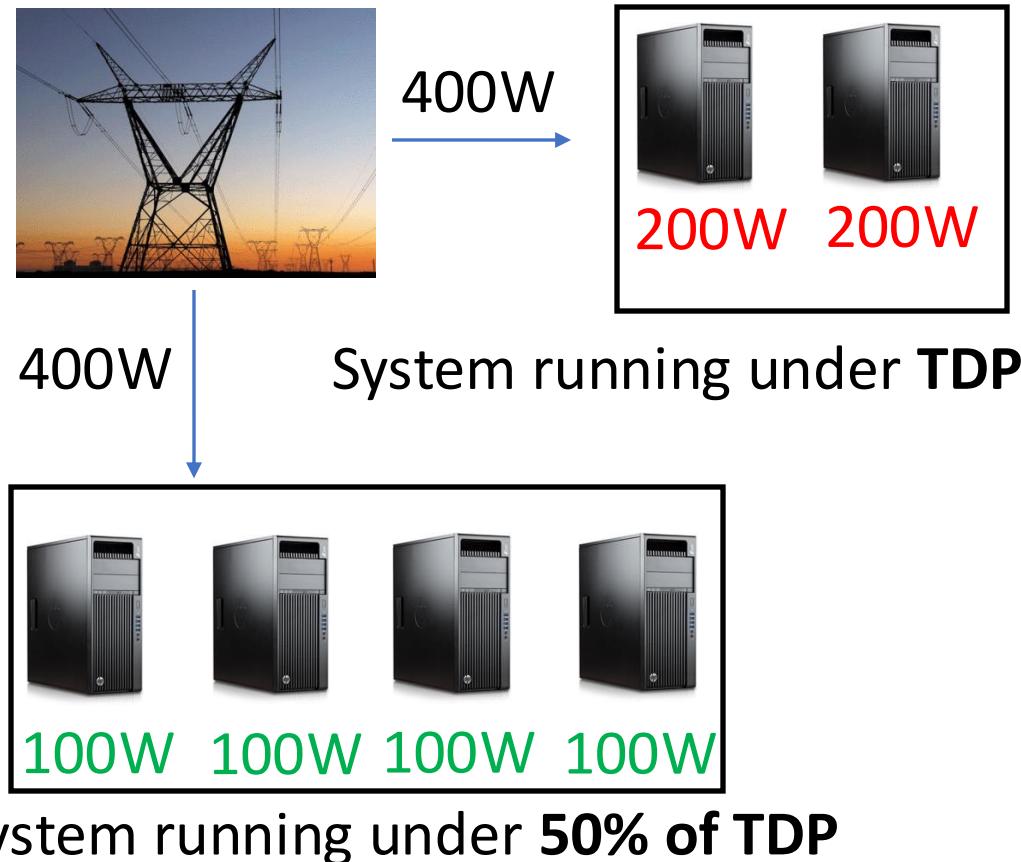
Power usage at supercomputers



1. <https://www.iea.org/reports/energy-and-ai>
2. Patki et.al. [ICS2025]

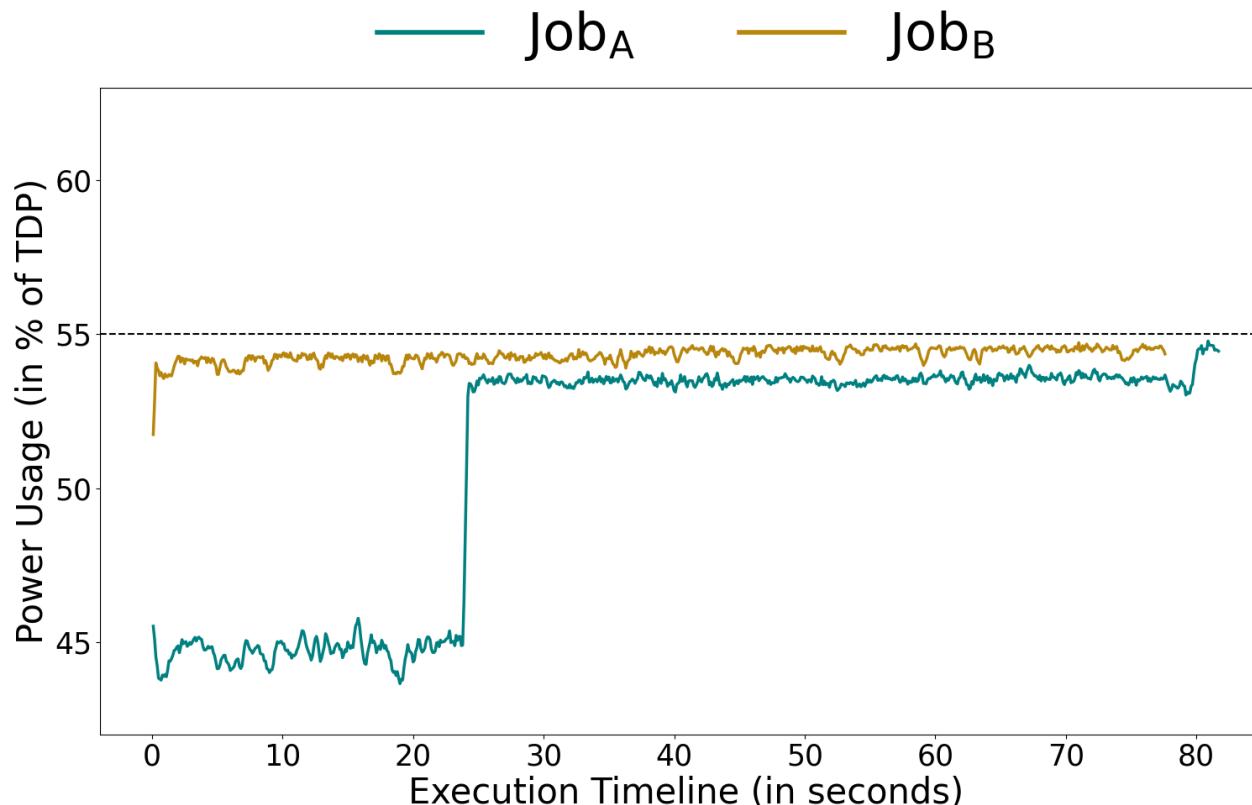
It is extremely essential to improve power efficiency

Hardware Overprovisioning using Power cap



- Servers are designed to operate within the Thermal Design Power (TDP) limit
 - TDP is the maximum power limit
- Power capping (PCAP) restricts power usage below TDP
 - Allows using more servers within the same power budget

Issues with Power Capping



Power usage changes throughout the application execution

Job_B fully utilizes the available power, whereas Job_A only partially utilizes it

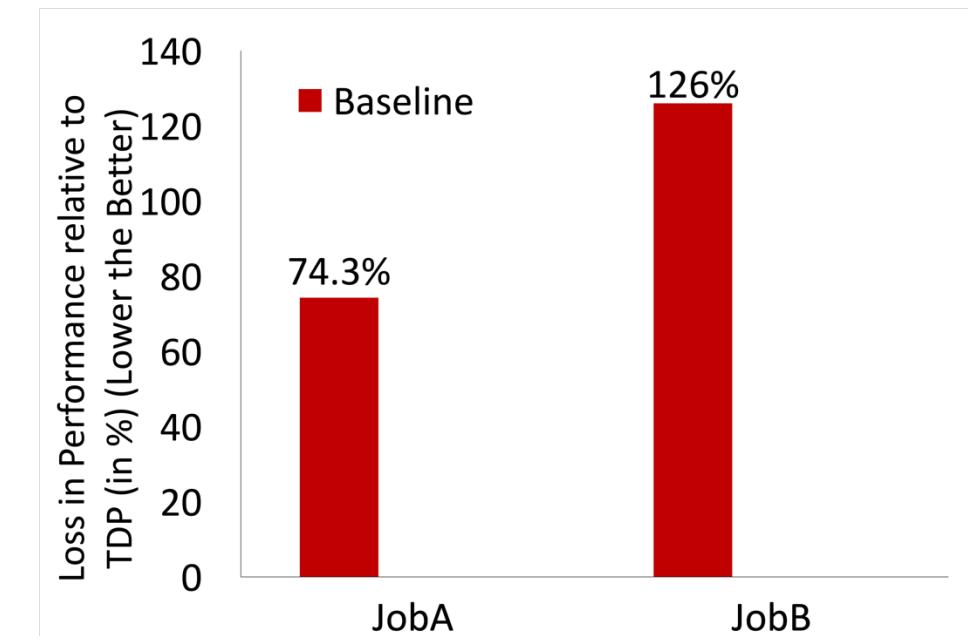
Issues with Power Capping



Power usage changes throughout the application execution

Job_B fully utilizes the available power, whereas Job_A only partially utilizes it

Issues with Power Capping

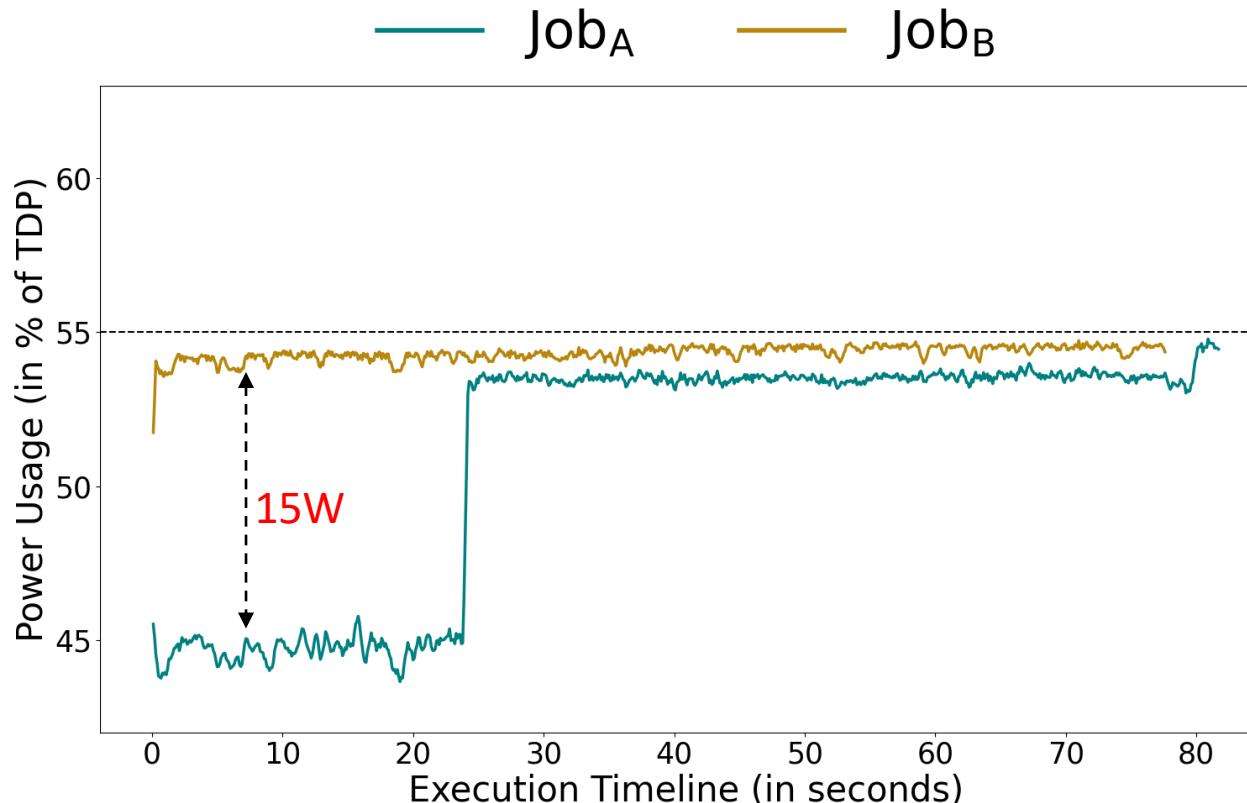


Power usage changes throughout the application execution

Job_B fully utilizes the available power, whereas Job_A only partially utilizes it

Existing Approaches

Improving Performance under PCAP



Overall system throughput improved by 3.3% (geometric mean of speedup of each application over baseline)

Co-running applications on multi-socket servers provides an opportunity to reduce power wastage

Improving Performance under PCAP

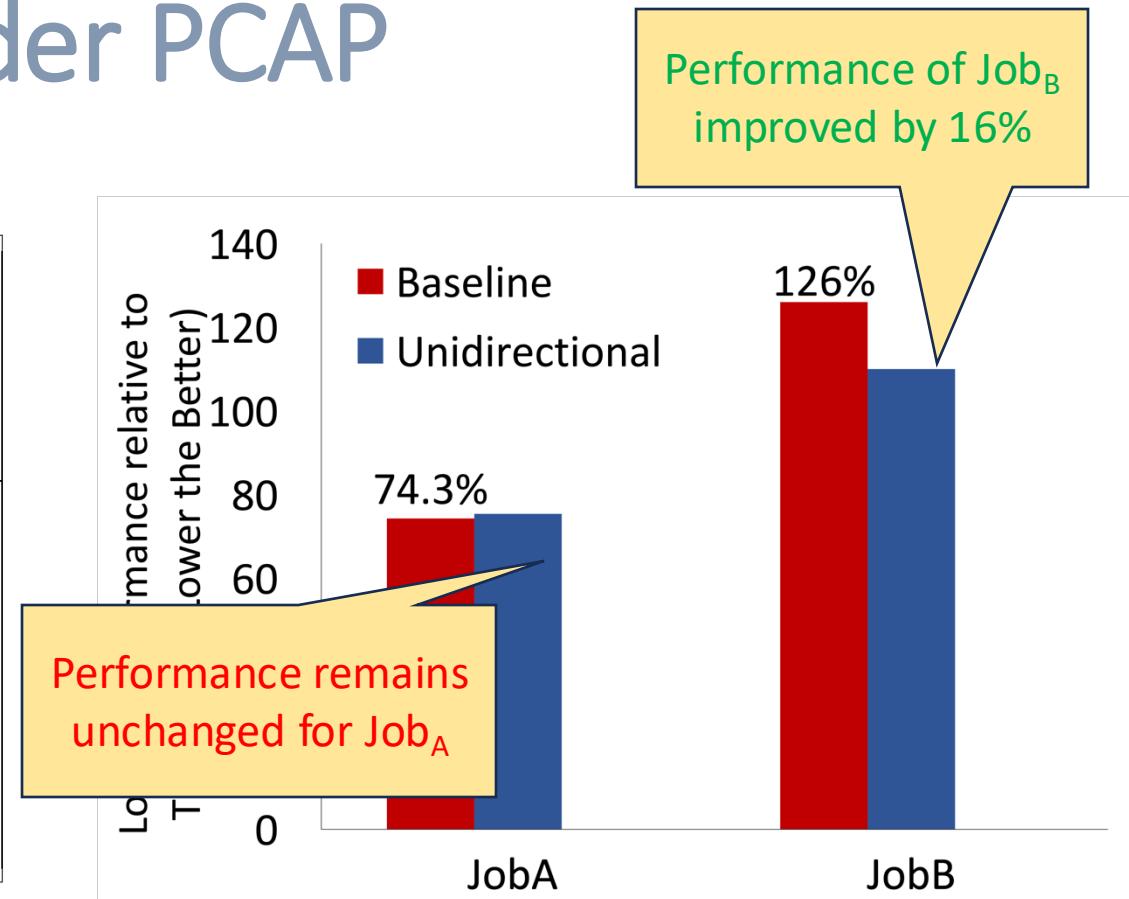
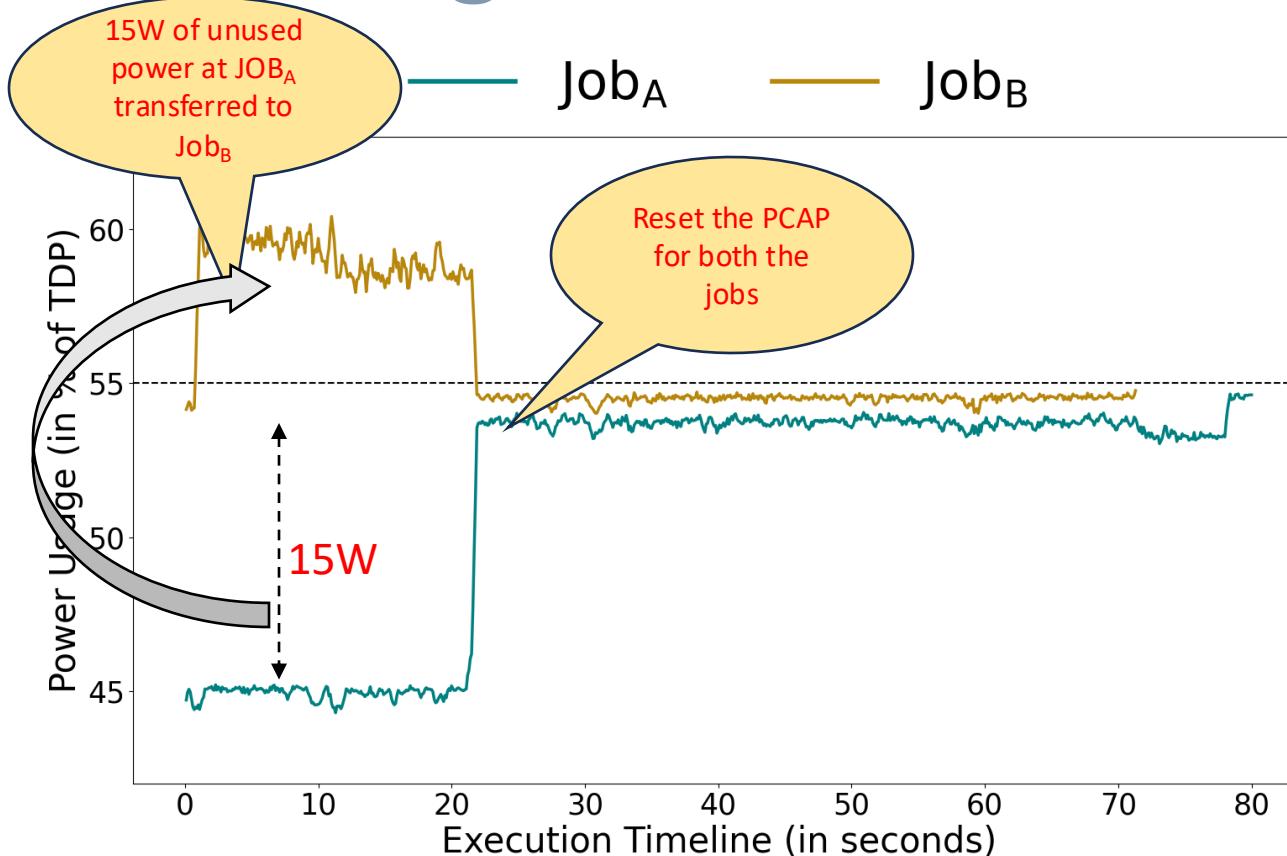


Overall system throughput improved by 3.3% (geometric mean of speedup of each application over baseline)

Co-running applications on multi-socket servers provides an opportunity to reduce power wastage

Existing Approaches

Improving Performance under PCAP

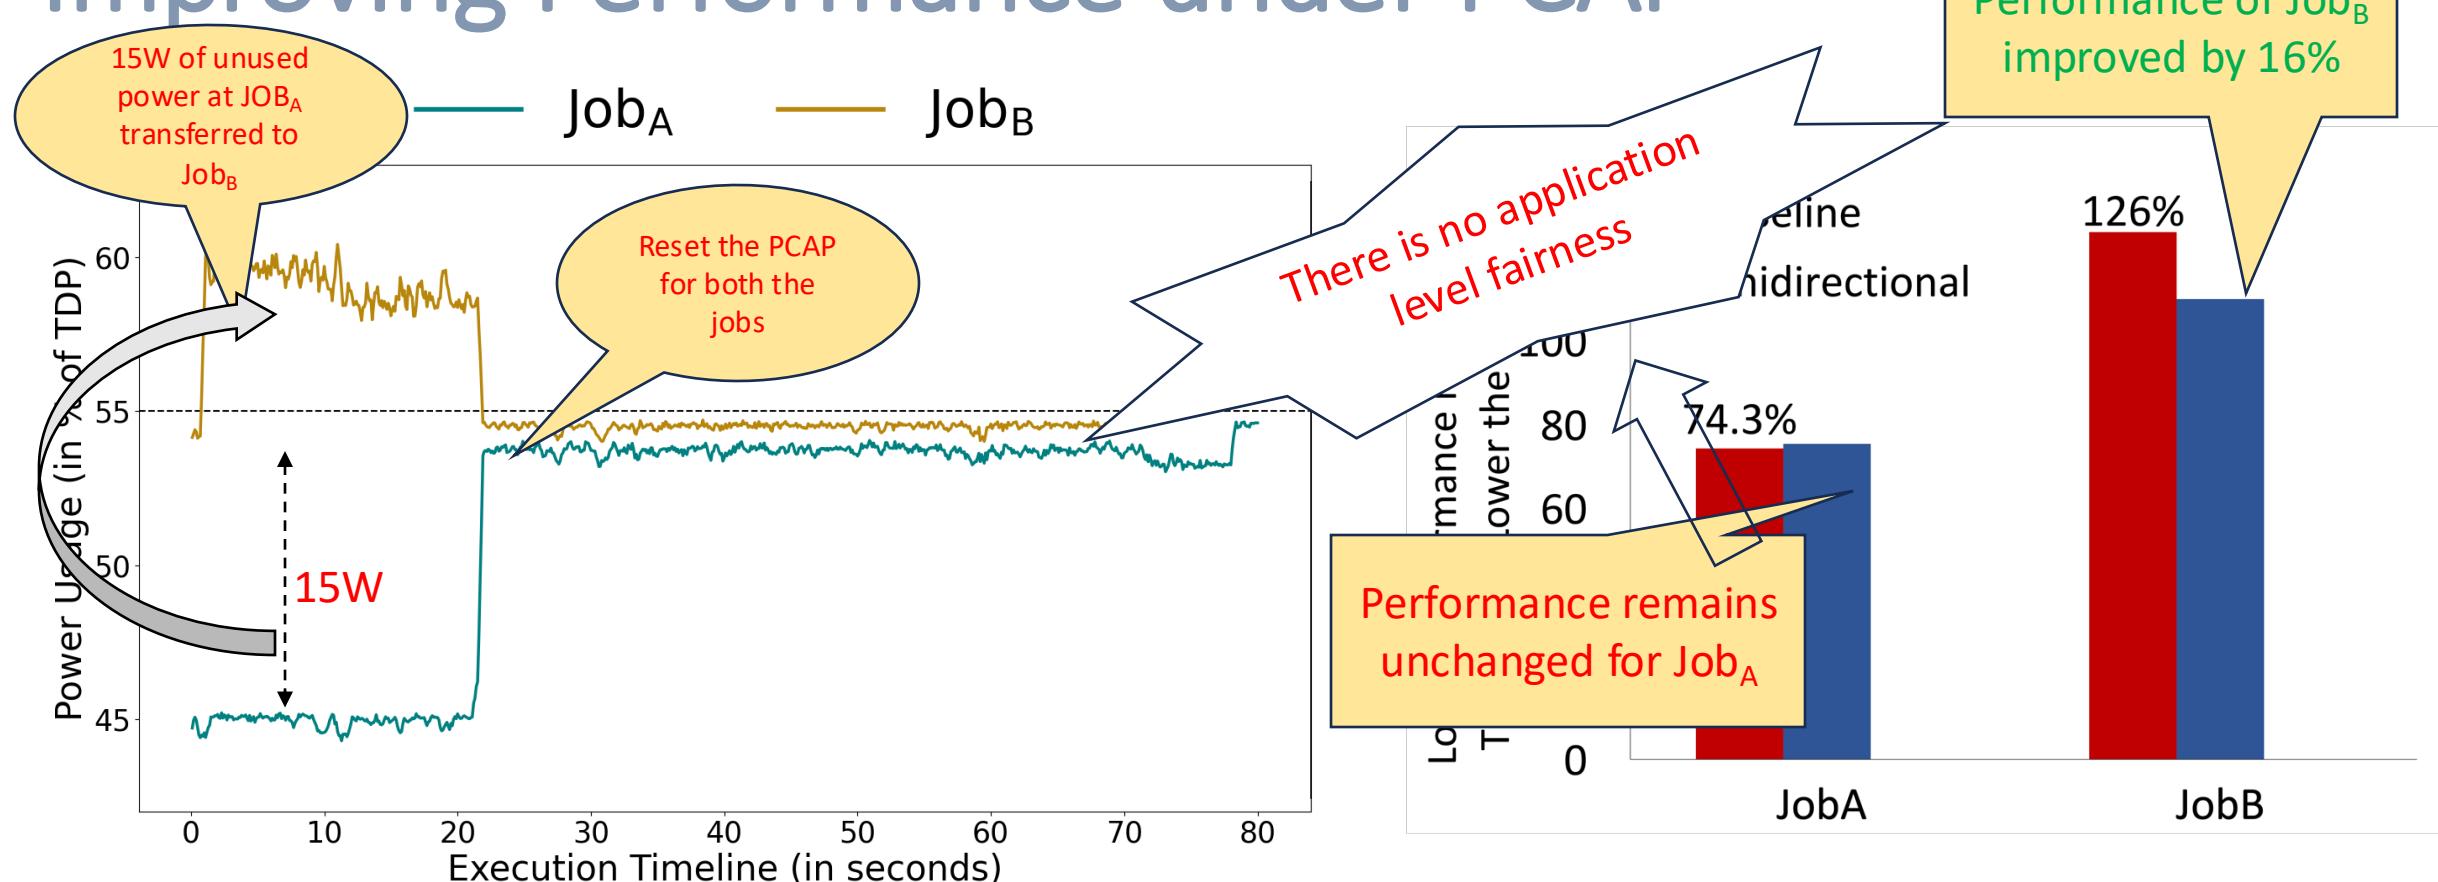


Overall system throughput improved by 3.3% (geometric mean of speedup of each application over baseline)

Co-running applications on multi-socket servers provides an opportunity to reduce power wastage

Existing Approaches

Improving Performance under PCAP



Overall system throughput improved by 3.3% (geometric mean of speedup of each application over baseline)

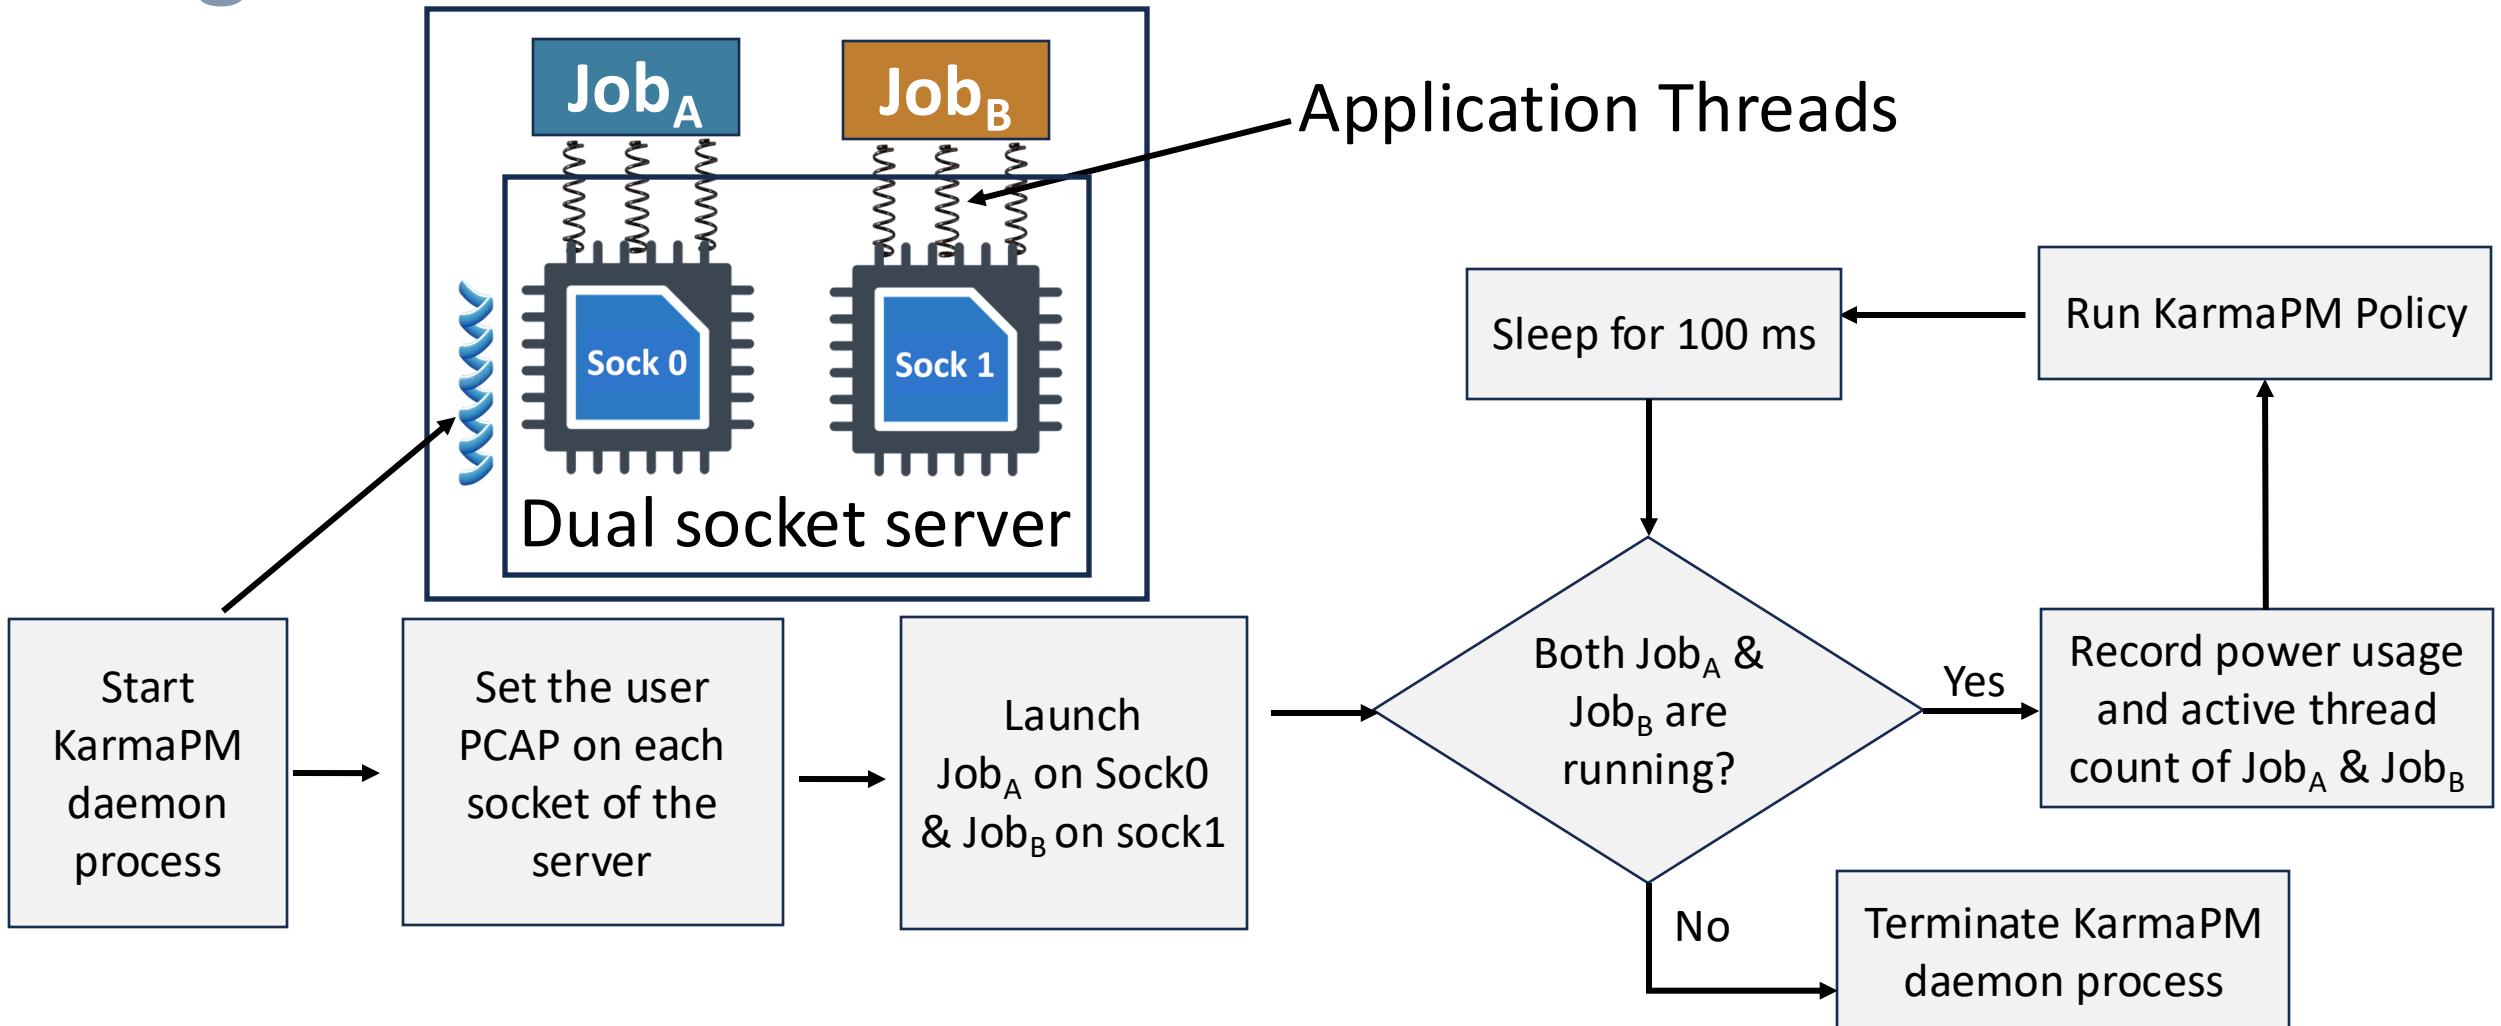
Co-running applications on multi-socket servers provides an opportunity to reduce power wastage

Contributions

- ✓ **KarmaPM: A library-based power management system**
 - ✓ A light-weight daemon that dynamically reallocates power by profiling hardware performance counters
 - ✓ ML model-free and oblivious to the parallel programming models
- ✓ **Enables bi-directional power transfer between co-running jobs**
 - ✓ A novel reward mechanism that improves both throughput and fairness
- ✓ **Experimental evaluations on a quad-socket 72-core Intel Xeon processor**
 - ✓ Using several exascale proxy applications (MPI, OpenMP and Kokkos)
- ✓ **Results**
 - ✓ Our results show that KarmaPM can substantially improve the system throughput and application-level fairness.

Implementation

High-level Architecture of KarmaPM



Implementation

KarmaPM Policy



Implementation

KarmaPM Policy



- Job_A starts with low power usage, Job_B with high power usage
- KarmaPM transfers **15W** surplus power from Socket0 to Socket1 at $t=0$
 - *Similar to existing approaches*
- After **22s**, Job_A resumes using full power
 - KarmaPM resets PCAP on the server
- KarmaPM rewards Job_A by returning 50% of the previously transferred power to Job_B for the same duration (next 22s)
 - **Provides application-level fairness**

Implementation

KarmaPM Policy



- Job_A starts with low power usage, Job_B with high power usage
- KarmaPM transfers **15W** surplus power from Socket0 to Socket1 at $t=0$
 - *Similar to existing approaches*
- After **22s**, Job_A resumes using full power
 - KarmaPM resets PCAP on the server
- KarmaPM rewards Job_A by returning 50% of the previously transferred power to Job_B for the same duration (next 22s)
 - **Provides application-level fairness**
- Execution continues with the user-set PCAP at JobA and JobB after $t=44$

Experimental Methodology

Exascale OpenMP proxy applications

- ✓ Pennant ✓ SimpleMOC (MPI)
- ✓ MiniFE (Kokkos) ✓ PathFinder
- ✓ Quicksilver ✓ RS Bench
- ✓ CoMD (Kokkos) ✓ CG (NPB suite)

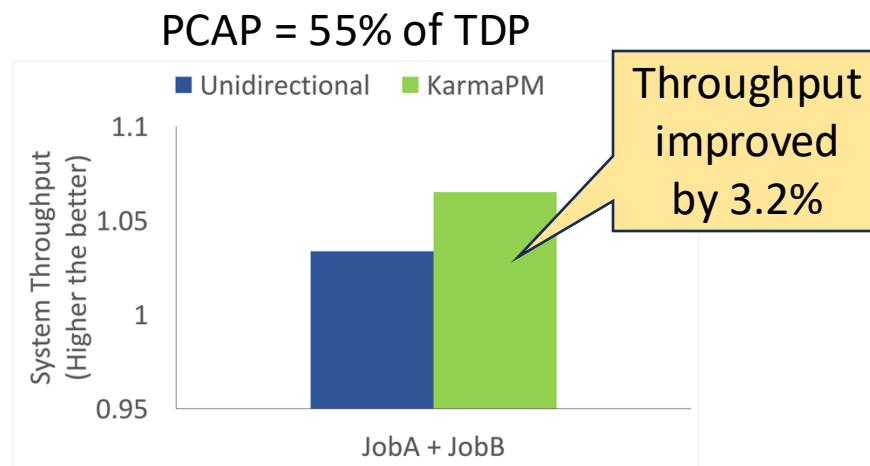
Hardware platform

- ✓ Quad socket Intel Xeon Cooper Lake
- ✓ **18** cores per socket
- ✓ TDP per socket = 150 Watts

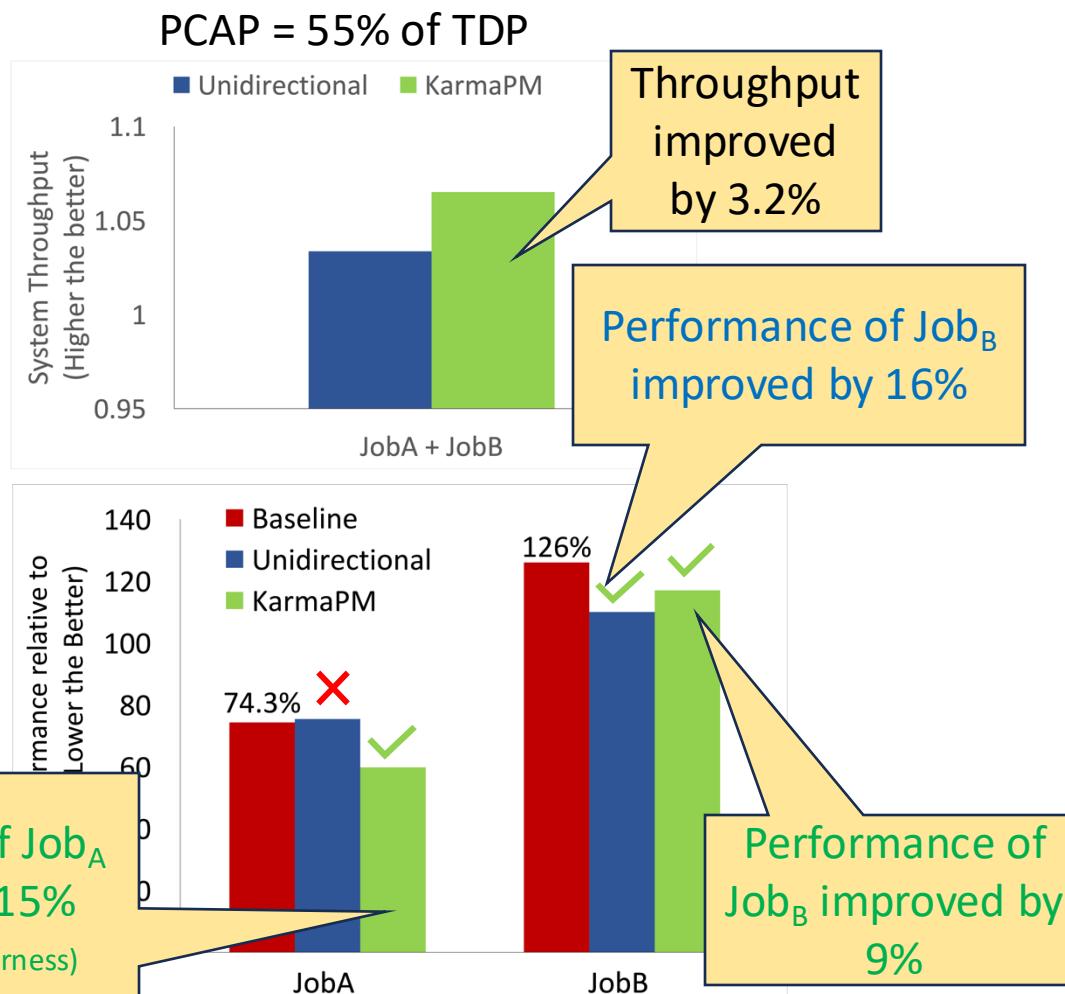
Mix Type	Number of Mixes	Socket Binding
4 Applications	5	Each application uses one socket
2 Applications	1	Each application uses two socket

Evaluated using three PCAP settings, 55%, 65%, and 75% of TDP

Throughput and Fairness from KarmaPM

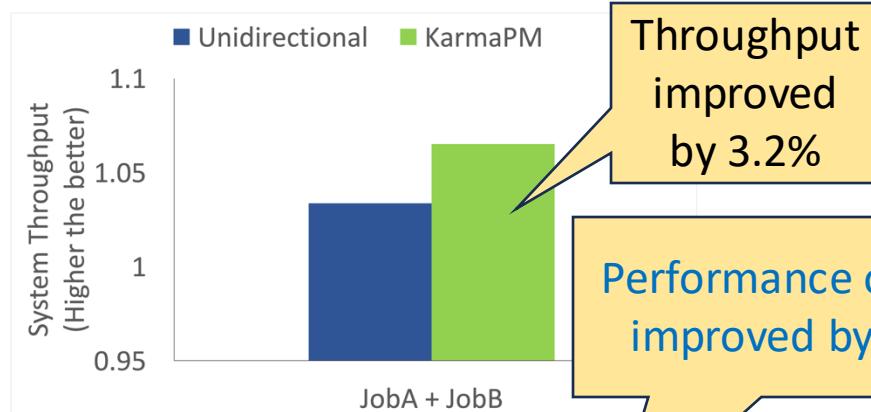


Throughput and Fairness from KarmaPM



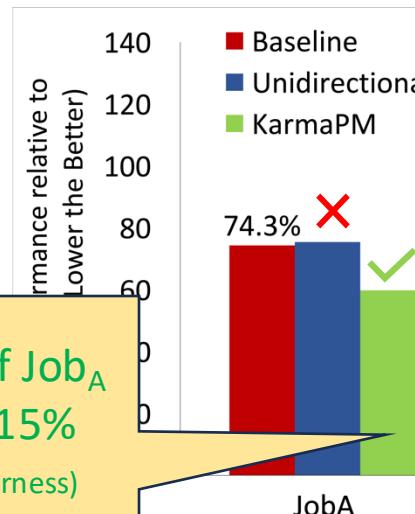
Throughput and Fairness from KarmaPM

PCAP = 55% of TDP



Throughput improved by 3.2%

Performance of Job_B improved by 16%



Performance of Job_A improved by 15% (application-level-fairness)

126%

Performance of Job_B improved by 9%

126%

Performance of Job_A improved by 4% (application-level-fairness)

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

0

140

120

100

80

60

40

20

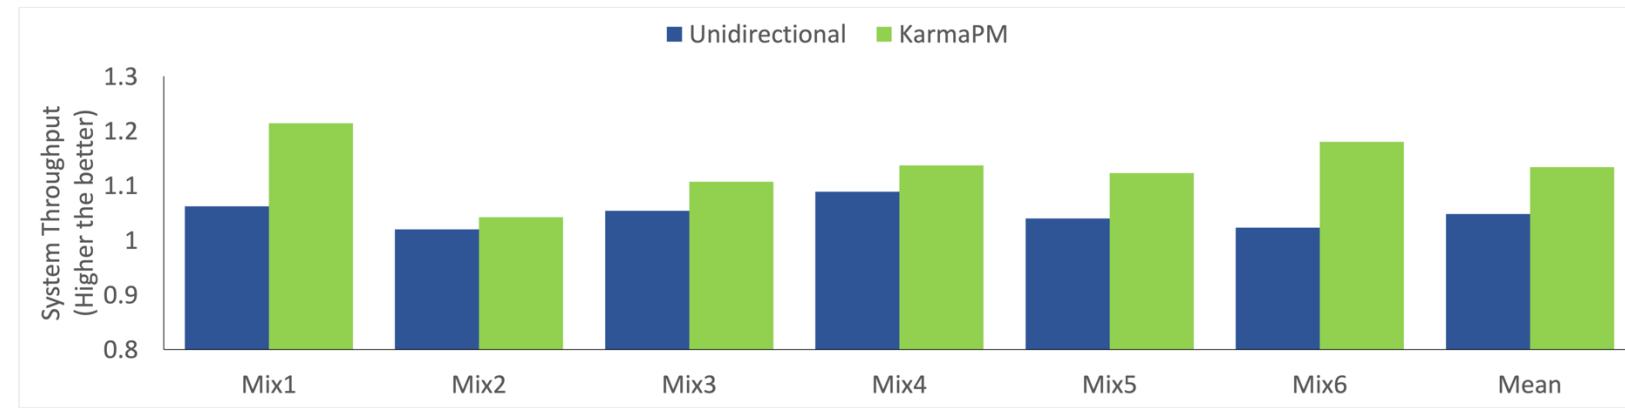
0

140

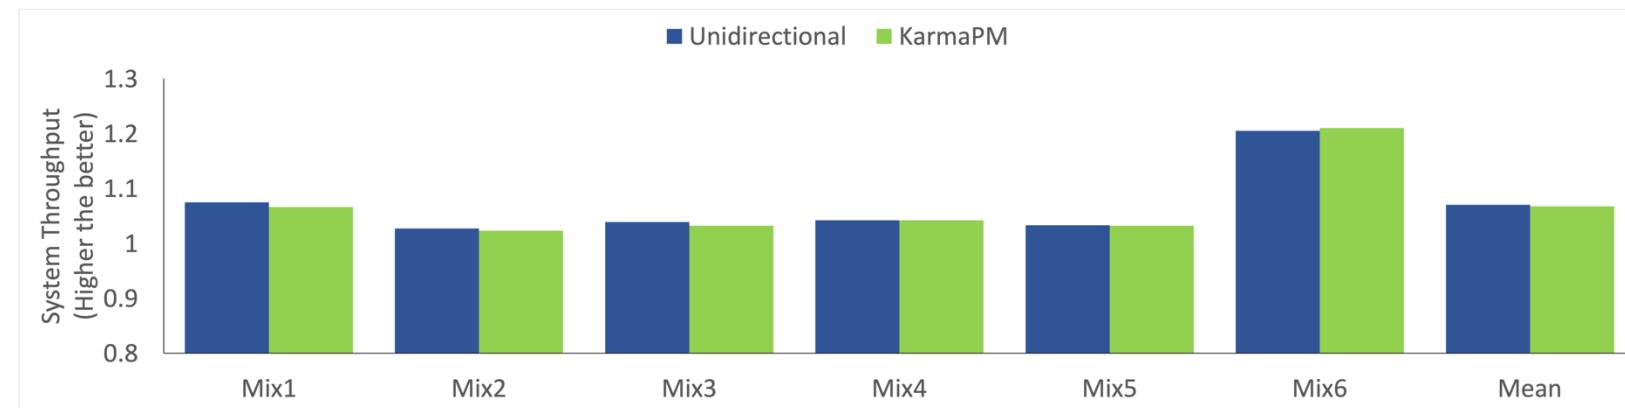
120

100

System Throughput from KarmaPM



PCAP=55% of TDP



PCAP=65% of TDP

- KarmaPM improves both throughput and fairness at low PCAP
- At higher PCAPs, KarmaPM improves fairness without affecting throughput

Conclusion and Future Work

Summary

- Hardware overprovisioning using power capping addresses the increasing computing power demand
 - However, PCAP degrades the application performance
- Running applications in pairs on a single server provides an opportunity to reduce power wastage by transferring unused power from one application to other
 - However, this approach does not support application-level fairness
- **KarmaPM** uses a novel reward-driven bi-directional power transfer mechanism that improves both throughput and application-level fairness
- In future, we plan to extend KarmaPM for heterogeneous architecture (CPU+GPU)

Q&A

Acknowledgement

This research is supported by Google PhD Fellowship 2022