

# Energy-Aware Runtime Resource Harmonizer for Co-running Applications

Vanshika Jain<sup>1</sup>, Varun Parashar<sup>1</sup>, Vivek Kumar<sup>1</sup>, Chiranjib Sur<sup>2</sup>

1 IIIT Delhi, New Delhi, India

2 Shell India Markets Pvt. Ltd, India



# Outline

- ▶ Introduction
- ▶ Motivation
- ▶ Existing Approaches
- ▶ Contributions
- ▶ Implementation
- ▶ Results
- ▶ Conclusion

# Resource Utilization in the Exascale Era

**Increasing number of sockets  
and cores per node**

Rank of Top500 (November 2025)	Sockets Per Node	Cores Per Node
1	<b>4</b>	96
2	1	64
3	2	104
4	<b>4</b>	<b>288</b>
5	2	96

1. <https://top500.org/lists/top500/>




# Resource Utilization in the Exascale Era

## Increasing number of sockets and cores per node

Rank of Top500 (November 2025)	Sockets Per Node	Cores Per Node
1	<b>4</b>	96
2	1	64
3	2	104
4	<b>4</b>	<b>288</b>
5	2	96

## Power usage at supercomputers



1. <https://top500.org/lists/top500/>


2. Patki et.al. [ICS2025]

# Resource Utilization in the Exascale Era

## Increasing number of sockets and cores per node

Rank of Top500 (November 2025)	Sockets Per Node	Cores Per Node
1	<b>4</b>	96
2	1	64
3	2	104
4	<b>4</b>	<b>288</b>
5	2	96

## Power usage at supercomputers




It is critical to improve resource utilization for achieving energy efficiency

1. <https://top500.org/lists/top500/>


2. Patki et.al. [ICS2025]

# Improving Resource Utilization via Co-location




Applications A, B & C to be executed on a quad-socket system

# Improving Resource Utilization via Co-location




# Improving Resource Utilization via Co-location




Application C  
waiting for the CPUs

# Improving Resource Utilization via Co-location




Each application completed their execution one by one

# Improving Resource Utilization via Co-location




Batch Execution




Co-running Execution

Threads of Application  
A, B & C  
running in parallel


# Improving Resource Utilization via Co-location




# Improving Resource Utilization via Co-location



# Improving Resource Utilization via Co-location




# Improving Resource Utilization via Co-location





Type: **Block-Cyclic**

# Improving Resource Utilization via Co-location





Type: **Block-Cyclic**



# Improving Resource Utilization via Co-location



# Improving Resource Utilization via Co-location






# Improving Resource Utilization via Co-location






Type: **Interleaved**

# Improving Resource Utilization via Co-location





A  
B  
C

# Improving Resource Utilization via Co-location




Type: **Block-Interleaved**

# Improving Resource Utilization via Co-location




Choosing optimal thread placement over Batch execution improves EDP by up to 81%

# Achieving Energy Efficiency on Multicores




# Achieving Energy Efficiency on Multicores





- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level

# Achieving Energy Efficiency on Multicores



- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level
- Uncore Frequency Scaling (UFS)
  - Socket-level



# Achieving Energy Efficiency on Multicores



- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level
- Uncore Frequency Scaling (UFS)
  - Socket-level

Heatmap represents the change in EDP with a particular combination of core-uncore frequency relative to default settings



# Achieving Energy Efficiency on Multicores



- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level
- Uncore Frequency Scaling (UFS)
  - Socket-level

Heatmap represents the change in EDP with a particular combination of core-uncore frequency relative to default settings



# Achieving Energy Efficiency on Multicores



- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level
- Uncore Frequency Scaling (UFS)
  - Socket-level

Heatmap represents the change in EDP with a particular combination of core-uncore frequency relative to default settings


# Achieving Energy Efficiency on Multicores



- Dynamic Voltage and Frequency Scaling (DVFS)
  - Core-level
- Uncore Frequency Scaling (UFS)
  - Socket-level


Heatmap represents the change in EDP with a particular combination of core-uncore frequency relative to default settings

# Achieving Energy Efficiency on Multicores





- Dynamic Concurrency Throttling (DCT)
  - Adjusts the application level parallelism by controlling core allocation

# Achieving Energy Efficiency on Multicores



- Dynamic Concurrency Throttling (DCT)
  - Adjusts the application level parallelism by controlling core allocation
  - Thread packing and unpacking technique provides runtime independence



# Achieving Energy Efficiency on Multicores



- Dynamic Concurrency Throttling (DCT)
  - Adjusts the application level parallelism by controlling core allocation
  - Thread packing and unpacking technique provides runtime independence

**Heatmap represents the change in EDP by changing the core count relative to default with maximum core allocation**

# Achieving Energy Efficiency on Multicores



- Dynamic Concurrency Throttling (DCT)
  - Adjusts the application level parallelism by controlling core allocation
  - Thread packing and unpacking technique provides runtime independence

**Heatmap represents the change in EDP by changing the core count relative to default with maximum core allocation**


# Insights

- Choosing optimal thread placement improves resource utilization for co-running applications
- There is a strong correlation between application behavior and resource requirement

**Cache-Sensitive and Neutral**



**Memory-bound**



# Existing Approaches for Co-Running Applications

Categories of Resource Management Techniques	DCT only	DVFS only	UFS only	DVFS+ UFS	DCT+ DVFS+ UFS
Thread Placement for contention reduction	✓	✗	✗	✗	✗
Runtime Oblivious	✓	✓	✗	✗	✗
Model Free	✓	✓	✓	✓	✗

# Existing Approaches for Co-Running Applications


Categories of Resource Management Techniques	DCT only	DVFS only	UFS only	UFS	DCT+ DVFS+ UFS
Thread Placement for contention reduction	✓	✗	✗	✗	✗
Runtime Oblivious	✓	✓	✗	✗	✗
Model Free	✓	✓	✓	✓	✗

Our Focus

# Contributions

- ✓ Harmonizer: A library-based resource management framework for co-running applications on multicore multi-socket servers
  - ✓ Model-free and runtime oblivious
- ✓ Dynamically manages thread placement, core frequency, uncore frequency and core allocation
  - ✓ Uses a lightweight daemon for online profiling of hardware PMCs
- ✓ Experimental Evaluations on a quad-socket 72-core Intel Xeon processor
  - ✓ Using several exascale proxy applications (OpenMP, Kokkos and HCLib)
- ✓ Results
  - ✓ Demonstrating substantially energy savings and performance gains

## High-level Architecture of Harmonizer




## High-level Architecture of Harmonizer

A


B

C




Harmonizer  
Daemon  
Process


# High-level Architecture of Harmonizer




# High-level Architecture of Harmonizer




# High-level Architecture of Harmonizer



# High-level Architecture of Harmonizer




# Harmonizer Policy



- Classify each application's Memory Access Pattern (**MAP**)
  - Core-level PMCs
    - Cache misses
    - Cache accesses
  - Uncore PMCs (Socket-level)
    - Integrated Memory Controller (IMC) accesses

# Harmonizer Policy





A

B

C

- Classify each application's Memory Access Pattern (**MAP**)
  - Core-level PMCs
    - Cache misses
    - Cache accesses
  - Uncore PMCs (Socket-level)
    - Integrated Memory Controller (IMC) accesses

# Harmonizer Policy



- Classify each application's Memory Access Pattern (**MAP**)
  - Core-level PMCs
    - Cache misses
    - Cache accesses
  - Uncore PMCs (Socket-level)
    - Integrated Memory Controller (IMC) accesses

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



Optimal placement for a particular mix

**Cache Sensitive – Neutral – Cache Sensitive**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



Optimal placement for a particular mix

**Cache Sensitive – Neutral – Cache Sensitive  
(Block-Cyclic to minimize LLC sharing)**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



**UFS is used to explore optimal UF**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



**UFS exploration possible only over two sockets in this mix because UFS can be applied at socket-level**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



Harmonizer rearranges threads over sockets to maximize application isolation while retaining the behaviour of Block-cyclic placement

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



**UFS exploration now possible on three sockets instead of two**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



UFS exploration now possible on three sockets instead of two

**Reduced exploration space based on MAP identified over each socket**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



DVFS is used to explore optimal CF for each application

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



DVFS is used to explore optimal CF for each application

**Reduced exploration space based on MAP identified for each application**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS



Explore DCT



**Uniform DVFS settings on each core of sockets hosting a single application**

# Harmonizer Policy

Identify MAP of each Co-Runners



Set Optimal Thread Placement




Explore UFS



Explore DVFS




Explore DCT




**Non-uniform DVFS setting at socket hosting multiple application's threads**

# Harmonizer Policy



# Harmonizer Policy



# Experimental Methodology

## Exascale proxy applications

Type of Applications	Application
Cache Sensitive	SimpleMOC (OpenMP) MinTally (OpenMP) XSbench (OpenMP)
Memory Bound	HPCCG (OpenMP) MiniFE ( <b>Kokkos</b> )
Neutral	CoHMM ( <b>HCLib</b> ) CoMD (OpenMP)

# Experimental Methodology

## Exascale proxy applications

Type of Applications	Application
Cache Sensitive	SimpleMOC (OpenMP) MinTally (OpenMP) XS Bench (OpenMP)
Memory Bound	HPCCG (OpenMP) MiniFE ( <b>Kokkos</b> )
Neutral	CoHMM ( <b>HCLib</b> ) CoMD (OpenMP)

Number of Applications in a Mix	Number of Mixes
3	6
4	3

# Experimental Methodology

## Exascale proxy applications

Type of Applications	Application
Cache Sensitive	SimpleMOC (OpenMP) MinTally (OpenMP) XS Bench (OpenMP)
Memory Bound	HPCCG (OpenMP) MiniFE ( <b>Kokkos</b> )
Neutral	CoHMM ( <b>HCLib</b> ) CoMD (OpenMP)

Number of Applications in a Mix	Number of Mixes
3	6
4	3

## Hardware Platform

- Quad socket Intel Xeon 5318H Cooper Lake
- 18 cores per socket, Total 72 cores (144 CPUs)

# Experimental Methodology

## Exascale proxy applications

Type of Applications	Application
Cache Sensitive	SimpleMOC (OpenMP) MinTally (OpenMP) XS Bench (OpenMP)
Memory Bound	HPCCG (OpenMP) MiniFE ( <b>Kokkos</b> )
Neutral	CoHMM ( <b>HCLib</b> ) CoMD (OpenMP)

Number of Applications in a Mix	Number of Mixes
3	6
4	3


## State-of-the-Art used for comparison

- Mapper (TACO'22)
- NuPoCo (PACT'18)


## Hardware Platform

- Quad socket Intel Xeon 5318H Cooper Lake
- 18 cores per socket, Total 72 cores (144 CPUs)

# EDP of Harmonizer Relative to Default




# EDP of Harmonizer Relative to Default



Harmonizer Policy	Mean EDP Improvement (Mix1- Mix3)
Thread Placement	7.3%
UFS	3%
DVFS	N/A
DCT	N/A

Improvement in EDP from individual policies


# EDP of Harmonizer Relative to Default



Harmonizer Policy	Mean EDP Improvement (Mix4- Mix6)
Thread Placement	26.6%
UFS	N/A
DVFS	3%
DCT	13.7%

Improvement in EDP from individual policies


# EDP of Harmonizer Relative to Default



Harmonizer Policy	Mean EDP Improvement (Mix7- Mix9)
Thread Placement	14.6%
UFS	N/A
DVFS	3.2%
DCT	9.3%

Improvement in EDP from individual policies

# System Throughput Relative to Default



## System Throughput

Geometric mean of speedup  
of each application

## Summary

- Effective system utilization is key to improving energy efficiency in the exascale era
  - Co-running applications can improve system utilization by complementing each other's resource requirements
- Harmonizer dynamically profiles the core and uncore PMCs to characterize the behaviour of co-running applications
  - It applies optimal thread placement for improving the system utilization
  - Dynamically tunes each socket's core and uncore frequencies, and application level core allocation to enhance energy efficiency
- Future Work
  - We plan to extend Harmonizer to handle dynamically varying memory access patterns in applications and scale it to cluster-level environments

# Thank You



Scan to access the Harmonizer artifact